


FROM STOCHASTIC CHAOS TO DETERMINISTIC CERTAINTY

AI for Critical Industrial Infrastructure

STOCHASTIC FIELD · DETERMINISTIC ISLANDS



DETERMINISTIC ISLAND · AUDITABLE BOUNDARY

GRIGORIOS AGATHANGELIDIS

NEXUS-1 Series · Critical infrastructure, AI, and deterministic engineering

From Stochastic Chaos to Deterministic Certainty

AI for Critical Industrial Infrastructure

Grigorios Agathangelidis

Target Audience: Engineers, Data Scientists, Physicists, and Technology
Executives with
University-Level Education in the Exact Sciences

Standing Rule

STANDING RULE - NEXUS-1 SERIES

Nothing in this book claims to be exact that is not. When a result is an approximation, the approximation is named. When a model is a stepping stone, the step is named. When a number is measured rather than derived, its source is named.

- Standing Rule, NEXUS-1 Series

© Grigorios Agathangelidis. All rights reserved.

Companion volumes: From Grid to Core; From Core to Quantum.

The Bridge from Quantum to Intelligence

This volume is the third in the NEXUS-1 series. Where the first book, *From Grid to Core*, walked inward from the switchyard to the reactor core, and the second, *From Core to Quantum*, continued inward to the quantum rules that govern matter and back out to software, this volume walks from the quantum foundations of information to the highest expression of computation applied to critical systems: Artificial Intelligence that can be trusted where failure is not an option.

The central thesis of this book is that AI in critical mission applications cannot be a black box of probabilities. It must be built upon **Deterministic Islands** — points of reference where behaviour is reproducible, auditable, and certifiable, even while the broader system operates in a stochastic regime. This philosophy permeates every chapter that follows.

The arc of the series is now visible: from the grid, to the core, to the quantum world beneath it, to the intelligence that lets us model and control any of it.

CONTENTS

Introduction: The Bridge from Quantum to Intelligence	iii
I Foundations — Information, Uncertainty, Probability	1
1 Information, Uncertainty, and the Shannon Limit	2
1.1 What You Already Know	2
1.2 The Birth of Information Theory	2
1.2.1 The Problem of Compression and Transmission	2
1.2.2 Information as “Surprise”	2
1.2.3 Entropy	3
1.3 The Shannon Theorem — A Detailed Treatment	4
1.3.1 The Noisy Channel Coding Theorem	4
1.3.2 Interpretation	4
1.3.3 The Proof, Skeletally	5
1.3.4 The Converse Theorem	5
1.3.5 Worked Examples: AWGN, BSC, and the Shannon Limit in dB . .	5
1.3.6 Connection to Critical Infrastructure	6
1.4 Rate-Distortion Theory	6
1.4.1 The Theorem	7
1.4.2 Lossy Compression	7
1.4.3 Connection to Feature Extraction	7
1.5 The Shannon Theorem as the Forerunner of AI	7
1.5.1 Entropy as a Loss Function	7
1.5.2 Mutual Information as a Feature Selection Criterion	8
1.5.3 The Channel as a Neuron Model	8
1.5.4 The Shannon Limit as a Learning Limit	8
1.5.5 The Information Bottleneck	8
1.6 Worked Example: Shannon in an Industrial Sensor	9
2 Probability and Statistics — A Staged Introduction	11
2.1 What You Already Know	11
2.2 Level 1 — Intuition: Sample Spaces, Events, and Kolmogorov’s Axioms	11
2.3 Level 2 — Conditional Probability and Bayes’ Theorem	12
2.4 Level 3 — Random Variables and Distributions	13
2.5 Level 4 — Multivariate Distributions	14
2.6 Level 5 — Statistical Inference	14
2.7 Level 6 — Hypothesis Testing	14
2.8 Level 7 — Stochastic Processes	15
2.9 Level 8 — Information Theory Meets Statistics	15
2.10 Level 9 — Towards Machine Learning	16
2.11 Worked Examples	17

3	From Data to Models	19
3.1	What You Already Know	19
3.2	Maximum Likelihood and Maximum A Posteriori Estimation, Applied	19
3.3	The Bias-Variance Tradeoff	20
3.4	Cross-Validation and Regularization	21
3.4.1	Cross-Validation	21
3.4.2	Regularization	21
3.4.3	The Closed-Form Ridge Solution	21
3.4.4	Why L1 Produces Sparsity: The Geometric Picture	22
3.5	The Minimum Description Length Principle	22
3.6	Learning as Compression: The Shannon Connection Revisited	23
3.7	The Curse of Dimensionality and the No Free Lunch Theorem	24
3.7.1	The Curse of Dimensionality	24
3.7.2	The No Free Lunch Theorem	25
4	Towards Machine Learning	26
4.1	What You Already Know	26
4.2	Linear Regression as a Single Neuron	26
4.3	Logistic Regression as a Classifier	27
4.3.1	The Log-Odds Interpretation	28
4.3.2	The Cross-Entropy Loss, Concretely	29
4.4	Gradient Descent: From Calculus to Optimization	29
4.4.1	Batch, Stochastic, and Mini-Batch Gradient Descent	30
4.5	The Problem of Non-linearity: Why a Single Layer Is Not Enough	31
4.6	Introduction to Neural Networks: The Perceptron as a Building Block	32
II	Neural Networks — From Perceptron to Deep Networks	34
5	The Perceptron — The Foundation	35
5.1	What You Already Know	35
5.2	The Biological Neuron	35
5.3	The McCulloch-Pitts Model	36
5.4	The Artificial Neuron (Perceptron)	36
5.4.1	Inputs, Weights, and Bias	36
5.4.2	Linear Combination	36
5.4.3	Activation: The Step Function	36
5.4.4	Output	36
5.5	The Perceptron Learning Algorithm	37
5.5.1	The Update Rule	37
5.5.2	Geometric Interpretation: Rotating the Hyperplane	37
5.5.3	The Perceptron Convergence Theorem	37
5.6	Worked Example: Linear Separation in Industrial Data	38
5.7	The Limitation: The XOR Problem	39
5.7.1	The Perceptron Cannot Solve XOR	40
5.7.2	Proof: XOR Is Not Linearly Separable	40
5.7.3	Geometric Interpretation, Restated	40
5.7.4	Historical Note: Minsky, Papert, and the First AI Winter	40

6	Solving XOR — Hidden Layers and Non-linearity	42
6.1	What You Already Know	42
6.2	The Solution to XOR	42
6.2.1	XOR as a Composition of Simpler Gates	42
6.2.2	Two-Layer Implementation	43
6.2.3	Visualization: Transforming the Input Space	44
6.3	Non-linear Activation Functions	44
6.3.1	Why Non-linearity Is Essential	45
6.4	The Universal Approximation Theorem	45
6.4.1	Statement	45
6.4.2	Proof Sketch, via Step Functions	45
6.5	Worked Example: XOR in C#	46
6.6	From XOR to Deep Networks	48
6.6.1	Hierarchical Feature Learning	48
6.6.2	Each Layer, a More Abstract Representation	49
6.6.3	Examples: From Pixels to Objects	49
7	Backpropagation — The Engine of Learning	50
7.1	What You Already Know	50
7.2	The Backpropagation Algorithm	50
7.2.1	Forward Pass: Computing the Output	50
7.2.2	Loss Functions: MSE and Cross-Entropy, Recalled	50
7.2.3	Backward Pass: Computing Gradients	51
7.2.4	Weight Update	51
7.3	Detailed Derivation for a 2-2-1 Network	51
7.3.1	Step-by-Step Gradient Computation	51
7.3.2	Table of Variables and Partial Derivatives	52
7.4	Worked Example: Training XOR with Backprop in C#	52
7.4.1	Code, Annotated	52
7.4.2	Visualization of the Loss Evolution	54
7.4.3	Final Weight Values	54
7.5	Challenges	55
7.5.1	Vanishing and Exploding Gradients	55
7.5.2	Local Minima	55
7.5.3	Overfitting	56
7.5.4	Learning Rate Selection	56
7.6	Optimization Techniques	57
7.6.1	Momentum and Nesterov Momentum	57
7.6.2	AdaGrad, RMSProp, and Adam	57
7.6.3	Batch Normalization	57
7.6.4	Dropout	57
8	Modern Architectures	59
8.1	What You Already Know	59
8.2	Convolutional Neural Networks (CNNs)	59
8.2.1	Locality, Weight Sharing, and Pooling	59
8.2.2	Applications: Image and Signal Processing in Industry	61
8.3	Recurrent Neural Networks (RNNs)	61
8.3.1	Sequences, Hidden State, and Backpropagation Through Time	61
8.3.2	LSTM and GRU: Gates and Memory	61
8.3.3	Applications: Time-Series Prediction and Anomaly Detection	62
8.4	The Attention Mechanism	62

8.4.1	The Key-Value-Query Formulation	62
8.4.2	Self-Attention and Multi-Head Attention	63
8.4.3	Positional Encoding	63
8.5	The Transformer Architecture	64
8.5.1	Encoder-Decoder Structure	64
8.5.2	Why Transformers Dominate Modern AI	64
8.5.3	Applications: Language Models, Time-Series, and Multimodal Data	64
8.6	Autoencoders and Variational Autoencoders (VAEs)	65
8.7	Generative Adversarial Networks (GANs)	65
8.8	Architecture Selection for Industrial Data	66
9	Supervised, Unsupervised, and Reinforcement Learning	67
9.1	What You Already Know	67
9.2	Supervised Learning	67
9.3	Unsupervised Learning	68
9.3.1	Clustering	68
9.3.2	Dimensionality Reduction	69
9.3.3	Anomaly Detection	69
9.4	Reinforcement Learning	70
9.4.1	Markov Decision Processes	70
9.4.2	Q-Learning and Policy Gradients	70
9.4.3	Applications in Control: When RL Is Appropriate, and When It Is Dangerous	71
9.5	Semi-supervised and Self-supervised Learning	71
9.6	Transfer Learning and Fine-Tuning	72
9.7	Applications in Critical Infrastructure	72
III	The Industry Language Model (ILM)	74
10	Why an ILM?	75
10.1	What You Already Know	75
10.2	The Limits of General LLMs in Critical Applications	75
10.2.1	Hallucinations: The Model "Confabulates" Facts	75
10.2.2	Lack of Domain Knowledge	75
10.2.3	Non-Deterministic Behaviour	76
10.2.4	Inability to Verify	76
10.2.5	Safety: One Mistake Can Be Catastrophic	77
10.3	Definition of the ILM (Industry Language Model)	77
10.3.1	Data	77
10.3.2	Goal	77
10.3.3	Requirements: Reliability, Verifiability, Determinism Where Needed	77
10.4	Architecture of the ILM	77
10.4.1	Transformer-Based	78
10.4.2	Adaptations	78
10.5	Training Data	80
10.5.1	Sources	80
10.5.2	Preprocessing	80
10.5.3	Handling Sparse Data and Class Imbalance	80
10.5.4	Data Ethics and Security	81

11 Training and Adaptation of the ILM	82
11.1 What You Already Know	82
11.2 Pre-training	82
11.3 Masked Language Modeling and Causal Language Modeling	82
11.4 Fine-Tuning	83
11.5 Parameter-Efficient Fine-Tuning: LoRA and Adapters	83
11.6 Retrieval-Augmented Generation	84
11.7 Prompt Engineering for Industrial Contexts	85
11.8 Evaluation: Perplexity, BLEU, ROUGE, and Human Evaluation	85
11.9 Ethical Issues: Bias, Fairness, and Transparency	86
12 Deterministic Islands	88
12.1 What You Already Know	88
12.2 The Problem	88
12.2.1 Neural Networks Are Stochastic	88
12.2.2 Two Identical Inputs, Two Different Outputs	88
12.2.3 In Critical Applications, This Is Unacceptable	88
12.3 The Deterministic Islands	89
12.3.1 Static Vault: Hash-Based Routing	89
12.3.2 Frozen Snapshot: Locked Software, Seed, and CPU Inference	90
12.3.3 Intra-Vector Routing: Islands the Model Creates for Itself	91
12.4 Architecture of the Islands	93
12.4.1 How They Are Integrated into the ILM	93
12.4.2 Code Examples: The Arbiter	93
12.4.3 Conflict Resolution Between Islands	94
12.4.4 Updating Islands: When and How	94
12.5 Worked Example: An Island in Industrial Control	95
13 Evaluation and Verification of the ILM	96
13.1 What You Already Know	96
13.2 Test Sets, Validation Sets, Cross-Validation	96
13.3 Metrics: Accuracy, Precision, Recall, F1, and AUC-ROC	96
13.4 Robustness: Adversarial Examples and Distribution Shift	98
13.5 Explainability: SHAP, LIME, and Attention Visualization	99
13.6 Uncertainty Quantification	99
13.7 Formal Verification: Specifications and Model Checking	100
13.8 Certification: Standards	100
14 Ethics, Safety, and Governance	102
14.1 What You Already Know	102
14.2 Bias and Fairness	102
14.3 Transparency and Explainability	103
14.4 Security: Adversarial Attacks and Data Poisoning	103
14.5 Privacy: Federated Learning and Differential Privacy	104
14.6 Regulatory Framework	104
14.7 Human Oversight	105
14.8 Responsibility and Accountability	105
IV Applications in Critical Infrastructure	107
15 Predictive Maintenance	108
15.1 What You Already Know	108

15.2	Data: Vibration, Temperature, Pressure, Current	108
15.3	Feature Engineering: Time-Domain and Frequency-Domain	109
15.3.1	Time-Domain Features	109
15.3.2	Frequency-Domain Features	110
15.4	Models: RNN, LSTM, and Transformer	110
15.5	Remaining Useful Life (RUL) Estimation	110
15.6	Examples: Turbines, Pumps, and Transformers	111
15.7	Predictive Maintenance Inside the ILM	111
16	Anomaly Detection	113
16.1	What You Already Know	113
16.2	What Is an Anomaly? Point, Contextual, Collective	113
16.3	Methods: Statistical, Distance-Based, Density-Based, Reconstruction-Based	113
16.3.1	Statistical Methods	114
16.3.2	Distance-Based and Density-Based Methods	114
16.3.3	Reconstruction-Based Methods	114
16.4	Autoencoders for Anomaly Detection	115
16.5	One-Class SVM and Isolation Forest	115
16.6	Deep Learning: GANs and VAEs	116
16.7	Evaluation: Precision@k and AUPRC	116
16.8	Examples: Leak Detection, Sensor Failure, and Cyberattack	117
16.9	Anomaly Detection Inside the ILM	117
17	Control and Optimization	119
17.1	What You Already Know	119
17.2	Model Predictive Control (MPC) with Neural Networks	119
17.3	Reinforcement Learning for Control	120
17.4	Safe RL: Constrained RL and Shielding	120
17.5	Digital Twins	122
17.6	Examples: Reactor Control, Combustion, and Energy Management	122
18	Decision Support	123
18.1	What You Already Know	123
18.2	Decision Support Systems	123
18.3	Human-AI Collaboration	124
18.4	Explainable AI for Decisions	124
18.5	Scenario Analysis	124
18.6	Examples: Fault Diagnosis, Maintenance Planning, and Crisis Management	125
19	Safety and Compliance	127
19.1	What You Already Know	127
19.2	Regulations: Nuclear (IAEA), Aerospace (FAA, EASA), Medical (FDA)	127
19.3	Certification of AI Systems	128
19.4	Safety Cases	128
19.5	Auditing and Traceability	129
19.6	Examples: AI Certification for Medical Devices and Autonomous Vehicles	129


V Software Modeling – DDD in C#	131
20 Introduction to Software Modeling for Engineers	132
20.1 What You Already Know	132
20.2 Why Model Physics and AI with Code?	132
20.2.1 Simulations and Digital Twins	132
20.2.2 Educational Tools and Interactive Learning	132
20.2.3 Data Analysis and Experimental Validation	132
20.2.4 Control Automation and Safety	133
20.3 Basic Principles of Object-Oriented Programming (OOP) in C#	133
20.4 Introduction to Domain-Driven Design (DDD)	134
20.4.1 Ubiquitous Language	134
20.4.2 Bounded Contexts	134
20.4.3 Entities vs. Value Objects	134
20.4.4 Aggregates and Aggregate Roots	135
20.4.5 Repositories	136
20.4.6 Domain Services	136
20.4.7 Domain Events	136
20.5 The Translation Table: Physics → DDD	137
21 Bounding the Quantum Context (Revisited)	138
21.1 What You Already Know	138
21.2 Value Objects: StateVector, Operator, Basis	138
21.3 Entities: Qubit and QuantumSystem	139
21.4 Modeling Quantum Collapse and Entanglement	140
21.5 Code Example: Gate Application, Measurement, and Time Evolution	141
22 Bounding the AI Context	143
22.1 What It Includes: Models, Training, Inference, Prompts, Deterministic Islands	143
22.2 Domain Terminology	143
22.3 Value Objects	143
22.4 Entities	144
22.5 Aggregate and the ILM Aggregate Root	145
22.6 Domain Services	146
22.7 Domain Events	147
22.8 Code Example: Building a Minimal ILM in C#	148
23 Integration and Application Architecture	150
23.1 Connecting the Three Contexts	150
23.2 Use of Repositories	150
23.3 Application Creation	151
23.3.1 Application Layer: Coordination of Domain Services	151
23.3.2 Infrastructure Layer: Repository Implementation, Logging, Monitoring	152
23.3.3 Presentation Layer: CLI or Web API	152
23.4 Design Patterns Used	152
23.5 Creating a Console or API for Interacting with the Model	153
23.6 Testing	154
23.7 Future Extensions: Digital Twin of the Reactor	155

VI Towards Artificial Certainty	157
24 Deterministic Islands — Deep Dive	158
24.1 What You Already Know	158
24.2 Mathematical Foundations	158
24.2.1 Hashing: SHA-256 and Collision Resistance	158
24.2.2 Reproducibility: Seeds, Determinism, and Bitwise Accuracy	159
24.2.3 Formal Definition of a Deterministic Island	159
24.3 Implementation, Revisited Through This Lens	160
24.4 Management	160
24.4.1 Updating and Extending Islands	160
24.4.2 Conflict Resolution, Recalled	160
24.5 Code Examples in C#	161
24.6 Case Study: A Deterministic Island for Reactor Safety	161
25 Governance and Regulation	163
25.1 What You Already Know	163
25.2 Regulatory Frameworks, Mapped to This Book's Architecture	163
25.2.1 EU AI Act	163
25.2.2 NIST AI Risk Management Framework	163
25.2.3 Sector-Specific Regulations	164
25.3 Responsibility and Accountability, Made Queryable	164
25.4 Ethical Issues, Extended: Autonomous Decision-Making	165
26 The Future — Artificial Certainty	166
26.1 What You Already Know	166
26.2 The Vision: AI That Is Not Merely Intelligent, but Reliable	166
26.3 Deterministic Islands as the Foundation of Trustworthy AI	166
26.4 The Neural Constitution: Rules, Constraints, Values	167
26.5 Ethical AI	168
26.6 The Future of Critical Infrastructure	168
26.7 Quantum Neural Networks with Deterministic Islands	168
26.7.1 Introduction to Quantum Neural Networks (QNNs)	168
26.7.2 The Challenge of Noise and Decoherence	168
26.7.3 Deterministic Islands in the Quantum Regime	169
26.7.4 A Hybrid Architecture: Classical, Quantum, and Deterministic	169
26.7.5 Potential Applications	170
26.7.6 The Road Ahead	170
26.8 The Long View: From Artificial Intelligence to Artificial Certainty	171
Epilogue: From Certainty to Responsibility	172
Appendix A.1: Mathematical Review	174
Appendix A.2: Table of Physical Constants	176
Appendix A.4: Keepin Six-Group Delayed Neutron Data	177
Appendix A.5: Architecture of a Deterministic Island	178
Appendix A.6: Comparison Table of AI Paradigms in Critical Systems	179

Appendix A.7: Installation and Execution Guide for the Code (C#)	180
BACK MATTER	
Glossary of Terms (Appendix A.3)	182
Index	188
Sources and Reference Notes (Bibliography)	190

PART I

FOUNDATIONS - INFORMATION, UNCERTAINTY, PROBABILITY



FOUNDATIONS · INFORMATION, UNCERTAINTY, PROBABILITY

1. Information, Uncertainty, and the Shannon Limit

1.1 What You Already Know

Every engineer who has worked with a noisy sensor loop already has an intuition for the subject of this chapter, even without the vocabulary for it. You know that a 4–20 mA loop carries less detail than a 16-bit digital transmitter. You know that a noisy thermocouple signal needs averaging before it is trustworthy. You know that a narrow-bandwidth radio channel cannot carry a high-definition video feed, no matter how clever the modulation scheme. And you know, from bitter experience, that there is a point past which turning up the gain on a noisy signal amplifies the noise along with the signal and buys you nothing.

What you may not have is the underlying quantity that makes all of these intuitions precise: a numerical measure of *how much is knowable*, given the noise, the bandwidth, and the precision you demand. That quantity is **information**, and the man who made it rigorous was Claude Shannon. This chapter builds his theory from the ground up, in the language of an engineer, and then turns it toward the specific question this book is about: how much of what a sensor, a log file, or a fault record tells us can an AI system actually learn, and how much is fundamentally unlearnable, no matter how large the model?

1.2 The Birth of Information Theory

1.2.1 The Problem of Compression and Transmission

In 1948, working at Bell Labs, Claude Shannon published *A Mathematical Theory of Communication*. The practical problem in front of him was the telephone network: how do you send a message over a wire (or, later, a radio channel) as efficiently as possible, and how do you protect it from the noise that the wire inevitably adds? Shannon's insight was to strip the problem of everything that is not essential to it. He was not interested in the *meaning* of a message, only in how surprising, and therefore how compressible, it is. This deliberate indifference to meaning is what let him build a theory that applies with equal force to English prose, DNA sequences, sensor telemetry, and the weights of a neural network.

1.2.2 Information as “Surprise”

Shannon's first move was to define the information content of a single outcome x of a random variable X , with probability $P(x)$, as a decreasing function of that probability:

Self-Information

$$I(x) = -\log_2 P(x)$$

measured in **bits**.

This definition has three properties that any reasonable measure of “surprise” must have. An event that is certain ($P(x) = 1$) carries zero information: $-\log_2 1 = 0$. An event that is rare carries a large amount of information: as $P(x) \rightarrow 0$, $I(x) \rightarrow \infty$. And the information of two independent events, observed together, adds: $I(x, y) = -\log_2 (P(x)P(y)) = I(x) + I(y)$, which matches the everyday sense that two unrelated pieces of surprising news are more surprising, in total, than either alone.

1.2.3 Entropy

Self-information tells you the surprise of one outcome. **Entropy** is the *average* surprise across the whole distribution: it is the expected number of bits needed, on average, to describe an outcome drawn from X .

Shannon Entropy

$$H(X) = -\sum_i P(x_i) \log_2 P(x_i) = \mathbb{E}[-\log_2 P(X)]$$

Entropy is a property of the distribution, not of any single outcome. It is maximised when all outcomes are equally likely (maximum uncertainty) and is zero exactly when the outcome is certain (no uncertainty at all). This is the single most important idea in this chapter, and it will recur, in one disguise or another, in nearly every remaining chapter of this book: as the loss function that trains a neural network (Chapter 4 and Chapter 7), as the criterion behind feature selection (Section 1.5.2), and as the reason a language model’s output is fundamentally a probability distribution rather than a lookup table (Chapter 10).

Fair Coin, Loaded Coin, Deterministic Event

Consider a binary source with $P(\text{heads}) = p$ and $P(\text{tails}) = 1 - p$. Its entropy is the **binary entropy function**

$$H_b(p) = -p \log_2 p - (1 - p) \log_2 (1 - p).$$

- **Fair coin** ($p = 0.5$): $H_b(0.5) = 1$ bit. This is the maximum possible entropy for a binary source — you learn a full bit every time you observe the outcome, because you had no way to predict it.
- **Loaded coin** ($p = 0.9$): $H_b(0.9) \approx 0.469$ bits. You can usually guess “heads” and be right, so each observation carries less new information.
- **Deterministic event** ($p = 1$ or $p = 0$): $H_b(p) = 0$. There is nothing to learn; the outcome was never in doubt.

This is already the first appearance of a theme that runs through the whole book: a **Deterministic Island** (Chapter 12) is, in information-theoretic terms, a region of the system’s behaviour that has been engineered to have $H = 0$ — zero entropy, by construction, regardless of how stochastic the surrounding model is.




Figure 1.1: The binary entropy function $H_b(p)$. Entropy is symmetric around $p = 0.5$, where it peaks at exactly 1 bit, and falls to zero at the two deterministic extremes $p = 0$ and $p = 1$.

1.3 The Shannon Theorem — A Detailed Treatment

1.3.1 The Noisy Channel Coding Theorem

The result that made Shannon’s paper a foundation of the digital age is the **Noisy Channel Coding Theorem**. It answers a question every communications engineer (and, as we will see, every designer of a sensor network) needs answered: given a channel that corrupts what you send with noise, what is the highest rate at which you can send information and still recover it, in principle, with an error probability as small as you like?

Channel Capacity

$$C = \max_{P(x)} I(X; Y) = \max_{P(x)} [H(X) - H(X | Y)]$$

where $I(X; Y)$ is the **mutual information** between the channel input X and output Y , and the maximisation is over the input distribution $P(x)$.

Read this equation as a statement about how much of the uncertainty in X survives the trip through the channel. $H(X)$ is the uncertainty you start with. $H(X | Y)$ is what remains *after* you have observed the (possibly corrupted) output Y — this is sometimes called the **equivocation**. The difference is exactly how much the channel let you learn about X . Capacity is this quantity maximised over every possible way of encoding the input.

1.3.2 Interpretation

The theorem’s content is deceptively simple, and it is worth stating in words rather than symbols, because the words are what an engineer actually needs to remember: for any transmission rate $R < C$, there exists a coding scheme that achieves an error probability as small as you like. For any rate $R > C$, no coding scheme, however clever, can drive the error probability to zero. Capacity is a hard wall, not a soft recommendation. You cannot engineer around it; you can only engineer to approach it.

1.3.3 The Proof, Skeletally

A full proof belongs in an information theory textbook (see the Bibliography), but the strategy is worth sketching because the intuition it builds is reused, almost unchanged, in Chapter 13 when we discuss what it means to *verify* a trained model.

Shannon's proof is a **random coding** argument, and it is famous for being non-constructive: it proves that good codes exist without exhibiting one. The steps are:

1. Generate a codebook by choosing 2^{nR} codewords of length n at random, according to the capacity-achieving input distribution.
2. For a long block length n , the **Asymptotic Equipartition Property** guarantees that the sequences actually produced by the source concentrate on a small set of **typical sequences**, each with probability close to $2^{-nH(X)}$.
3. Two independently generated codewords are, with high probability, **jointly typical** with the received sequence only if they were actually sent; the probability that a *wrong* codeword looks jointly typical with the received sequence shrinks exponentially in n as long as $R < C$.
4. Averaging the error probability over all random codebooks shows it can be made arbitrarily small; therefore *at least one* codebook achieves it.

The proof tells you a code exists. It does not tell you how to build one — that took another fifty years of coding theory (Turbo codes, LDPC codes, polar codes) to approach in practice. This gap between *existence* and *construction* is a pattern you will see again: the Universal Approximation Theorem (Chapter 6) guarantees a neural network can represent a function without telling you how to train one that does.

1.3.4 The Converse Theorem

The converse is the other half of the wall: it proves that no scheme can beat capacity. The argument uses Fano's inequality, which lower-bounds the probability of decoding error in terms of the equivocation $H(X | Y)$. If you attempt to transmit at a rate $R > C$, Fano's inequality forces the error probability away from zero as n grows, no matter what encoder and decoder you choose. Achievability (Section 1.3.3) and the converse together are what make C an *equality*, not just a bound: it is simultaneously the best you can do and the best anyone could ever do.

1.3.5 Worked Examples: AWGN, BSC, and the Shannon Limit in dB

The AWGN Channel

For a continuous channel with additive white Gaussian noise, power constraint P , noise power N , and bandwidth B , capacity is

$$C = B \log_2 \left(1 + \frac{P}{N} \right) \text{ bits/second.}$$

This is the formula every RF and instrumentation engineer has met under the name **Shannon-Hartley**. It says capacity grows only *logarithmically* with signal-to-noise ratio (SNR), but *linearly* with bandwidth — which is why modern systems chase more spectrum far more aggressively than they chase more power.

The Binary Symmetric Channel

For a channel that flips each transmitted bit independently with probability p (the Binary Symmetric Channel, BSC), capacity is

$$C = 1 - H_b(p) \quad \text{bits per channel use,}$$

where H_b is the binary entropy function of Figure 1.1. At $p = 0$ or $p = 1$ (a perfectly reliable, if inverted, channel), $C = 1$: every bit gets through. At $p = 0.5$, the channel output is pure noise independent of the input, $H_b(0.5) = 1$, and $C = 0$: nothing gets through, no matter how cleverly you encode.

The Shannon Limit in Decibels

A classic figure of merit in communications engineering is the minimum energy-per-bit-to-noise ratio E_b/N_0 needed to communicate reliably at all, as bandwidth $B \rightarrow \infty$. Taking the AWGN capacity formula to the infinite-bandwidth limit gives the celebrated **Shannon limit**:

$$\left. \frac{E_b}{N_0} \right|_{\min} = \ln 2 \approx -1.6 \text{ dB.}$$

No communication system — regardless of modulation, coding, or cleverness — can operate reliably below this bound. Real systems (LDPC and Turbo codes) now operate within a few tenths of a dB of it. This is worth remembering the next time a vendor claims a “breakthrough” modulation scheme: the wall is a law of information, not a limitation of current technology.

1.3.6 Connection to Critical Infrastructure

Every sensor in a nuclear plant, a chemical reactor, or a power grid is, in Shannon’s sense, a noisy channel: the physical quantity is X , the digitised reading is Y , and the transmitter, cabling, and analogue-to-digital converter are the channel. Three questions from this section translate directly into plant engineering:

- **How many bits can a sensor transmit?** Bounded above by $C = B \log_2(1 + \text{SNR})$ for the physical link, and in practice further bounded by the resolution of the analogue-to-digital converter.
- **How much information is lost to noise?** Exactly the equivocation $H(X | Y)$ — the residual uncertainty about the true plant state after the noisy reading has been observed.
- **How does this bound what AI can learn?** An ILM (Chapter 10) trained on sensor data can never recover information that the sensor channel itself destroyed. This is not a modelling failure; it is a hard floor set by Section 1.3.1, and no amount of additional training data or model capacity can push through it. Data engineering (better sensors, higher sampling rates, less noisy wiring) is the only way to raise the ceiling.

1.4 Rate-Distortion Theory

The Noisy Channel Coding Theorem answers “how much can I send losslessly?” A closely related question, and the more practically important one for industrial

telemetry, is: “how much can I *throw away*, and how much distortion does that cost me?” Rate-distortion theory, also due to Shannon, answers exactly this.

1.4.1 The Theorem

Rate-Distortion Function

$$R(D) = \min_{P(\hat{x}|x): \mathbb{E}[d(x, \hat{x})] \leq D} I(X; \hat{X})$$

$R(D)$ is the minimum number of bits per symbol needed to represent the source X by a reconstruction $\hat{X}$, subject to the average distortion $d(x, \hat{x})$ (for example, squared error) staying below a budget D . It is the natural extension of channel capacity to *lossy* compression: capacity asked how much can pass through a noisy channel with zero error; rate-distortion asks how much needs to be sent to tolerate a chosen, non-zero amount of error.

1.4.2 Lossy Compression

At $D = 0$, $R(D)$ reduces to $H(X)$: to reconstruct the source perfectly, you need its full entropy, no more and no less. As you permit more distortion, $R(D)$ decreases — this is the formal reason that JPEG, MP3, and every other lossy codec can achieve compression ratios far beyond what lossless coding allows: they are deliberately trading bits for tolerable distortion, and rate-distortion theory tells you the best possible trade.

1.4.3 Connection to Feature Extraction

The same logic applies to feature extraction and dimensionality reduction, which is the everyday business of a machine learning pipeline. When you reduce a 200-channel vibration spectrum to ten summary features, you are choosing a point on that system’s own rate-distortion curve: you are accepting some reconstruction distortion (information about the raw signal you can no longer recover) in exchange for a lower “rate” (fewer numbers to store, transmit, and feed to a model). Principal Component Analysis, an autoencoder’s bottleneck layer (Chapter 4), and a hand-engineered feature set are all, from this point of view, doing the same thing: approximating an unknown rate-distortion function with a particular family of encoders.

1.5 The Shannon Theorem as the Forerunner of AI

It is not an exaggeration to say that modern machine learning is applied information theory. This section makes five of those connections explicit, each of which will be developed fully in later chapters.

1.5.1 Entropy as a Loss Function

The **cross-entropy loss**, used to train nearly every classifier in this book (Chapter 4 onward), is

$$H(P, Q) = - \sum_i P(x_i) \log_2 Q(x_i),$$

where P is the true label distribution and Q is the model's predicted distribution. Minimising cross-entropy is exactly minimising the extra bits needed to encode the true labels using a code optimised for the model's (possibly wrong) beliefs Q instead of the true distribution P . A perfectly trained classifier, in this view, is one whose beliefs require no extra bits at all.

1.5.2 Mutual Information as a Feature Selection Criterion

Mutual information $I(X; Y)$, introduced in Section 1.3.1 as the quantity a channel preserves, doubles as a feature-selection score in machine learning: a candidate feature X is useful for predicting a target Y exactly to the degree that $I(X; Y)$ is large. Features with near-zero mutual information with the target are, by this measure, theoretically uninformative, regardless of how intuitive they seem to a domain expert.

1.5.3 The Channel as a Neuron Model

A single artificial neuron (formalised in Chapter 5) computes a weighted sum of its inputs, adds noise (from dropout, from finite-precision arithmetic, from stochastic gradient estimates during training), and passes the result through a non-linearity. This is structurally identical to Shannon's channel model: input, corrupting noise, output. Reading a neuron this way is not just a metaphor; it is literally how researchers derive information-theoretic bounds on what a given network architecture can, and cannot, represent.

1.5.4 The Shannon Limit as a Learning Limit

Section 1.3.6 argued that sensor noise imposes a hard ceiling on what any model can extract from plant telemetry. The same argument applies one level up: a training dataset itself has a finite mutual information with the underlying phenomenon it was sampled from. No architecture, however large, and no amount of training compute, can extract more information about the phenomenon than the dataset actually contains. This is the single most important **Honest Boundary** of this entire book, and it will be repeated, in increasingly concrete forms, through Chapters 10 to 14.

1.5.5 The Information Bottleneck

The **Information Bottleneck** principle (Tishby, Pereira, and Bialek) formalises what a good internal representation T of an input X , aimed at predicting a target Y , should do:

Information Bottleneck Objective

$$\min_{P(t|x)} I(X; T) - \beta I(T; Y)$$

The first term pushes the representation T to discard as much irrelevant detail about X as possible (compression); the second term pulls it to retain as much relevance to Y as possible (prediction); β trades one against the other. This is precisely the objective an autoencoder or a variational autoencoder (Chapter 8) is implicitly approximating, and it gives a rigorous, information-theoretic answer to the otherwise vague question of what a "good" learned representation even means.

1.6 Worked Example: Shannon in an Industrial Sensor

Information Budget of a Temperature Transmitter

Consider a temperature sensor with range 0–1000 °C, stated accuracy 0.1 °C, sampled at 1 kHz.

Step 1 — Distinguishable levels. The accuracy figure means there are, at most,

$$N = \frac{1000}{0.1} = 10,000 \text{ distinguishable levels.}$$

Step 2 — Worst-case (maximum-entropy) bits per sample. If every level were equally likely — the maximum-entropy assumption, and hence an *upper bound* on the true entropy — each sample would carry

$$H_{\max} = \log_2(10,000) \approx 13.29 \text{ bits.}$$

Step 3 — Raw information rate. At 1000 samples per second,

$$R_{\text{raw}} \approx 13.29 \times 1000 \approx 13,290 \text{ bits/s} \approx 1.66 \text{ kB/s.}$$

Over 24 hours, that is roughly 1.15×10^9 bits, or about **143 MB per day**, per sensor, at full nominal precision.

Step 4 — Compare against the physical channel's capacity. Suppose the transmitter's physical link has an effective bandwidth of $B = 500$ Hz (a reasonable Nyquist figure for a 1 kHz sample rate) and an operating SNR = 40 dB ($= 10^4$ linear). By Shannon-Hartley,

$$C = 500 \times \log_2(1 + 10^4) \approx 500 \times 13.29 \approx 6,640 \text{ bits/s.}$$

This is *below* the 13,290 bits/s that full nominal precision would require at 1 kHz. In words: at these bandwidth and SNR figures, the physical channel cannot losslessly carry the full stated resolution of the sensor at the full sample rate. Something must give — lower sample rate, coarser resolution, better SNR, or wider bandwidth — and rate-distortion theory (Section 1.4) is precisely the tool for choosing which, and by how much, with a quantified cost.

Step 5 — How much of this is actually useful? Step 2's figure is a *worst case*: it assumes consecutive readings are independent and uniformly distributed across all 10,000 levels. A real process temperature changes slowly relative to a 1 kHz sample rate; consecutive samples are highly correlated, so the *true* entropy rate of the signal (the information genuinely needed to predict the next sample given the recent past) is far lower than $H_{\max}$ — plausibly in the low single-digit bits per sample for a slowly drifting process, though the exact figure is an empirical property of the specific process and must be measured, not assumed. This gap between raw quantization bits and true entropy rate is exactly the opportunity that feature extraction and compression (Section 1.4.3) exploit, and it is the quantitative reason that a failure-prediction model does not need anywhere near 143 MB/day of raw telemetry to do its job well: most of those bits are redundant, not informative.

Honest Boundary

The Shannon theorem is **asymptotic**: its guarantees hold in the limit of infinite block length $n \rightarrow \infty$. Real systems operate with finite, often short, blocks, and pay a rate penalty for it.

The theorem proves a good code **exists**; it does not construct one. Coding theory, a separate and still-active field, is the discipline of closing that gap.

Entropy and capacity presuppose a **known probability distribution**. In practice, that distribution is estimated from data, and the estimate itself carries uncertainty that these formulas do not account for.

The channel model — a fixed conditional distribution $P(y | x)$ — is a simplification. Real industrial links have time-varying noise, drift, and failure modes that a static AWGN or BSC model does not capture. Treat every capacity figure in this chapter as a ceiling computed under stated assumptions, not a measurement of the actual installed system.

2. Probability and Statistics — A Staged Introduction

2.1 What You Already Know

Every engineer already draws a sharp, instinctive line between two kinds of behaviour. A brake must be deterministic: press the pedal, and the same force must produce the same deceleration, every time, under specification. The rain tomorrow, by contrast, is stochastic: you can forecast it, bound it, quote a probability for it, but you cannot pin down its exact amount in advance, and no engineering effort will change that. The difference is not a difference in how well we understand each system. It is a difference in *kind*. A brake's behaviour is, in principle, fully determined by its design once manufacturing tolerances are accounted for; weather is the visible output of a chaotic system whose microscopic state we can never fully observe.

This distinction is the spine of the entire book. Chapter 1 measured uncertainty in bits. This chapter builds the mathematical machinery that *describes* uncertainty, so that later chapters can decide, case by case, which parts of an AI system may be allowed to behave like the rain, and which parts must be engineered to behave like the brake. The chapter is organised as nine escalating **levels**, each one a self-contained tool, each one a prerequisite for the next.

2.2 Level 1 — Intuition: Sample Spaces, Events, and Kolmogorov's Axioms

A **sample space** Ω is the set of everything that could happen in a given experiment: every possible sequence of sensor readings, every possible failure pattern of a redundant system, every possible outcome of a die roll. An **event** A is a subset of Ω — a well-defined collection of outcomes, such as “the pressure sensor reads above 120 bar” or “at least two of the three pumps have failed.”

Andrey Kolmogorov, in 1933, reduced the entire mathematics of probability to three axioms that any assignment of probabilities $P(A)$ to events must satisfy:

Kolmogorov's Axioms

1. **Non-negativity:** $P(A) \geq 0$ for every event A .
2. **Normalisation:** $P(\Omega) = 1$.
3. **Countable additivity:** for any sequence of *mutually exclusive* events $A_1, A_2, \dots$,

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i).$$

Everything else in probability theory — every formula in this chapter, and every loss function in the rest of the book — is a theorem derived from these three rules. It is worth sitting with that fact: the entire statistical apparatus behind a modern neural network rests on three axioms an engineer can state from memory. The complexity that follows is complexity of *consequence*, not of foundation.

2.3 Level 2 — Conditional Probability and Bayes' Theorem

Conditional probability $P(A | B)$ is the probability of event A , given that event B is known to have occurred:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}, \quad P(B) > 0.$$

Rearranging this definition two ways and equating the results gives one of the most consequential formulas in applied science.

Bayes' Theorem

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)}$$

Fault Diagnosis Given a Sensor Reading

Suppose a bearing fails (event F) in 1 in 500 operating hours on average, i.e. $P(F) = 0.002$ for a given hour of operation. A vibration sensor flags an anomaly (event A) in 95% of the hours in which the bearing is actually failing: $P(A | F) = 0.95$. But the sensor is imperfect and also flags an anomaly in 3% of healthy hours due to noise: $P(A | \neg F) = 0.03$.

Given an alarm, what is the probability the bearing is actually failing? By the law of total probability,

$$P(A) = P(A | F)P(F) + P(A | \neg F)P(\neg F) = (0.95)(0.002) + (0.03)(0.998) \approx 0.0318.$$

Then, by Bayes' theorem,

$$P(F | A) = \frac{(0.95)(0.002)}{0.0318} \approx 0.0597.$$

Despite a sensor that looks quite good on paper (95% detection, only 3% false-alarm rate), a raised alarm corresponds to only about a **6% chance of a genuine fault**. This is not a flaw in the sensor; it is a direct, unavoidable consequence of the base rate $P(F)$ being small. This exact calculation — and this exact trap for the unwary — reappears throughout the book every time an anomaly-detection

or fault-diagnosis system is evaluated (Chapter 16), and it is the single most common source of disappointment when a “95% accurate” model is deployed against a rare event.

2.4 Level 3 — Random Variables and Distributions

A **random variable** X maps outcomes in Ω to numbers. **Discrete** random variables (a count of failed sensors, a number of alarms per shift) are described by a probability mass function $P(X = x)$. **Continuous** random variables (a temperature, a pressure, a time-to-failure) are described by a probability density function $f(x)$, with $P(a \leq X \leq b) = \int_a^b f(x) dx$. Three named distributions do most of the work in industrial applications, and each has a distinct physical story behind it.

Three Canonical Distributions

Exponential (time between failures of a component with a constant hazard rate λ):

$$f(t) = \lambda e^{-\lambda t}, \quad t \geq 0, \quad \mathbb{E}[T] = \frac{1}{\lambda}$$

Gaussian / Normal (measurement noise, the sum of many small independent effects):

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

Poisson (count of rare, independent events in a fixed interval, e.g. alarms per shift):

$$P(X = k) = \frac{\mu^k e^{-\mu}}{k!}, \quad k = 0, 1, 2, \dots$$

The Exponential distribution has a defining property worth internalising: it is **memoryless** — $P(T > s + t \mid T > s) = P(T > t)$. A component whose failure time is truly Exponential is, at every instant it is still working, exactly as likely to fail in the next hour as it was on the day it was installed. This is a strong, often unrealistic, modelling assumption (most real components wear out, meaning their hazard rate increases with age), and it is exactly the kind of assumption this book insists on naming rather than smuggling in silently.




Figure 2.1: The three distributions that do most of the modelling work in this book: Exponential for time-to-failure, Gaussian for measurement noise, Poisson for rare event counts.

2.5 Level 4 — Multivariate Distributions

Real plants are never described by one signal in isolation. Temperature, pressure, and flow move together, and the relationships between them often carry more diagnostic information than any single signal alone. The **joint distribution** $P(X, Y)$ of two random variables describes their behaviour together; the **marginal** distribution $P(X) = \sum_y P(X, Y = y)$ recovers the distribution of one variable, averaging out the other; and **covariance** and **correlation** quantify how strongly, and in what direction, two variables move together:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mu_X)(Y - \mu_Y)], \quad \rho_{X,Y} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y} \in [-1, 1].$$

The natural generalisation of the Gaussian to d dimensions is the **multivariate Gaussian**, fully specified by a mean vector μ and a covariance matrix Σ :

Multivariate Gaussian

$$f(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^\top \Sigma^{-1}(\mathbf{x} - \mu)\right)$$

Geometrically, the contours of constant density are ellipsoids centred at μ , whose axes point along the eigenvectors of Σ and whose lengths scale with the square roots of the corresponding eigenvalues. This geometric picture is not a mere curiosity: it is exactly what Principal Component Analysis computes (it diagonalises Σ), and it is exactly the shape that a Mahalanobis-distance anomaly detector (Chapter 16) draws around “normal” operation.

2.6 Level 5 — Statistical Inference

Sections 1 through 4 assumed a known distribution. In practice, the distribution is never known; it must be **estimated** from a finite sample. An **estimator** $\hat{\theta}$ is a function of the data that produces a guess at a true, unknown parameter θ (e.g. the sample mean $\bar{x}$ as an estimator of the true mean μ). A **confidence interval** quantifies the estimator’s uncertainty: a 95% confidence interval is an interval constructed by a procedure that, if repeated over many independent samples, would contain the true parameter in 95% of those repetitions.

This is the quality-control mindset that separates engineering statistics from a single point estimate: a number without an uncertainty band is not yet a measurement. Every reliability figure, every mean-time-between-failures estimate, and every model-accuracy number quoted later in this book should be read as carrying an implicit confidence interval, even where space does not allow it to be written out explicitly.

2.7 Level 6 — Hypothesis Testing

Hypothesis testing formalises the question “is this difference real, or could it plausibly be noise?” A **null hypothesis** H_0 (typically, “no effect,” e.g. “the new sensor calibration has not changed the mean reading”) is tested against data; the **p-value** is the probability, *assuming H_0 is true*, of observing data at least as extreme as what was actually observed. **Power** is the probability of correctly rejecting H_0 when it is in fact false — the test’s ability to detect a real effect when one exists.

A long-running methodological dispute, directly relevant to how Chapter 13 approaches model evaluation, is the **Bayesian versus Frequentist** divide. The Frequentist view treats the true parameter as a fixed, unknown constant and probability as a long-run frequency over repeated experiments; confidence intervals and p-values belong to this world. The Bayesian view treats the parameter itself as a random variable with a **prior** distribution, updates that prior with observed data via Bayes' theorem (Section 2.3) to obtain a **posterior**, and reports probabilities as direct statements of belief ("there is a 90% probability the true failure rate lies in this interval"). Neither camp is simply wrong; they answer different questions, and this book will use both, chosen by which question is actually being asked.

2.8 Level 7 — Stochastic Processes

A **stochastic process** is a sequence of random variables indexed by time, $\{X_t\}$. Three specific processes recur throughout the applications half of this book (Part IV).

Markov Chains, Poisson Processes, Queues

A **Markov chain** has the property that the future depends on the past only through the present state:

$$P(X_{t+1} | X_t, X_{t-1}, \dots, X_0) = P(X_{t+1} | X_t).$$

A **Poisson process** counts the occurrences of independent, rare events over continuous time at a constant rate λ ; the number of events in any interval of length t is Poisson-distributed with mean λt , and the *gaps* between consecutive events are, as a direct consequence, Exponentially distributed — the process that connects the two distributions of Section 2.4. A **queue** (e.g. the M/M/1 model: Poisson arrivals, Exponential service times, one server) describes items or alarms waiting for processing, and gives closed-form results for expected wait time and queue length as a function of the arrival rate and service rate.




Figure 2.2: A three-state Markov chain modelling equipment condition per inspection cycle: Healthy, Degraded, Failed. Only state-*changing* transitions are drawn; the remaining probability mass at each state (0.97, 0.37, and 1.00 respectively) is the probability of staying put for another cycle. Transition probabilities depend only on the current state — the Markov property — not on how long the equipment has already been in it. "Failed" is an absorbing state: once reached, the chain stays there with probability 1.

2.9 Level 8 — Information Theory Meets Statistics

Chapter 1 defined entropy for a single distribution. Comparing *two* distributions — the true data-generating process against a model's belief about it — requires one

more tool.

Kullback-Leibler Divergence

$$D_{\text{KL}}(P \parallel Q) = \sum_i P(x_i) \log_2 \frac{P(x_i)}{Q(x_i)}$$

$D_{\text{KL}}(P \parallel Q)$ measures the extra bits wasted by encoding data that actually comes from P using a code optimised for Q instead. It is always non-negative, and zero exactly when $P = Q$. It is *not* symmetric ($D_{\text{KL}}(P \parallel Q) \neq D_{\text{KL}}(Q \parallel P)$ in general), which matters: it is not a true distance, only a measure of divergence, and the direction of comparison changes what it penalises most. Mutual information, introduced in Chapter 1 as what a channel preserves, has a compact re-statement in this language:

$$I(X; Y) = D_{\text{KL}}(P(x, y) \parallel P(x)P(y)),$$

the divergence between the true joint distribution and the joint distribution the two variables *would* have if they were independent. This single equation is the bridge connecting Chapter 1's information theory to this chapter's statistics, and it reappears directly in Section 2.11's third worked example.

2.10 Level 9 — Towards Machine Learning

Three estimation principles, all direct consequences of Bayes' theorem, are the last rung of this chapter's ladder and the first rung of Chapter 3.

MLE, MAP, and Bayesian Inference

Maximum Likelihood Estimation chooses the parameter that makes the observed data most probable:

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} P(\text{data} \mid \theta).$$

Maximum A Posteriori estimation adds a prior belief $P(\theta)$ and maximises the posterior instead:

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} P(\text{data} \mid \theta) P(\theta) = \arg \max_{\theta} [\log P(\text{data} \mid \theta) + \log P(\theta)].$$

Bayesian inference does not collapse to a single point estimate at all; it keeps and reports the entire posterior distribution $P(\theta \mid \text{data})$.

The MAP formula deserves a second look, because it is quietly the origin story of a technique every reader who has trained a model has already used. If the prior $P(\theta)$ is chosen to be Gaussian centred at zero, $\log P(\theta)$ is (up to constants) $-\theta^2$: MAP estimation under a Gaussian prior is *exactly* maximum likelihood plus an L2 penalty on the parameters — what every deep learning practitioner calls **weight decay** or **ridge regularisation**, introduced formally in Chapter 3. A Laplace prior gives an L1 penalty instead. Regularisation, in other words, is not an ad-hoc trick bolted onto neural network training; it is Bayesian inference with a specific, nameable prior belief about what a "reasonable" set of weights looks like. This is the bridge the title of this section promises, and it is the first of many places in this book where a technique from deep learning turns out, on inspection, to be classical statistics wearing new notation.

2.11 Worked Examples

Probability of Failure in a Redundant System

A safety function uses **2-out-of-3 voting**: it requires at least two of three identical, independent sensors to agree before it trips. Each sensor has an independent failure probability $q = 0.02$ per demand. The voting system itself fails (fails to trip when it should) only if *at least two* of the three sensors fail simultaneously. Let X be the number of failed sensors, $X \sim \text{Binomial}(3, q)$.

$$P(X \geq 2) = \binom{3}{2} q^2 (1 - q) + \binom{3}{3} q^3 = 3(0.02)^2 (0.98) + (0.02)^3$$

$$= 3(0.0004)(0.98) + 0.000008 = 0.001176 + 0.000008 = 0.001184.$$

The voted system's failure probability is about **0.118%**, roughly 17 times lower than a single sensor's own 2% failure probability. This is the quantitative payoff of redundancy — and the Honest Boundary closing this section will immediately qualify it: the entire calculation rests on the sensors failing *independently*, an assumption that is frequently false in practice (shared power supply, shared calibration error, common-cause failure from a single flood or fire).

Estimating Mean Lifetime from Censored Data

Ten identical units are placed on a reliability test that runs for a fixed duration $T = 1000$ hours (Type I, or "time," censoring). By the end of the test, $d = 6$ units have failed, at recorded times summing to $\sum t_i = 3,200$ hours; the remaining 4 units are still running at $T = 1000$ hours each and are recorded as censored. Assuming Exponential lifetimes (Section 2.4), the maximum likelihood estimator of the failure rate λ under Type I censoring uses the **total time on test**: the sum of the observed failure times *plus* the censored units' full running time, since every unit, failed or not, contributed information for as long as it was observed.

$$\text{Total time on test} = \sum t_i + (n - d) T = 3,200 + (4)(1000) = 7,200 \text{ hours.}$$

$$\hat{\lambda}_{\text{MLE}} = \frac{d}{\text{total time on test}} = \frac{6}{7,200} \approx 8.33 \times 10^{-4} \text{ hr}^{-1}, \quad \widehat{\text{MTTF}} = \frac{1}{\hat{\lambda}} \approx 1,200 \text{ hours.}$$

Notice what this estimator gets right, and what an untrained instinct gets wrong: simply averaging the six *observed* failure times ($3,200/6 \approx 533$ hours) would badly *underestimate* the true mean lifetime, because it silently discards the four units that survived the whole test — exactly the most reliable units in the batch. Censoring must be accounted for in the likelihood itself, not patched afterward.

Detecting Distribution Shift with KL Divergence

A sensor's readings are binned into five ranges. Over a long base-line period, the historical (reference) distribution across bins is $P = (0.10, 0.25, 0.35, 0.22, 0.08)$. Over the most recent monitoring window, the observed distribution is $Q = (0.05, 0.15, 0.30, 0.30, 0.20)$ — visibly shifted toward the higher bins.

$$D_{\text{KL}}(Q \parallel P) = \sum_i Q_i \log_2 \frac{Q_i}{P_i}$$

$$\begin{aligned}
&= 0.05 \log_2 \frac{0.05}{0.10} + 0.15 \log_2 \frac{0.15}{0.25} + 0.30 \log_2 \frac{0.30}{0.35} + 0.30 \log_2 \frac{0.30}{0.22} + 0.20 \log_2 \frac{0.20}{0.08} \\
&\approx (-0.050) + (-0.111) + (-0.067) + (0.134) + (0.264) \approx 0.171 \text{ bits.}
\end{aligned}$$

A single divergence number, tracked over successive monitoring windows, is a standard, principled **concept-drift detector**: as long as D_{KL} stays near the noise floor established during stable operation, the process is behaving as trained; a sustained rise above a calibrated threshold flags that the distribution the model was trained on no longer matches the distribution it is now seeing — a warning sign that its predictions can no longer be trusted, addressed directly in Chapter 13.

Honest Boundary

The choice of **prior** in Bayesian analysis (Section 2.9) is inescapably **subjective**. Two honest analysts, given the same data, can reach different posteriors if they started from different, individually defensible priors. This is not a defect to be engineered away; it is a property of the framework, and it must be disclosed, not hidden, whenever Bayesian inference informs a safety-relevant conclusion.

P-values are systematically misinterpreted, even by experienced practitioners. A p-value of 0.03 is *not* “a 3% probability the null hypothesis is true”; it is the probability of data this extreme *given that the null hypothesis is true* — a subtly but fundamentally different statement. Confusing the two is one of the most persistent errors in applied statistics.

Independence is rare in practice. Nearly every formula in this chapter — the redundancy calculation in Section 2.11, the additivity of variance, the simple form of a likelihood as a product — assumed independence somewhere. Real industrial components share power supplies, share environments, share maintenance crews, and share calibration drift. Treat every independence assumption in this book, and in any statistical argument you are handed, as a claim to be checked against the physical system, not a convenience to be accepted by default.

3. From Data to Models

3.1 What You Already Know

Every engineer has, at some point, fitted a straight line through a scatter of noisy data points — a calibration curve, a trend line through a set of measurements, a linear fit relating valve position to flow rate. You minimised the sum of squared vertical distances between the line and the points (the **residuals**), because it was the method your spreadsheet or your textbook handed you, and because the resulting line looked reasonable. What this chapter makes explicit is that the **least-squares method** is not an arbitrary convention. It is the single most important special case of a much more general principle — Maximum Likelihood Estimation, from Chapter 2, Level 9 — and understanding it that way is what lets the same underlying logic scale from a two-parameter calibration line to a billion-parameter neural network.

3.2 Maximum Likelihood and Maximum A Posteriori Estimation, Applied

Recall from Section 2.10 that MLE chooses the parameter θ that maximises $P(\text{data} | \theta)$. Fitting a model $y = f(x; \theta) + \varepsilon$ to data is exactly this, once you commit to a noise model for ε .

Least Squares Is Maximum Likelihood

Assume the noise on each observation is independent and Gaussian, $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$. The likelihood of the full dataset $\{(x_i, y_i)\}_{i=1}^n$ is

$$P(\text{data} | \theta) = \prod_{i=1}^n \frac{1}{\sigma \sqrt{2\pi}} \exp\left(-\frac{(y_i - f(x_i; \theta))^2}{2\sigma^2}\right).$$

Taking the log (Chapter 1's move: turning products into sums) and discarding the additive constants that do not depend on θ ,

$$\log P(\text{data} | \theta) = -\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - f(x_i; \theta))^2 + \text{const.}$$

Maximising this log-likelihood over θ is identical to *minimising* $\sum_i (y_i - f(x_i; \theta))^2$ — the ordinary least-squares objective. The method every engineer already knows is, without exception, maximum likelihood estimation under an assumption of Gaussian noise. Change the noise assumption, and the “natural” loss function changes with it: Laplace-distributed noise leads to minimising *absolute* error, not squared error; a Poisson noise model leads to the loss function used

to fit event-count data. There is no such thing as a loss function chosen for free; every loss function is a noise model in disguise, whether or not the person who chose it thought of it that way.

As Section 2.10 already showed, adding a prior over θ turns MLE into MAP, and a Gaussian prior on θ turns plain least squares into **ridge regression**. Section 3.4 returns to this in full.

3.3 The Bias-Variance Tradeoff

Every model's expected error on new, unseen data decomposes, exactly, into three non-negative pieces:

Bias-Variance Decomposition

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{(\text{Bias}[\hat{f}(x)])^2}_{\text{underfitting}} + \underbrace{\text{Var}[\hat{f}(x)]}_{\text{overfitting}} + \underbrace{\sigma^2}_{\text{irreducible noise}}$$

Bias is the error from a model too simple to capture the true relationship — fitting a straight line to a curved process, no matter how much data you collect, will never remove this term. **Variance** is the error from a model so flexible that it fits the noise in this *particular* training set, producing wildly different predictions if the training data were resampled. **Irreducible noise** σ^2 is the floor set by Section 1.3.6 and Section 1.3.1: the entropy the measurement process itself injects, which no model, however well chosen, can predict away.




Figure 3.1: Training error falls monotonically with model complexity; test error falls, reaches a minimum, then rises again as variance takes over. The gap between the two curves past the sweet spot is overfitting, made visible.

The practical consequence, visible in Figure 3.1, is that **training error is not a trustworthy guide to model quality**: it decreases monotonically with model complexity, and a model with zero training error is not a good model, it is a memoriser. Only error measured on data the model did not see during fitting — the subject of Section 3.4 — can reveal the sweet spot.

3.4 Cross-Validation and Regularization

3.4.1 Cross-Validation

k -fold cross-validation estimates test error without sacrificing a separate hold-out set permanently: the training data is split into k equal folds; the model is trained on $k - 1$ folds and evaluated on the remaining one; this is repeated k times, rotating which fold is held out, and the k resulting error estimates are averaged. It is, in effect, a way of simulating “unseen data” k times over using only the data already in hand — valuable precisely because industrial datasets, unlike academic benchmarks, are often too small to afford a large permanent hold-out set.

3.4.2 Regularization

Section 3.2 showed that a Gaussian prior on the parameters turns MLE into MAP, and that MAP turns into a penalised least-squares problem. Made explicit, with a tunable strength λ :

Ridge, Lasso, and Elastic Net

$$\text{Ridge (L2): } \min_{\theta} \sum_i (y_i - f(x_i; \theta))^2 + \lambda \sum_j \theta_j^2$$

$$\text{Lasso (L1): } \min_{\theta} \sum_i (y_i - f(x_i; \theta))^2 + \lambda \sum_j |\theta_j|$$

$$\text{Elastic Net: } \min_{\theta} \sum_i (y_i - f(x_i; \theta))^2 + \lambda_1 \sum_j |\theta_j| + \lambda_2 \sum_j \theta_j^2$$

Ridge shrinks every parameter toward zero smoothly, and is the right choice when you believe most inputs contribute a little. Lasso, owing to the geometry of the L1 penalty (its contours have corners exactly on the coordinate axes), drives some parameters to *exactly* zero, performing automatic feature selection — the right choice when you believe only a handful of inputs actually matter and you want the model to tell you which. Elastic Net interpolates between the two when you are not sure which story applies, which is most of the time in industrial data. The regularization strength λ itself is not chosen by inspection; it is chosen by cross-validation, closing the loop between the two halves of this section.

3.4.3 The Closed-Form Ridge Solution

Unlike Lasso, which has no closed form and must be solved numerically, ridge regression on a linear model $f(x; \theta) = X\theta$ can be solved exactly. Writing the objective in matrix form and setting its gradient to zero,

$$\nabla_{\theta} [\|y - X\theta\|^2 + \lambda \|\theta\|^2] = 0 \implies \hat{\theta}_{\text{ridge}} = (X^{\top}X + \lambda I)^{-1} X^{\top}y.$$

Compare this to the ordinary least-squares solution $\hat{\theta}_{\text{OLS}} = (X^{\top}X)^{-1} X^{\top}y$: ridge simply adds λ to the diagonal of $X^{\top}X$ before inverting. This has a second, entirely practical benefit beyond regularisation: when input features are highly correlated (a near-universal condition in plant telemetry, where temperature, pressure, and flow are physically coupled), $X^{\top}X$ is nearly singular and $\hat{\theta}_{\text{OLS}}$ is numerically unstable — small changes in the data produce wild swings in the fitted coefficients. Adding

λI before inverting fixes this directly, which is why ridge regression is often worth using purely for numerical stability, independent of any belief about the true sparsity of the underlying relationship.

3.4.4 Why L1 Produces Sparsity: The Geometric Picture

The difference between Ridge and Lasso is easiest to see, not to compute. Both methods minimise the same squared-error term subject to a budget on the penalty; geometrically, this means finding the point where the elliptical contours of constant squared error — centred on the unconstrained least-squares solution $\hat{\theta}_{\text{OLS}}$ — first touch the constraint region defined by the penalty.




Figure 3.2: Why L1 regularisation induces sparsity and L2 does not. Both panels show elliptical contours of constant squared-error (orange), centred on the unconstrained solution $\hat{\theta}_{\text{OLS}}$, and the constraint region of the penalty (teal), centred on the origin. The regularised solution is the first point where the contours touch the constraint region. For the L2 ball (**left**), that point is generically interior to an edge, away from either axis: both θ_1 and θ_2 shrink, but neither is driven to zero. For the L1 diamond (**right**), the constraint region has sharp corners sitting exactly on the axes; contours frequently make first contact *at* a corner, where one coordinate is exactly zero. This is a geometric fact about diamonds versus circles, not a numerical accident.

Figure 3.2 makes visible what Section 2.10 stated algebraically: a Laplace prior (Lasso) concentrates probability mass along the axes of parameter space, in a way a Gaussian prior (Ridge) does not, and optimisation under that prior inherits the same preference for corner solutions.

3.5 The Minimum Description Length Principle

Jorma Rissanen's **Minimum Description Length** (MDL) principle reframes model selection as a direct application of Chapter 1's compression theory: the best model is the one that lets you *describe the data using the fewest total bits*, where the total cost has two parts.

MDL Objective

$$L(\text{model}) = \underbrace{L(\theta)}_{\substack{|\{Z\}| \\ \text{bits to describe the model}}} + \underbrace{L(\text{data} \mid \theta)}_{\substack{|\{Z\}| \\ \text{bits to describe the residuals, given the model}}}$$

A model that is too simple (high bias) fits poorly, so $L(\text{data} \mid \theta)$ is large: the residuals are large and themselves hard to compress. A model that is too complex (high variance) has a large $L(\theta)$: it takes many bits just to write down its parameters, or, in a real sense, it has memorised the noise, which is not compressible *as* noise. MDL therefore recovers the bias-variance sweet spot of Section 3.3 from an entirely different starting point — description length rather than test error — and the fact that two independently motivated principles converge on the same answer is itself evidence that something real is being measured.




Figure 3.3: The MDL tradeoff. The cost of describing the model itself, $L(\theta)$, grows with complexity; the cost of describing what the model fails to explain, $L(\text{data} \mid \theta)$, shrinks with it. Their sum is minimised at an intermediate complexity — the same sweet spot Figure 3.1 located using test error instead of bits.

3.6 Learning as Compression: The Shannon Connection Revisited

Section 1.5’s five connections between Shannon’s theory and machine learning can now be drawn together into a single sentence: a model that generalises well is, in every formal sense developed so far, a model that has found a short description of the data’s regularities, rather than a long description of the data’s noise. Cross-entropy loss (Section 1.5.1) is the expected code length under the model’s beliefs; the Information Bottleneck (Section 1.5.5) explicitly trades compression against prediction; and MDL, just introduced, makes the connection to code length completely literal rather than metaphorical. This is not three separate coincidences. It is one idea, stated three times in three different specialised vocabularies, and recognising the underlying unity is what lets a result from information theory transfer directly into an intuition about why a particular neural network architecture generalises, or fails to.

3.7 The Curse of Dimensionality and the No Free Lunch Theorem

3.7.1 The Curse of Dimensionality

Intuitions built in two or three dimensions actively mislead in high dimensions. As the number of input features d grows, the volume of space grows exponentially, and the data needed to cover it densely grows with it; equivalently, in high dimensions, essentially all of the volume of a hypercube concentrates near its corners, and essentially all points become approximately equidistant from one another, which quietly breaks any method (including the nearest-neighbour and clustering methods used throughout Part IV) that relies on the notion of "nearby" points being meaningfully more similar than "far" points.

The Shrinking Inscribed Sphere

Consider a d -dimensional hypercube of side length 2 (volume 2^d) and its inscribed hypersphere of radius 1. The volume of a unit-radius d -ball is $V_d = \pi^{d/2} / \Gamma(d/2 + 1)$. The fraction of the cube's volume occupied by the sphere is:

d	Sphere / cube volume ratio
2	0.785
10	0.00249
50	1.54×10^{-28}

At $d = 2$, the inscribed circle fills 78.5% of the square: geometric intuition holds. At $d = 50$, a modest number of features by modern standards, the inscribed sphere occupies a fraction of the cube so small it is not meaningfully different from zero: essentially all of the volume, and hence essentially all randomly sampled data, sits in the corners, far from the center. Any algorithm whose logic implicitly assumes data fills the space uniformly and densely quietly stops working somewhere in this range, usually without announcing that it has done so.




Figure 3.4: The sphere/cube volume ratio computed in the worked example above, shown across the full range $d = 1$ to 60 on a logarithmic axis. Dashed lines mark $d = 2$, $d = 10$, and $d = 50$. Linear intuition (built at $d = 2$ or $d = 3$) suggests the inscribed sphere should always occupy “most” of the cube; the true curve falls by roughly one order of magnitude every two to three added dimensions, crossing 23 orders of magnitude by $d = 60$.

3.7.2 The No Free Lunch Theorem

Wolpert and Macready’s **No Free Lunch Theorem** for optimisation and learning proves a result that is easy to state and uncomfortable to accept: averaged over *all possible* problems, every learning algorithm performs identically, including random guessing. An algorithm that outperforms another on one class of problems must underperform it on some other class, as a mathematical necessity, not an accident of current technology.

Honest Boundary

The No Free Lunch theorem is frequently misquoted as “all algorithms are equally good in practice,” which is not what it says. It says that averaged over *all mathematically possible* problems, performance is equal. Real industrial problems are drawn from a tiny, highly structured corner of that space — physical processes obey conservation laws, smoothness, and locality — and algorithms that exploit that structure (the inductive biases of Chapter 5 onward) legitimately, consistently, and provably outperform generic ones on the problems that actually occur in engineering. The theorem is a statement about the absence of a universal free lunch, not a statement that no lunch is ever better than another.

No model is universally best; the No Free Lunch theorem makes this a mathematical fact, not a matter of current technology or insufficient research effort. All learning involves **inductive bias** — an assumption, built into the choice of model family, about what kinds of patterns are more likely a priori. Linear regression assumes linearity; a convolutional network (Chapter 8) assumes local, translation-invariant structure. There is no bias-free learning algorithm, and claiming one has been found is itself a signal that the inductive bias has simply gone unexamined, not that it is absent.

4. Towards Machine Learning

4.1 What You Already Know

By this point in the book you have, without perhaps naming it this way, already met every mathematical ingredient a neural network is built from. You know linear regression: fit a line, or a hyperplane, through data by minimising squared error. You very likely know logistic regression too, even if only as “the thing you use instead of linear regression when the output is yes/no” — a classifier that outputs a probability rather than a raw number. And Chapter 3 handed you gradient descent implicitly, the moment it handed you MLE: somewhere, a computer had to actually search for the θ that maximises a likelihood or minimises a loss, and for anything beyond the simplest linear-Gaussian case, that search is gradient descent.

This chapter’s entire purpose is to take these three familiar tools and show that they are not three separate techniques that happen to sit near each other in a statistics course. They are the same mathematical object, viewed at increasing resolution: linear regression *is* the simplest possible artificial neuron; logistic regression *is* that same neuron with one extra ingredient (a non-linear activation function); and gradient descent *is* the one algorithm that trains both, and will turn out, completely unchanged in its core logic, to be the algorithm that trains a network with a billion parameters in Chapter 11. Nothing new needs to be invented to get from a calibration line to a neural network. What is needed is composition, which is exactly the subject of Chapter 5 onward, starting from the single building block this chapter now takes apart in full detail.

4.2 Linear Regression as a Single Neuron

Strip linear regression down to its computational skeleton. Given inputs $x_1, \dots, x_d$, the model computes

The Linear Neuron

$$\hat{y} = w_1 x_1 + w_2 x_2 + \dots + w_d x_d + b = \mathbf{w}^T \mathbf{x} + b$$

and that is, in its entirety, both “multiple linear regression” and the computation performed by a single artificial neuron with a **linear (identity) activation function**. The weights $w_1, \dots, w_d$ say how strongly, and in which direction, each input pushes the prediction; the bias b is the prediction when every input is zero — geometrically, it is what lets the fitted hyperplane sit anywhere in space rather than being pinned through the origin.




Figure 4.1: Linear regression, redrawn as a single artificial neuron. Each input is multiplied by a weight, the results (plus a bias, treated here as the weight on a constant input of 1) are summed to give $z = \mathbf{w}^T \mathbf{x} + b$, and an **identity** activation passes z through unchanged: $\hat{y} = z$. Chapter 5 keeps this exact diagram and simply replaces the activation box.

Fitting a Cooling-Flow Compensation Line

A pump's discharge temperature rise y (in °C above a reference) is measured against the deviation x of its cooling-water flow rate from nominal (in %), across five operating points:

x (% flow deviation)	-10	-5	0	5	10
y (°C rise)	9.1	6.4	4.2	1.9	-0.6

Because $\bar{x} = 0$ here, the least-squares formulas of Chapter 3 simplify: $b = \bar{y} = 4.2$, and

$$w = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2} = \frac{-4.78 \times 100}{5 \times 50} = -0.478.$$

The fitted neuron is $\hat{y} = -0.478x + 4.2$: every 1% increase in cooling flow above nominal reduces predicted discharge-temperature rise by about 0.48 °C. The residuals (0.12, -0.19, 0.00, 0.09, -0.02) sum-of-squares to $\text{SSE} \approx 0.059$, small relative to the spread of y , indicating a tight linear fit — exactly the kind of physically motivated, nearly-linear relationship (heat rejection scales close to linearly with flow over a modest operating range) that makes a single linear neuron, with no hidden layers at all, a perfectly adequate model.

4.3 Logistic Regression as a Classifier

Linear regression predicts an unbounded real number. Many industrial questions, though, are inherently binary: will this bearing fail in the next month, is this reading anomalous, should this alarm be trusted? A raw linear output $\mathbf{w}^T \mathbf{x} + b$ can be any real number, but a probability must live in $[0, 1]$. **Logistic regression** bridges the two with one added ingredient: a squashing function applied to the linear neuron's output.

The Sigmoid / Logistic Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}, \quad \hat{p} = \sigma(\mathbf{w}^T \mathbf{x} + b)$$




Figure 4.2: The sigmoid function. It maps any real-valued z onto a valid probability in $(0, 1)$, is symmetric about $z = 0 \Rightarrow p = 0.5$, and saturates: far from zero, it becomes flat, so large swings in z far from the decision boundary barely change the predicted probability.

4.3.1 The Log-Odds Interpretation

The sigmoid has an exact algebraic inverse that gives logistic regression its second, equally important reading. Solving $p = \sigma(z)$ for z :

$$z = \log\left(\frac{p}{1-p}\right) \quad (\text{the } \mathbf{logit}, \text{ or log-odds}).$$

This means logistic regression is not merely “linear regression squashed into $[0, 1]$ ”; it is, exactly, **linear regression on the log-odds of the outcome**:

$$\log\left(\frac{p}{1-p}\right) = \mathbf{w}^T \mathbf{x} + b.$$

Every coefficient w_j has a precise, defensible engineering meaning under this reading: a one-unit increase in x_j changes the *log-odds* of the outcome by exactly w_j , and multiplies the *odds* by e^{w_j} . This is what lets a trained logistic regression model be reported to a safety reviewer not as an opaque function but as a table of odds ratios — a form of interpretability that survives all the way from a two-variable calibration to a component of the ILM’s architecture in Chapter 10.

Log-Odds of Bearing Failure vs. Vibration Level

A logistic regression model, already fit to historical failure records, relates standardised vibration level x (z-score units) to one-month failure probability via $z = 0.8x - 1.2$:

x (vibration z-score)	$z = 0.8x - 1.2$	$\hat{p} = \sigma(z)$
-1	-2.0	0.119
0	-1.2	0.231
1	-0.4	0.401
1.5	0.0	0.500
2	0.4	0.599
3	1.2	0.769

The decision boundary $\hat{p} = 0.5$ occurs exactly where $z = 0$, at $x = -b/w = 1.2/0.8 = 1.5$ standard deviations above the mean vibration level. Compare the

per-unit change in $\hat{p}$ far from the boundary against near it: from $x = -1$ to $x = 0$, $\hat{p}$ rises by only 0.112; from $x = 1$ to $x = 2$, straddling the boundary, it rises by 0.198 — nearly double. This is the sigmoid's saturation behaviour from Figure 4.2 made concrete: the model is most sensitive to a change in its input precisely where it is least certain, and gracefully insensitive to input changes once it is already confident.

4.3.2 The Cross-Entropy Loss, Concretely

Section 1.5.1 introduced cross-entropy loss in the abstract; here it takes its most common concrete form, for a single binary-labelled example $(\mathbf{x}, y)$ with $y \in \{0, 1\}$:

$$\mathcal{L}(y, \hat{p}) = -[y \log \hat{p} + (1 - y) \log(1 - \hat{p})].$$

This is exactly the negative log-likelihood of a Bernoulli observation, so minimising it over a dataset is, once again, maximum likelihood estimation (Section 3.2) — the same underlying principle as ordinary least squares, just with a Bernoulli noise model standing in for the Gaussian one.

4.4 Gradient Descent: From Calculus to Optimization

Chapter 3 defined MLE and MAP as the θ that maximises or minimises some objective. It did not say *how* to find that θ when, unlike ordinary least squares, no closed-form solution exists — which is the case for logistic regression, and for every neural network in this book. **Gradient descent** is the answer, and its logic is nothing more than first-semester calculus applied iteratively.

Gradient Descent Update Rule

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

where η (the **learning rate**) controls the step size, and $\nabla_{\theta} \mathcal{L}$ is the vector of partial derivatives of the loss with respect to every parameter.

The gradient $\nabla_{\theta} \mathcal{L}$ points, at any given θ , in the direction of *steepest increase* of the loss. Subtracting a small multiple of it therefore takes a step in the direction of steepest *decrease* — downhill. Repeated enough times, with a well-chosen η , this walk converges to a point where the gradient is zero: a minimum (or, as Chapter 7 will qualify carefully, a saddle point or local minimum, in the non-convex landscapes neural networks actually present).

Gradient Descent, Step by Step, on an Elliptical Bowl

Consider the toy loss $\mathcal{L}(w_1, w_2) = (w_1 - 3)^2 + 4(w_2 + 1)^2$, whose minimum is obviously at $(w_1, w_2) = (3, -1)$ by inspection, but whose *path* to that minimum under gradient descent is instructive. The gradient is

$$\nabla \mathcal{L} = (2(w_1 - 3), 8(w_2 + 1)).$$

Starting at $(w_1, w_2) = (0, 0)$ with a fixed learning rate $\eta = 0.1$:

Iteration	w_1	w_2
0	0.000	0.000
1	0.600	-0.800
2	1.080	-0.960
3	1.464	-0.992
4	1.771	-0.998
5	2.017	-1.000
6	2.214	-1.000

Notice that w_2 is essentially converged after two iterations, while w_1 is still visibly short of 3 after six. This is not a coincidence of this particular loss; it is a direct, general consequence of **curvature**. The loss curves eight times more steeply in the w_2 direction than the w_1 direction (coefficients 8 vs. 2), so the *same* learning rate η that is well tuned for w_2 is far too small for efficient progress along w_1 . This single, simple example is the entire motivation for the adaptive-learning-rate optimisers (AdaGrad, RMSProp, Adam) introduced formally in Section 7.6: they exist specifically to give each parameter its own effective step size, precisely because real loss landscapes are never as conveniently symmetric as a textbook bowl.




Figure 4.3: The gradient descent path from the worked example, plotted against the loss contours. The path takes large strides toward the minimum in the steeply-curved w_2 direction and small, slow strides in the shallow w_1 direction, visibly bending rather than heading straight for the minimum — the geometric signature of an ill-conditioned loss landscape.

4.4.1 Batch, Stochastic, and Mini-Batch Gradient Descent

The worked example computed the exact gradient using the full loss at every step — **batch gradient descent**. When the loss is a sum over a large training set, computing the exact gradient requires touching every example on every single step, which is prohibitively slow for the dataset sizes Chapter 10 and Chapter 11 work with. **Stochastic gradient descent** (SGD) estimates the gradient from a single randomly chosen example (or, in practice, **mini-batch gradient descent**, from a

small random subset) at each step. The resulting gradient estimate is noisy — an unbiased but high-variance estimate of the true gradient, in the exact statistical sense of Chapter 2, Section 2.6 — but is vastly cheaper per step, and the noise itself is often beneficial, helping the optimiser escape shallow local minima that would trap an exact, deterministic descent.

4.5 The Problem of Non-linearity: Why a Single Layer Is Not Enough

Every model built so far in this chapter, linear or logistic, computes a single linear combination of its inputs before any squashing. This has a hard geometric consequence: the decision boundary of a logistic regression classifier (the set of points where $\hat{p} = 0.5$, i.e. where $\mathbf{w}^T \mathbf{x} + b = 0$) is always exactly a hyperplane — a straight line in two dimensions, a flat plane in three. No choice of weights can bend it.

A Task No Linear Model Can Solve

Consider a safety interlock fed by two redundant, independent sensor flags $x_1, x_2 \in \{0, 1\}$. The interlock must trigger an alarm precisely when the sensors *disagree* (one reads a fault, the other does not) — a classic exclusive-or (XOR) requirement, used deliberately in redundant instrumentation to flag sensor *discrepancy* rather than a specific fault state:

x_1	x_2	Alarm
0	0	0
0	1	1
1	0	1
1	1	0




Figure 4.4: The XOR problem. Open circles = no alarm, filled circles = alarm. Two representative attempted linear boundaries (dashed) are shown; every possible straight line necessarily puts at least one point on the wrong side, because the two classes are not linearly separable — no matter how the weights of a single logistic neuron are chosen, this exact pattern cannot be learned by it.

No setting of w_1 , w_2 , and b makes $\sigma(w_1 x_1 + w_2 x_2 + b)$ reproduce this table,

and Figure 4.4 shows why geometrically: the two “alarm” points and the two “no alarm” points cannot be separated by any single straight line. This is not a failure of optimisation — no amount of additional training data or cleverer gradient descent variant fixes it, because the function class itself (a single linear boundary) does not contain the needed solution. It is a failure of **representational capacity**, and it is historically significant: this exact limitation, identified rigorously by Minsky and Papert in 1969, is largely what caused the first “AI winter,” a decade-long collapse in neural network research funding, because at the time no one had yet shown how to train networks with more than one layer. Chapter 6 is the direct resolution of this cliffhanger.

4.6 Introduction to Neural Networks: The Perceptron as a Building Block

Frank Rosenblatt’s 1958 **perceptron** is, computationally, almost exactly Figure 4.1 with one change: the identity activation is replaced with a hard step function,

$$\hat{y} = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} + b \geq 0 \\ 0 & \text{otherwise,} \end{cases}$$

trained not by gradient descent on a smooth loss (the step function’s derivative is zero almost everywhere, which breaks Section 3.2’s machinery outright) but by its own dedicated update rule, examined in full in Section 5.5. The perceptron, logistic regression, and linear regression are, in this light, three points on a single continuum, differing only in which activation function sits in the box in Figure 4.1:

Model	Activation	Typical loss
Linear regression	identity	squared error
Logistic regression	sigmoid	cross-entropy
Perceptron	hard step	perceptron criterion


Every one of them, alone, inherits the XOR limitation of Section 4.5, because that limitation is a property of the single linear combination $\mathbf{w}^T \mathbf{x} + b$ shared by all three, not of any particular activation choice. What differs between them is only what happens *after* that linear combination is formed. The resolution, developed in full starting in Chapter 5, is to stop using a single one of these units in isolation and instead wire many of them together in layers — at which point, as Chapter 6 shows concretely, the XOR problem stops being a problem at all.

Honest Boundary

Machine learning, at every level examined in this chapter, is not magic; it is **statistics at scale**. Linear regression, logistic regression, and the perceptron are all instances of the general estimation framework built in Chapter 2 and Chapter 3: choose a model family, choose a loss function that corresponds to a noise assumption, and find the parameters that minimise it. A neural network with a billion parameters, introduced starting in Chapter 5, changes the size and the composition of the model family and the optimiser’s engineering demands enormously. It does not change the underlying logic of estimation in any way that has been, or needs to be, exempted from anything already established in this book. Treat any claim that a modern AI system operates by fundamentally different principles than the ones in this chapter with the same scepticism you

would apply to any other extraordinary claim: the burden of proof is on the claim, and so far, it has not been met.

NEURAL NETWORKS - FROM PERCEPTRON TO DEEP NETWORKS



5. The Perceptron — The Foundation

5.1 What You Already Know

Three pieces of mathematics you already carry with you turn out to be the entire geometric content of this chapter. You know the **inner product** $\mathbf{w}^\top \mathbf{x} = \sum_j w_j x_j$, and you know that it measures how much one vector points in the direction of another. You know that the set of points satisfying $\mathbf{w}^\top \mathbf{x} + b = 0$ is a **hyperplane** — a line in two dimensions, a plane in three, a flat $(d-1)$ -dimensional sheet cutting through d -dimensional space, always perpendicular to $\mathbf{w}$. And you know, intuitively, what it means for two clouds of points to be **linearly separable**: that some straight cut divides one cloud entirely from the other, with nothing on the wrong side. Chapter 4 built logistic regression and linear regression on exactly this geometry. This chapter asks a narrower, older, and in some ways more fundamental question: what is the *simplest* possible learning machine that finds such a separating hyperplane by itself, from data, with no calculus and no gradient at all?

5.2 The Biological Neuron

Artificial neurons take their name, and their original inspiration, from an actual cell. A biological neuron collects electrical and chemical signals through branching **dendrites**; these signals are integrated in the cell body, or **soma**; if the accumulated signal crosses a threshold, the neuron fires an **action potential** — an all-or-nothing electrical pulse — down a long fibre, the **axon**, to its endpoint; and at a **synapse**, that pulse is passed, with some strength that can itself change over time (synaptic plasticity, the biological basis of learning), to the dendrites of the next neuron in the chain.




Figure 5.1: The biological neuron, schematically. Dendrites collect incoming signals; the soma integrates them; if threshold is crossed, an action potential propagates down the axon (the hatch marks represent the myelin sheath that speeds this propagation); the synapse passes the signal onward, with adjustable strength. Every element of this picture has a direct, named counterpart in the artificial neuron of Section 5.4.

5.3 The McCulloch-Pitts Model

In 1943, neurophysiologist Warren McCulloch and logician Walter Pitts proposed the first mathematical model of a neuron, years before any computer existed to run it. Their neuron summed binary inputs, each either excitatory or inhibitory, and fired (output 1) if the sum met or exceeded a threshold, and stayed silent (output 0) otherwise. Despite its extreme simplicity, McCulloch and Pitts proved that networks of such units could, in principle, compute any logical proposition expressible in propositional logic — the first formal demonstration that a network of simple, threshold-based units could carry out arbitrarily complex computation. Every neuron introduced in this book, all the way through the Transformer of Chapter 8, is a direct mathematical descendant of this one 1943 idea.

5.4 The Artificial Neuron (Perceptron)

Frank Rosenblatt's 1958 **perceptron** adds one crucial ingredient the McCulloch-Pitts neuron lacked: a rule for *learning* the threshold and the input weights from data, rather than having them hand-set. Its computation, formalised:

5.4.1 Inputs, Weights, and Bias

Given a real-valued input vector $\mathbf{x} = (x_1, \dots, x_d)$, the perceptron holds a weight vector $\mathbf{w} = (w_1, \dots, w_d)$ and a bias b — exactly the parameters of Figure 4.1 in Chapter 4.

5.4.2 Linear Combination

Perceptron Pre-Activation

$$z = \mathbf{w}^T \mathbf{x} + b$$

5.4.3 Activation: The Step Function

$$\text{step}(z) = \begin{cases} +1 & z \geq 0 \\ -1 & z < 0 \end{cases}$$

5.4.4 Output

Perceptron Output

$$\hat{y} = \text{sign}(z) = \text{sign}(\mathbf{w}^T \mathbf{x} + b) \in \{-1, +1\}$$

Geometrically, $\hat{y}$ answers a single question: which side of the hyperplane $\mathbf{w}^T \mathbf{x} + b = 0$ does $\mathbf{x}$ fall on? $\mathbf{w}$ is the vector normal to that hyperplane, and b shifts it away from the origin. Training a perceptron, in this geometric language, means searching for a hyperplane that puts every positively-labelled point on one side and every negatively-labelled point on the other — that is, searching for a linear separator, if one exists.

5.5 The Perceptron Learning Algorithm

5.5.1 The Update Rule

Perceptron Update Rule

For each training example $(\mathbf{x}_i, y_i)$, with $y_i \in \{-1, +1\}$:

$$\hat{y}_i = \text{sign}(\mathbf{w}^\top \mathbf{x}_i + b), \quad \mathbf{w} \leftarrow \mathbf{w} + \eta (y_i - \hat{y}_i) \mathbf{x}_i, \quad b \leftarrow b + \eta (y_i - \hat{y}_i)$$

Because $y_i, \hat{y}_i \in \{-1, +1\}$, the correction factor $(y_i - \hat{y}_i)$ is exactly zero whenever the point is already correctly classified (no update happens) and exactly ± 2 whenever it is misclassified. The rule is therefore often written in the equivalent, simpler **mistake-driven** form used throughout this chapter: *do nothing when a point is classified correctly; on a mistake, update $\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i$ and $b \leftarrow b + \eta y_i$.*

5.5.2 Geometric Interpretation: Rotating the Hyperplane

Each mistake-driven update nudges $\mathbf{w}$ directly toward $y_i \mathbf{x}_i$ — that is, it rotates the separating hyperplane so that the misclassified point moves toward the correct side of it. A single update rarely fixes everything at once (it may even newly misclassify a point that was previously correct), which is exactly why the algorithm must sweep over the data repeatedly. What it guarantees, made precise in the next section, is that this sequence of small rotations cannot continue forever if a perfect separator exists at all. Section 5.6's worked example makes this rotation literally visible, not just described.

5.5.3 The Perceptron Convergence Theorem

Perceptron Convergence Theorem — Novikoff, 1962

Suppose the training data is linearly separable with margin γ : there exists a unit vector $\mathbf{w}^{*'}$ (in the bias-augmented space, $\mathbf{x}' = (\mathbf{x}, 1)$) such that $y_i(\mathbf{w}^{*'} \cdot \mathbf{x}'_i) \geq \gamma > 0$ for every training example. Suppose also every augmented input satisfies $\|\mathbf{x}'_i\| \leq R$. Then the mistake-driven perceptron algorithm makes at most

$$\left(\frac{R}{\gamma}\right)^2$$

mistakes in total, regardless of the order in which examples are presented, and therefore halts — finds a separating hyperplane — in a finite number of steps.

Proof sketch. Track two quantities across updates: the alignment $\mathbf{w}'_k \cdot \mathbf{w}^{*'}$ and the squared norm $\|\mathbf{w}'_k\|^2$, where $\mathbf{w}'_k$ is the augmented weight vector after k mistakes. Each mistake on example i increases the alignment by at least γ (by the margin assumption) and increases the squared norm by at most R^2 (a calculation using $\|\mathbf{x}'_i\| \leq R$ and the fact that a mistake means $y_i(\mathbf{w}'_k \cdot \mathbf{x}'_i) \leq 0$, so the cross-term in $\|\mathbf{w}'_k + y_i \mathbf{x}'_i\|^2$ cannot be positive). After k mistakes, therefore, $\mathbf{w}'_k \cdot \mathbf{w}^{*'} \geq k\gamma$, while $\|\mathbf{w}'_k\| \leq \sqrt{k} R$. Cauchy-Schwarz forces $\mathbf{w}'_k \cdot \mathbf{w}^{*'} \leq \|\mathbf{w}'_k\|$ (since $\mathbf{w}^{*'}$ is a unit vector), so $k\gamma \leq \sqrt{k} R$, which rearranges to $k \leq (R/\gamma)^2$. The number of mistakes is bounded,

independent of the total number of training examples or how many times each is revisited. ■

The theorem's content is worth restating in an engineer's words: if a perfect linear separator for the data exists at all, the perceptron algorithm is *guaranteed* to find one (not necessarily the same one Section 4.4's gradient descent would find) in a finite, boundable number of corrections. This is the oldest example in this book of the **Deterministic Island** philosophy that gives the series its name: a learning procedure whose termination is not a hope or an empirical observation but a proven, certifiable guarantee, under stated and checkable assumptions.

5.6 Worked Example: Linear Separation in Industrial Data

Learning a Good/Bad Operation Boundary

Six historical operating points are labelled by a plant engineer as **good** ($y = +1$) or **bad** ($y = -1$) operation, based on two readings: temperature deviation x_1 and pressure deviation x_2 , both in engineering units above a reference baseline:

(x_1, x_2)	Label	Operation
(5, 5)	+1	good
(6, 4)	+1	good
(7, 6)	+1	good
(1, 2)	-1	bad
(2, 1)	-1	bad
(1, 1)	-1	bad

Starting from $\mathbf{w} = (0, 0)$, $b = 0$, $\eta = 1$, and sweeping repeatedly through the six points in order, the mistake-driven update rule of Section 5.5.1 fires eight times before converging, in epoch 5, to a hyperplane with zero mistakes:

Update #	Epoch	Misclassified point	y_i	New $\mathbf{w}$	New b
1	1	(1, 2)	-1	(-1, -2)	-1
2	2	(5, 5)	+1	(4, 3)	0
3	2	(1, 2)	-1	(3, 1)	-1
4	2	(2, 1)	-1	(1, 0)	-2
5	3	(2, 1)	-1	(-1, -1)	-3
6	4	(5, 5)	+1	(4, 4)	-2
7	4	(1, 2)	-1	(3, 2)	-3
8	4	(2, 1)	-1	(1, 1)	-4

The final boundary is $\mathbf{w} = (1, 1)$, $b = -4$, i.e. $x_1 + x_2 = 4$ — a clean, physically readable rule: flag operation as bad whenever the combined temperature-and-pressure deviation exceeds 4 units.




Figure 5.2: The perceptron’s separating hyperplane, mid-training (after update 6, still misclassifying two “bad” points) and at convergence (after update 8). The boundary visibly slides and rotates across training, exactly as Section 5.5.2 described, until every point sits on its correct side.

How Loose Is the Convergence Bound, in Practice?

Using the augmented vectors $\mathbf{x}'_i = (x_1, x_2, 1)$, the largest norm in this dataset is $\|\mathbf{x}'\| = \sqrt{7^2 + 6^2 + 1^2} = \sqrt{86} \approx 9.274$, so $R = \sqrt{86}$. Using the final solution’s direction as a stand-in for the (unknown, non-unique) maximum-margin $\mathbf{w}^{*}$, normalised to a unit vector, $\|(1, 1, -4)\| = \sqrt{18} = 3\sqrt{2}$, the margin achieved is $\gamma = 1/(3\sqrt{2})$, so $1/\gamma^2 = 18$. Novikoff’s bound is therefore

$$\left(\frac{R}{\gamma}\right)^2 = R^2 \cdot \frac{1}{\gamma^2} = 86 \times 18 = 1,548$$

mistakes, worst case. The algorithm actually took **8**. This 190-fold gap between the guarantee and the observed behaviour is typical, not exceptional: Novikoff’s bound must hold for *every* linearly separable dataset satisfying the stated R and γ , including deliberately adversarial orderings of pathologically-shaped data, so it is necessarily loose for any one well-behaved dataset like this one. Treat convergence guarantees of this kind as certificates that training *will* terminate, not as predictions of *how quickly*.

5.7 The Limitation: The XOR Problem

5.7.1 The Perceptron Cannot Solve XOR

Chapter 4, Section 4.5 already showed geometrically that no linear boundary separates the XOR pattern. Now that Section 5.4 has fixed a precise convention ($y \in \{-1, +1\}$, threshold at $z \geq 0$), the same fact can be proven algebraically, with no appeal to a picture at all.

5.7.2 Proof: XOR Is Not Linearly Separable

Label the interlock-alarm task of Chapter 4 with $y = +1$ for “alarm” and $y = -1$ for “no alarm.” A perceptron correctly classifying all four XOR cases would require weights w_1, w_2 and bias b satisfying, simultaneously:

$$(0, 0) \rightarrow -1 : \quad b < 0$$

$$(0, 1) \rightarrow +1 : \quad w_2 + b \geq 0$$

$$(1, 0) \rightarrow +1 : \quad w_1 + b \geq 0$$

$$(1, 1) \rightarrow -1 : \quad w_1 + w_2 + b < 0$$

Adding the second and third inequalities gives $w_1 + w_2 + 2b \geq 0$, i.e. $w_1 + w_2 \geq -2b$. The fourth inequality gives $w_1 + w_2 < -b$. Combining, $-2b \leq w_1 + w_2 < -b$, which forces $-2b < -b$, i.e. $b > 0$. This directly contradicts the first inequality, $b < 0$. No values of w_1, w_2, b satisfy all four constraints simultaneously — the system is *algebraically* inconsistent, not merely geometrically awkward. ■

5.7.3 Geometric Interpretation, Restated

The algebra above is the exact symbolic mirror of Figure 4.4 in Chapter 4: the points $(0, 0)$ and $(1, 1)$ (both -1) and the points $(0, 1)$ and $(1, 0)$ (both $+1$) cannot be separated by any single straight line, because they form the vertices of a square with same-labelled points on opposite corners — a configuration no hyperplane in two dimensions can ever bisect correctly.

5.7.4 Historical Note: Minsky, Papert, and the First AI Winter

Marvin Minsky and Seymour Papert’s 1969 book *Perceptrons* proved this limitation rigorously and generally (for a wide class of predicates, not just XOR), and its reception is widely credited with precipitating the first “AI winter”: a period of roughly a decade in which funding for neural network research collapsed, in part because the field had not yet worked out how to train networks with more than one layer, which Section 6.2 shows resolves the XOR limitation completely. The historical lesson is not that Minsky and Papert were wrong — their mathematics was correct — but that a true limitation of one specific, simple model was, for a decade, widely misread as a limitation of the entire field of connectionist computation.

Honest Boundary

The perceptron is a **linear classifier** — limited, but foundational. Its convergence guarantee (Section 5.5.3) applies *only* when the data is linearly separable; if it is not (as for XOR, or for almost any real industrial dataset with any label noise at all), the mistake-driven algorithm as stated never terminates, and will cycle indefinitely without a stopping rule.

Every idea introduced in this chapter — the linear combination $\mathbf{w}^\top \mathbf{x} + b$, the geometric picture of a separating hyperplane, even the historical failure mode of XOR — survives entirely intact into every deeper network in this book. What changes, starting in the very next chapter, is only that many such units are composed together. The perceptron is not a historical curiosity being superseded; it is the atom every molecule in the rest of Part II is built from.

6. Solving XOR — Hidden Layers and Non-linearity

6.1 What You Already Know

You already know **function composition**: $f(g(x))$, applying one function to the output of another. And if you have ever touched digital logic, you know the basic **logic gates** — AND, OR, and NAND (NOT-AND) — and that any Boolean function whatsoever can be built by wiring enough of them together:

A	B	A AND B	A OR B	A NAND B
0	0	0	0	1
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0

Chapter 5 ended on a cliffhanger: no single perceptron, however its weights are chosen, can compute XOR. This chapter resolves that cliffhanger using nothing more exotic than these two familiar ideas — composition and logic gates — wired together in a very specific way.

6.2 The Solution to XOR

6.2.1 XOR as a Composition of Simpler Gates

The identity that unlocks everything is a piece of Boolean algebra:

XOR Decomposition

$$A \oplus B = (A \text{ OR } B) \text{ AND NOT}(A \text{ AND } B) = (A \text{ OR } B) \text{ AND } (A \text{ NAND } B)$$

Verify it directly: XOR is true exactly when *at least one* input is true (the OR condition) *and* they are *not both* true (the NAND condition). Each of OR and NAND, individually, is linearly separable — each was solved by a single perceptron back in Chapter 5's framework. XOR itself is not linearly separable, but it is a simple *logical combination* of two things that are.

6.2.2 Two-Layer Implementation

Translate the decomposition directly into a network of perceptrons, using the step-function convention of Chapter 5 (threshold at $z \geq 0$, output $\{0, 1\}$ for this construction):

Hand-Engineered XOR Network

Hidden neuron 1 (computes OR): $w = (1, 1)$, $b = -0.5$ **Hidden neuron 2 (computes NAND):** $w = (-1, -1)$, $b = 1.5$ **Output neuron (computes AND of the two hidden outputs):** $w = (1, 1)$, $b = -1.5$

Verifying the Hand-Engineered Network Solves XOR

Trace all four inputs through both layers:

x_1	x_2	$h_1=\text{OR}$	$h_2=\text{NAND}$	output	target
0	0	0	1	$\text{step}(0+1-1.5)=0$	0
0	1	1	1	$\text{step}(1+1-1.5)=1$	1
1	0	1	1	$\text{step}(1+1-1.5)=1$	1
1	1	1	0	$\text{step}(1+0-1.5)=0$	0

Every row matches. Three perceptrons, none individually capable of solving XOR, solve it exactly once composed in two layers. No training was needed for this particular network — the weights were derived directly from the Boolean identity above — but Section 6.5 shows the same architecture arriving at an equivalent solution purely by gradient descent, with no hand-engineering at all.

6.2.3 Visualization: Transforming the Input Space




Figure 6.1: The hidden layer as a space transformation. **Left:** the original input space, where the two “alarm” points (filled) and two “no-alarm” points (open) interleave at the corners of a square — not linearly separable. **Right:** the same four points, passed through the hidden layer $(h_1, h_2) = (\text{OR}, \text{NAND})$. The two inputs $(0, 1)$ and $(1, 0)$, both XOR-true, land on the exact same point $(1, 1)$ in hidden space; the two XOR-false inputs land at $(0, 1)$ and $(1, 0)$ in hidden space. In this transformed space, a single straight line ($h_1 + h_2 = 1.5$, exactly the output neuron’s decision boundary) separates them trivially. This is the general mechanism by which a hidden layer solves non-linear problems: it does not bend the decision boundary in the original space so much as *re-coordinatise* the data into a new space where a straight boundary suffices.

6.3 Non-linear Activation Functions

The hand-engineered network used a hard step function, exactly like Chapter 5’s perceptron, which is instructive but, as Chapter 5 noted, untrainable by gradient descent (its derivative is zero almost everywhere). Practical networks use smooth, differentiable alternatives.

Three Standard Activation Functions

Sigmoid (met already in Chapter 4):

$$\sigma(z) = \frac{1}{1 + e^{-z}}, \quad \text{range } (0, 1)$$

Tanh (hyperbolic tangent):

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} = 2\sigma(2z) - 1, \quad \text{range } (-1, 1)$$

ReLU (Rectified Linear Unit):

$$\text{ReLU}(z) = \max(0, z), \quad \text{range } [0, \infty)$$




Figure 6.2: The three activation functions compared. Sigmoid and tanh are both S-shaped and saturate at both ends (their derivative vanishes far from zero, the source of the “vanishing gradient” problem examined in Chapter 7); tanh is simply a rescaled, zero-centred sigmoid. ReLU does not saturate for positive z at all, is trivial to differentiate (1 for $z > 0$, 0 for $z < 0$), and is the default choice in most modern deep networks precisely because of these two properties.

6.3.1 Why Non-linearity Is Essential

The reason a hidden layer needs a *non-linear* activation, and not merely a second linear layer, is a two-line algebraic fact. Suppose two layers both used the identity activation:

$$\mathbf{h} = W_1 \mathbf{x} + \mathbf{b}_1, \quad \hat{\mathbf{y}} = W_2 \mathbf{h} + \mathbf{b}_2 = W_2 (W_1 \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2 = (W_2 W_1) \mathbf{x} + (W_2 \mathbf{b}_1 + \mathbf{b}_2).$$

$\underbrace{\quad\quad\quad}_{W'}$

$\underbrace{\quad\quad\quad}_{b'}$

The composition of two linear functions is itself just one linear function, $W' \mathbf{x} + b'$ — exactly the single-layer model of Chapter 4, with no additional representational power whatsoever. This holds for any number of stacked linear layers: a hundred linear layers compute nothing a single linear layer could not compute exactly. Depth without non-linearity is not depth at all; it is expensive redundancy. Every non-trivial capability of a “deep” network, starting with its ability to solve XOR at all, depends entirely on the non-linear activation between layers.

6.4 The Universal Approximation Theorem

6.4.1 Statement

Universal Approximation Theorem — Cybenko 1989, Hornik 1991

A feedforward network with a single hidden layer, a finite number of neurons, and a suitable non-linear activation (sigmoid, tanh, and most other common choices all qualify) can approximate any continuous function on a compact subset of $\mathbb{R}^d$ to arbitrary accuracy, given enough hidden neurons.

6.4.2 Proof Sketch, via Step Functions

The intuition is a staircase approximation, the same idea used to build the Riemann integral. Two sigmoid units, with large opposing weights and offset biases, can be

combined to produce a narrow “bump” function that is approximately zero everywhere except a chosen narrow interval, where it rises to approximately one. Summing many such bumps, each scaled to match the target function’s height over its own narrow interval, produces a piecewise-constant approximation of any continuous function — exactly a Riemann sum, built out of neurons instead of rectangles. As the number of hidden neurons grows, the intervals can be made arbitrarily narrow, and the approximation error shrinks toward zero, for any continuous function on a bounded domain.

Honest Boundary

The Universal Approximation Theorem guarantees **existence**: a network with enough hidden neurons *can* represent the target function to any desired accuracy. It says nothing at all about **how many** neurons are needed for a given function and a given accuracy (the number can be, and often is, astronomically large for anything but toy problems), and it says nothing about whether gradient descent (Chapter 4, Chapter 7) will actually *find* the right weights. This is the same existence-versus-construction gap already flagged for the Shannon coding theorem (Chapter 1) and the perceptron’s representational limits (Chapter 5): a proof that a solution exists is valuable and often surprising, but it is not, and was never meant to be, a training algorithm.

6.5 Worked Example: XOR in C#

The hand-engineered network of Section 6.2 solved XOR by direct construction, using Boolean algebra. Here, the same 2-2-1 architecture (two inputs, two hidden neurons, one output neuron) is trained **purely by gradient descent** on sigmoid activations, with no hand-derived weights at all, using nothing more than the chain-rule derivatives that a full, systematic treatment arrives at in Chapter 7.

XorNetwork.cs — training a 2-2-1 network on XOR

```

1 using System;
2
3 class XorNetwork
4 {
5     static double Sigmoid(double z) => 1.0 / (1.0 + Math.Exp(-z));
6
7     static void Main()
8     {
9         double[,] x = { {0,0}, {0,1}, {1,0}, {1,1} };
10        double[] y = { 0, 1, 1, 0 };
11
12        // Fixed, explicit initial weights -- reproducible, no RNG.
13        double[,] w1 = { {0.5, -0.6}, {-0.7, 0.4} };
14        double[] b1 = { 0.1, -0.2 };
15        double[] w2 = { 0.8, -0.9 };
16        double b2 = 0.1;
17
18        const double eta = 2.0;
19        const int epochs = 30000;
20
21        for (int epoch = 0; epoch < epochs; epoch++)
22        {

```

```

23     var h = new double[4, 2];
24     var o = new double[4];
25     var delta0 = new double[4];
26     var deltaH = new double[4, 2];
27
28     for (int i = 0; i < 4; i++)
29     {
30         for (int j = 0; j < 2; j++)
31         {
32             double pre = x[i, 0] * w1[0, j] + x[i, 1] * w1[1, j]
33                 + b1[j];
34             h[i, j] = Sigmoid(pre);
35         }
36         double oPre = h[i, 0] * w2[0] + h[i, 1] * w2[1] + b2;
37         o[i] = Sigmoid(oPre);
38         delta0[i] = (o[i] - y[i]) * o[i] * (1 - o[i]);
39     }
40
41     double gradW2A = 0, gradW2B = 0, gradB2 = 0;
42     var gradW1 = new double[2, 2];
43     var gradB1 = new double[2];
44
45     for (int i = 0; i < 4; i++)
46     {
47         gradW2A += h[i, 0] * delta0[i];
48         gradW2B += h[i, 1] * delta0[i];
49         gradB2 += delta0[i];
50
51         for (int j = 0; j < 2; j++)
52         {
53             deltaH[i, j] = delta0[i] * w2[j] * h[i, j] * (1 - h
54                 [i, j]);
55             gradW1[0, j] += x[i, 0] * deltaH[i, j];
56             gradW1[1, j] += x[i, 1] * deltaH[i, j];
57             gradB1[j] += deltaH[i, j];
58         }
59     }
60
61     w2[0] -= eta * gradW2A / 4; w2[1] -= eta * gradW2B / 4; b2
62     -= eta * gradB2 / 4;
63     for (int j = 0; j < 2; j++)
64     {
65         w1[0, j] -= eta * gradW1[0, j] / 4;
66         w1[1, j] -= eta * gradW1[1, j] / 4;
67         b1[j] -= eta * gradB1[j] / 4;
68     }
69
70     // Final forward pass for reporting omitted here;
71     // see the results table below.

```

Training Results

Loss (mean squared error over all four examples) against training epoch, and the final predictions:

Epoch	MSE Loss	x_1	x_2	Predicted	Target
0	0.2530	0	0	0.009	0
5,000	0.00071	0	1	0.989	1
10,000	0.00030	1	0	0.991	1
20,000	0.00013	1	1	0.008	0
30,000 (final)	0.0000863				

Every prediction lands within 1.2% of its target, with no XOR identity, no Boolean algebra, and no hand-set weight anywhere in the process — only the generic gradient descent machinery of Chapter 4, applied to a slightly deeper composition of functions. The final trained weights (W_1 approximately $\begin{pmatrix} 6.40 & 6.41 \\ -6.52 & -6.20 \end{pmatrix}$, $\mathbf{b}_1 \approx (-3.45, 3.11)$, $\mathbf{w}_2 \approx (10.70, -10.43)$, $b_2 \approx 4.98$) bear no resemblance to the clean ± 1 , ± 0.5 weights of Section 6.2's hand-engineered solution, yet compute an almost identical function. This is a first, concrete instance of a pattern that recurs throughout the rest of this book: gradient descent routinely finds *a* solution, but rarely finds the same, humanly-readable solution a domain expert would have written by hand — a central motivation, later, for the Deterministic Islands of Chapter 12.

6.6 From XOR to Deep Networks

6.6.1 Hierarchical Feature Learning

XOR needed exactly one hidden layer to become solvable. Real industrial perception and diagnosis problems — interpreting a vibration spectrum, a camera image of a component, or a multivariate sensor trace — are vastly more structured than XOR, and benefit from **multiple** hidden layers stacked in sequence, each building on the representation the previous layer constructed. This is **hierarchical feature learning**: rather than hand-engineering features (as a domain expert might design a vibration-envelope statistic), a deep network learns, layer by layer, its own internal features, each one a re-coordination of the last, in exactly the sense Figure 6.1 made concrete for a single hidden layer.

6.6.2 Each Layer, a More Abstract Representation




Figure 6.3: Hierarchical feature learning, schematically, for an image classifier. Early layers learn low-level, local features (edges, textures) that are useful for almost any visual task; deeper layers combine these into increasingly abstract, task-specific concepts (shapes, parts, whole objects). No layer is told what to represent; each learns whatever composition of the previous layer’s outputs best reduces the training loss. Chapter 8 formalises the specific architecture (convolutional networks) that makes this hierarchy efficient for spatial data.

6.6.3 Examples: From Pixels to Objects

This hierarchy is not a metaphor; when the internal representations of a trained image classifier are inspected directly, early layers are consistently found to have learned detectors resembling edges, colour blobs, and simple textures — features general enough to be useful for almost any visual task — while the deepest layers, just before the final classification, respond to highly specific, often human-recognisable concepts (a particular texture of corrosion, the silhouette of a specific component) that are useful only for the particular task the network was trained on. The same hierarchy appears, in different guises, in a vibration-based fault classifier (raw waveform → frequency-band energy → fault signatures → fault category) and in the ILM of Chapter 10 (tokens → local syntax → semantic content → task-specific decisions).

Honest Boundary

The Universal Approximation Theorem (Section 6.4) guarantees that a sufficiently wide single hidden layer can represent any continuous function. It does **not** guarantee **generalisation**: a network that perfectly represents, or even perfectly memorises, its training data may still fail badly on new data, for exactly the bias-variance reasons established in Chapter 3. Representational capacity and generalisation are different properties, established by different theorems, and possessing one is no guarantee of the other.

Deep networks are hard to train, in ways breadth alone does not suffer from: Chapter 7 examines vanishing and exploding gradients, the non-convexity of the loss landscape, and the practical optimisation difficulties that a purely existential theorem like Universal Approximation is silent about. “A network exists that solves your problem” and “gradient descent will find it in a practical amount of time and data” are two entirely different claims, and this book will never conflate them.

7. Backpropagation — The Engine of Learning

7.1 What You Already Know

Three tools you already have, in combination, are the entire content of this chapter. You know how to differentiate a function. You know the **chain rule**: the derivative of a composition $f(g(x))$ is $f'(g(x))g'(x)$ — multiply the derivatives along the chain. And Chapter 4 gave you gradient descent: repeatedly step the parameters in the direction that most reduces the loss. **Backpropagation** is not a fourth idea layered on top of these three. It is, exactly and only, the chain rule, applied systematically to a composition of functions with a particular shape — a neural network — organised into an algorithm efficient enough to run on networks with billions of parameters. Chapter 6's worked example already used backpropagation in full, without naming it; this chapter names it, derives it completely, and shows why it scales.

7.2 The Backpropagation Algorithm

7.2.1 Forward Pass: Computing the Output

A network computes its output by pushing the input through each layer in sequence — exactly the composition structure Chapter 6 built XOR from. For a network with one hidden layer,

$$\begin{aligned} \mathbf{z}^{(1)} &= W^{(1)}\mathbf{x} + \mathbf{b}^{(1)}, & \mathbf{a}^{(1)} &= \phi(\mathbf{z}^{(1)}), \\ z^{(2)} &= \mathbf{w}^{(2)\top}\mathbf{a}^{(1)} + b^{(2)}, & a^{(2)} &= \phi(z^{(2)}) = \hat{y}, \end{aligned}$$

where ϕ is the activation function (sigmoid, tanh, or ReLU from Chapter 6). This is the **forward pass**: every intermediate value, $\mathbf{z}^{(1)}$, $\mathbf{a}^{(1)}$, $z^{(2)}$, $a^{(2)}$, is computed once, left to right, and *cached* — the caching turns out to matter enormously for the backward pass.

7.2.2 Loss Functions: MSE and Cross-Entropy, Recalled

The forward pass ends at a loss, comparing $\hat{y}$ against the true target y . This book has already introduced both losses backpropagation is run against most often:

$$\text{MSE (Chapter 3): } L = \frac{1}{2}(y - \hat{y})^2,$$

$$\text{Cross-entropy (Chapter 4): } L = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})].$$

Nothing about backpropagation cares which loss is used; the loss is simply the last function in the chain the chain rule is applied through, and Section 7.3 derives the general case explicitly for MSE.

7.2.3 Backward Pass: Computing Gradients

The **backward pass** computes the loss's gradient with respect to *every* weight in the network, by walking the chain rule backward through the exact same computation graph the forward pass built, reusing the cached values at each step. The key algorithmic insight is that the chain rule, applied layer by layer, lets each layer's gradient be computed from only two things: the gradient "arriving" from the layer after it, and quantities already cached during the forward pass. This is what makes backpropagation efficient: without it, computing the gradient with respect to $W^{(1)}$ directly would require differentiating straight through every subsequent layer's function from scratch, redoing work the later layers had already effectively done once. Section 7.3 makes every one of these "arriving gradients" explicit, layer by layer, for the exact architecture of Chapter 6's worked example.

7.2.4 Weight Update

The gradients the backward pass produces are consumed exactly as Chapter 4 specified:

Weight Update, Recalled from Chapter 4

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L$$

applied independently to every weight matrix and bias vector in the network. Backpropagation computes $\nabla_{\theta} L$ efficiently; gradient descent (or one of Section 7.6's refinements of it) is what actually moves θ .

7.3 Detailed Derivation for a 2-2-1 Network

7.3.1 Step-by-Step Gradient Computation

Take the exact architecture Chapter 6 trained on XOR: two inputs, two sigmoid hidden units, one sigmoid output unit, MSE loss on a single example $(\mathbf{x}, y)$. Write $a_j^{(1)}$ for hidden unit j 's activation and $a^{(2)}$ for the output. The chain rule is applied in four steps, output to input.

Step 1 — output error signal. Define $\delta^{(2)} := \frac{\partial L}{\partial z^{(2)}}$. By the chain rule through $L = \frac{1}{2}(y - a^{(2)})^2$ and $a^{(2)} = \sigma(z^{(2)})$:

$$\delta^{(2)} = \frac{\partial L}{\partial a^{(2)}} \cdot \frac{\partial a^{(2)}}{\partial z^{(2)}} = -(y - a^{(2)}) \cdot \sigma'(z^{(2)}) = (a^{(2)} - y) a^{(2)} (1 - a^{(2)}),$$

using the sigmoid's own convenient derivative, $\sigma'(z) = \sigma(z)(1 - \sigma(z))$.

Step 2 — output-layer gradients. Since $z^{(2)} = \sum_j w_j^{(2)} a_j^{(1)} + b^{(2)}$:

$$\frac{\partial L}{\partial w_j^{(2)}} = \delta^{(2)} \cdot a_j^{(1)}, \quad \frac{\partial L}{\partial b^{(2)}} = \delta^{(2)}.$$

Step 3 — propagate the error backward. Define $\delta_j^{(1)} := \frac{\partial L}{\partial z_j^{(1)}}$. The chain rule

must now pass *through* $w_j^{(2)}$, because $z^{(2)}$ depends on $a_j^{(1)}$ with sensitivity $w_j^{(2)}$:

$$\delta_j^{(1)} = \frac{\partial L}{\partial z^{(2)}} \cdot \frac{\partial z^{(2)}}{\partial a_j^{(1)}} \cdot \frac{\partial a_j^{(1)}}{\partial z_j^{(1)}} = \delta^{(2)} \cdot w_j^{(2)} \cdot a_j^{(1)}(1 - a_j^{(1)}).$$

This single line is backpropagation's namesake move: the output layer's error signal $\delta^{(2)}$ is carried *backward* across the weight $w_j^{(2)}$ that connected the two layers in the forward pass, then corrected by the hidden unit's own local derivative.

Step 4 — hidden-layer gradients. Since $z_j^{(1)} = \sum_k W_{jk}^{(1)} x_k + b_j^{(1)}$:

$$\frac{\partial L}{\partial W_{jk}^{(1)}} = \delta_j^{(1)} \cdot x_k, \quad \frac{\partial L}{\partial b_j^{(1)}} = \delta_j^{(1)}.$$

Every one of these four steps is visible, unrolled and summed over all four training examples, in the C# code of Chapter 6, Section 6.5 — `delta0` is $\delta^{(2)}$, `deltaH` is $\delta^{(1)}$, and every gradient accumulation line matches one of the four boxed formulas above exactly.

7.3.2 Table of Variables and Partial Derivatives

Symbol	Meaning	Computed as
$z_j^{(1)}$	hidden pre-activation	$\sum_k W_{jk}^{(1)} x_k + b_j^{(1)}$
$a_j^{(1)}$	hidden activation	$\sigma(z_j^{(1)})$
$z^{(2)}$	output pre-activation	$\sum_j w_j^{(2)} a_j^{(1)} + b^{(2)}$
$a^{(2)}$	network output	$\sigma(z^{(2)})$
$\delta^{(2)}$	output error signal	$(a^{(2)} - y) a^{(2)}(1 - a^{(2)})$
$\delta_j^{(1)}$	hidden error signal	$\delta^{(2)} w_j^{(2)} a_j^{(1)}(1 - a_j^{(1)})$
$\partial L / \partial w_j^{(2)}$	output weight gradient	$\delta^{(2)} a_j^{(1)}$
$\partial L / \partial W_{jk}^{(1)}$	hidden weight gradient	$\delta_j^{(1)} x_k$

7.4 Worked Example: Training XOR with Backprop in C#

7.4.1 Code, Annotated

Chapter 6's implementation hand-unrolled every index for a fixed 2-2-1 shape. The version below is organised around the four derivation steps of Section 7.3.1 directly, so that each block of code is labelled with the formula it implements — the same training run, the same fixed initial weights, but written to be read alongside the mathematics rather than alongside the truth table.

XorBackprop.cs — the four derivation steps, in code

```

1 using System;
2
3 class XorBackprop
4 {
5     static double Sigmoid(double z) => 1.0 / (1.0 + Math.Exp(-z));
6     static double SigmoidPrime(double a) => a * (1 - a); // a already =
        sigmoid(z)

```

```

7
8 static void Main()
9 {
10     double[,] x = { {0,0}, {0,1}, {1,0}, {1,1} };
11     double[] y = { 0, 1, 1, 0 };
12
13     double[,] w1 = { {0.5, -0.6}, {-0.7, 0.4} };
14     double[] b1 = { 0.1, -0.2 };
15     double[] w2 = { 0.8, -0.9 };
16     double b2 = 0.1;
17     const double eta = 2.0;
18
19     for (int epoch = 0; epoch < 30000; epoch++)
20     {
21         var a1 = new double[4, 2];
22         var a2 = new double[4];
23         var delta2 = new double[4];
24         var delta1 = new double[4, 2];
25
26         // --- Forward pass (Section 7.2.1) ---
27         for (int i = 0; i < 4; i++)
28         {
29             for (int j = 0; j < 2; j++)
30             {
31                 double z1 = x[i, 0] * w1[0, j] + x[i, 1] * w1[1, j]
32                     + b1[j];
33                 a1[i, j] = Sigmoid(z1);
34             }
35             double z2 = a1[i, 0] * w2[0] + a1[i, 1] * w2[1] + b2;
36             a2[i] = Sigmoid(z2);
37
38             // --- Step 1: output error signal, delta^(2) ---
39             delta2[i] = (a2[i] - y[i]) * SigmoidPrime(a2[i]);
40
41             // --- Step 2: output-layer gradients, dL/dw2 and dL/db2
42             ---
43             double gradW2_0 = 0, gradW2_1 = 0, gradB2 = 0;
44             for (int i = 0; i < 4; i++)
45             {
46                 gradW2_0 += delta2[i] * a1[i, 0];
47                 gradW2_1 += delta2[i] * a1[i, 1];
48                 gradB2 += delta2[i];
49
50                 // --- Step 3: propagate backward, delta^(1) ---
51                 for (int j = 0; j < 2; j++)
52                     delta1[i, j] = delta2[i] * w2[j] * SigmoidPrime(a1[i, j]);
53             }
54
55             // --- Step 4: hidden-layer gradients, dL/dw1 and dL/db1
56             ---
57             var gradW1 = new double[2, 2];
58             var gradB1 = new double[2];
59             for (int i = 0; i < 4; i++)
60                 for (int j = 0; j < 2; j++)

```

```

60         gradW1[0, j] += delta1[i, j] * x[i, 0];
61         gradW1[1, j] += delta1[i, j] * x[i, 1];
62         gradB1[j] += delta1[i, j];
63     }
64
65     // --- Update (Section 7.2.4) ---
66     w2[0] -= eta * gradW2_0 / 4; w2[1] -= eta * gradW2_1 / 4;
67     b2 -= eta * gradB2 / 4;
68     for (int j = 0; j < 2; j++)
69     {
70         w1[0, j] -= eta * gradW1[0, j] / 4;
71         w1[1, j] -= eta * gradW1[1, j] / 4;
72         b1[j] -= eta * gradB1[j] / 4;
73     }
74 }
75 }

```

7.4.2 Visualization of the Loss Evolution




Figure 7.1: Training loss against epoch, on a logarithmic axis, for the XOR network of Section 7.4.1. Loss falls sharply between roughly epoch 800 and 2,000 — the network abruptly “finding” the XOR solution — then decreases slowly and smoothly for the remaining 28,000 epochs as it refines the margin around each of the four points. This two-phase shape (a sharp drop followed by a long, shallow tail) is typical of small networks trained past the point where they already classify every example correctly.

7.4.3 Final Weight Values

Parameter	Final value
$W^{(1)}$	$\begin{pmatrix} 6.40 & 6.41 \\ -6.52 & -6.20 \end{pmatrix}$
$\mathbf{b}^{(1)}$	$(-3.45, 3.11)$
$\mathbf{w}^{(2)}$	$(10.70, -10.43)$
$b^{(2)}$	4.98

These are, to the precision shown, identical to Chapter 6’s reported values — as

they must be, since both programs implement the same gradients from the same starting point; Section 7.3 simply made explicit, by name, the four chain-rule steps Chapter 6’s code had already been performing.

7.5 Challenges

7.5.1 Vanishing and Exploding Gradients

Look again at Section 7.3.1’s Step 3: the backward-propagated error $\delta_j^{(1)}$ is $\delta^{(2)}$ multiplied by a weight *and* by a local activation derivative. In a network with many hidden layers, this multiplication chains across every layer, and the sigmoid’s derivative $\sigma'(z) = \sigma(z)(1 - \sigma(z))$ has a maximum value of exactly 0.25, at $z = 0$ (visible in Figure 6.2 saturating toward zero away from the origin).

How Fast Gradients Vanish Through Sigmoid Layers

If each layer’s local derivative is at best 0.25, the gradient reaching a weight L layers back from the output is bounded, very roughly, by 0.25^L times what arrived at the output layer:

Layers back (L)	Gradient scale factor $\leq 0.25^L$
1	0.25
4	0.0039
8	1.53×10^{-5}
10	9.54×10^{-7}
12	5.96×10^{-8}

By ten or twelve sigmoid layers, the gradient reaching the earliest weights is, in the worst case, less than one part in a million of the gradient at the output — those early weights effectively stop learning at all, however long training continues. This single table is the entire, quantitative reason ReLU (Figure 6.2, constant derivative of exactly 1 for any positive input, not capped at 0.25) displaced sigmoid and tanh as the default hidden activation once networks grew beyond a handful of layers, and the direct motivation for the residual connections and normalisation layers examined in Chapter 8.

The mirror-image failure, **exploding gradients**, occurs when weights are large enough that the chained product grows rather than shrinks, sending gradients (and the resulting weight updates) to numerically unstable, enormous values; it is addressed in practice by gradient clipping (capping the gradient’s norm before the update) and by the same careful weight initialisation schemes that also help against vanishing gradients.

7.5.2 Local Minima

A loss surface with many parameters can, in principle, have many **local minima** — points where the gradient is zero but the loss is not globally lowest. This was a serious, practical worry in early neural network research. It has turned out to matter far less than feared, for a specific, quantifiable reason: at a critical point (zero gradient) in a d -dimensional space, being a genuine local minimum requires the loss to curve *upward* in *every one* of the d directions simultaneously (every eigenvalue of the Hessian matrix positive). As d grows into the thousands or millions, as it

does for any real neural network, the odds of a randomly encountered critical point having every single curvature direction positive, rather than some mix of positive and negative, shrink rapidly under quite general random-matrix assumptions. Most critical points a large network’s training actually encounters are, empirically and theoretically, **saddle points** — flat in some directions, curving down in others — not true local minima, and gradient descent (particularly its stochastic variant, Section 4.4.1, whose noise helps it slide off a flat saddle) typically escapes them without special intervention.

7.5.3 Overfitting

Chapter 3’s bias-variance tradeoff applies to neural networks with no modification: a sufficiently large network, trained long enough, can drive training loss arbitrarily close to zero (the Universal Approximation Theorem of Chapter 6 all but guarantees the capacity to do so) while generalising poorly to new data. Every regularisation tool from Chapter 3 — L1/L2 penalties, cross-validation — applies directly; Section 7.6.4 adds one more, dropout, specific to neural networks.

7.5.4 Learning Rate Selection

Chapter 4’s elliptical-bowl example already showed that a fixed learning rate is a compromise across every parameter direction at once. The same tension exists even in the simplest possible case, a single parameter:




Figure 7.2: Gradient descent on the toy loss $L(w) = (w - 3)^2$ (minimum at $w = 3$, dashed line) from $w_0 = 0$, for three learning rates. Too small ($\eta = 0.05$) crawls toward the minimum, wasting most of a fixed training budget. Well-tuned ($\eta = 0.4$) reaches the minimum in about five steps. Too large ($\eta = 1.05$) overshoots the minimum on every step by more than it started, and the oscillation’s amplitude grows without bound — numerically identical, in miniature, to the loss “exploding” during real network training.

The safe range for this toy loss ($0 < \eta < 1$, exactly, since the update is a linear recurrence $w_{t+1} = (1 - 2\eta)w_t + 6\eta$ that converges only while $|1 - 2\eta| < 1$) is a closed-form fact for a quadratic; for a real network’s loss landscape, the equivalent bound depends on the local curvature at every point along the trajectory, changes as training proceeds, and is never known in closed form, which is exactly why Section 7.6 abandons the fixed learning rate and lets the optimiser adapt it instead.

7.6 Optimization Techniques

7.6.1 Momentum and Nesterov Momentum

Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} + (1 - \beta) \nabla_{\theta} L, \quad \theta \leftarrow \theta - \eta \mathbf{v}$$

Momentum keeps a running, exponentially-decayed average $\mathbf{v}$ of past gradients (typically $\beta \approx 0.9$), rather than reacting only to the current step's gradient. This directly targets Chapter 4's ill-conditioned-bowl problem: in the shallow direction, consistent small gradients accumulate into a large, confident step; in a direction where the gradient keeps flipping sign (oscillation), the average partially cancels, damping the oscillation. **Nesterov momentum** refines this by evaluating the gradient at the position momentum is already about to carry the parameters to, rather than at the current position — a small look-ahead correction that measurably speeds convergence.

7.6.2 AdaGrad, RMSProp, and Adam

Adaptive Learning-Rate Optimisers

AdaGrad divides each parameter's learning rate by the square root of the sum of its past squared gradients, giving frequently and strongly updated parameters a shrinking effective rate. **RMSProp** replaces AdaGrad's ever-growing sum with an exponentially decayed moving average of squared gradients, so the effective rate does not decay to zero over a long training run. **Adam** (Adaptive Moment Estimation) combines momentum's running average of the gradient itself with RMSProp's running average of the squared gradient, giving each individual parameter both its own effective step size and its own momentum — directly solving the Chapter 4 elliptical-bowl problem, where a single global learning rate was, by construction, too small for one direction and too large for another.

Adam is, as of this writing, the default optimiser for the vast majority of neural network training in practice, including every network described from Chapter 10 onward in this book.

7.6.3 Batch Normalization

As training proceeds, the distribution of each layer's inputs shifts, because every earlier layer's weights are simultaneously changing too — a moving target sometimes called "internal covariate shift." **Batch normalization** re-standardises each layer's pre-activations (subtracting the batch mean, dividing by the batch standard deviation, then applying a small learned rescaling) before the activation function is applied, stabilising this shifting distribution and, empirically, allowing substantially higher learning rates and faster convergence.

7.6.4 Dropout

Dropout randomly zeroes out a fraction of neurons (typically 20-50%) on each training step, forcing the remaining neurons to not over-rely on any single other neuron's

exact output. It is best understood as an efficient, implicit form of the ensembling and regularisation ideas of Chapter 3: training with dropout approximates training a very large number of thinned sub-networks simultaneously and averaging their predictions, which directly attacks the variance term of Chapter 3's bias-variance decomposition.

Honest Boundary

Backpropagation is **just the chain rule** — nothing more. Every formula in Section 7.3 was ordinary calculus, applied mechanically; no additional mathematical machinery, and no conceptual leap beyond what Chapter 4 already introduced, was needed to derive it. Its practical importance comes entirely from *efficiency*: it computes every parameter's gradient in roughly the same time as one forward pass, rather than the far larger cost a naive, layer-by-layer re-differentiation would require.

Training a neural network is **non-convex optimisation with no convergence guarantees**. Unlike the perceptron's Novikoff theorem (Chapter 5) or the Shannon coding theorem (Chapter 1), nothing in this chapter promises that gradient descent, with or without momentum or Adam, will find a good set of weights, or any particular set of weights, in any bounded number of steps. What works in practice is supported by substantial empirical experience and a growing body of theory (Section 7.5.2's saddle-point argument among it), but it is not, and should never be represented as, a certified guarantee. Every claim of "the model learned X" in the chapters that follow should be read with this absence of guarantee in mind — and it is precisely this absence that motivates the Deterministic Islands introduced in Chapter 12, for the specific situations where a guarantee, not just a good empirical track record, is required.

8. Modern Architectures

8.1 What You Already Know

Every architecture in this chapter is built from the same two ideas you already have in hand. A **layer** is a function that takes a vector (or a grid, or a sequence of vectors) and produces another one — Chapter 4’s single neuron generalised to many neurons computed side by side. A **module** is a layer, or a small group of layers wired together, treated as a single reusable unit that can itself be composed with other modules — exactly Chapter 6’s lesson that composition, not any single exotic ingredient, is what turns simple pieces into powerful systems. Every architecture below (CNN, RNN, Transformer, autoencoder, GAN) is best understood not as a wholly new kind of mathematics, but as a new *wiring pattern* connecting the same neurons, trained by the same backpropagation, encoding a different assumption about the *structure* of the data it is shown.

8.2 Convolutional Neural Networks (CNNs)

8.2.1 Locality, Weight Sharing, and Pooling

A fully-connected layer (every input connected to every neuron, as in every network so far in this book) makes no assumption about spatial structure: pixel 1 and pixel 2 are treated exactly like any other pair of inputs, whether they are neighbours in an image or on opposite corners. For image and spatial-sensor data, this throws away exactly the structure that matters most: nearby pixels are typically strongly related, and a useful visual feature (an edge, a corner) can appear *anywhere* in the image. **Convolutional layers** encode both facts directly into the architecture, through three ideas:

Locality, Weight Sharing, Pooling

Locality: each output value depends only on a small, local patch of the input (a 3×3 or 5×5 window), not the entire input. **Weight sharing:** the *same* small filter (kernel) of weights is applied at every position of the input, rather than learning a separate weight for every pixel-position pair. A feature useful in one corner of an image is, by construction, detected identically anywhere else. **Pooling:** a fixed (non-learned) downsampling operation, typically taking the maximum or average value within a small window, that reduces spatial resolution and gives the network some tolerance to small shifts in where a feature appears.

A Convolution and a Max-Pool, by Hand

Take a small 4×4 input (perhaps one channel of a thermal image patch of a bearing housing) and a 3×3 vertical-edge-detecting kernel:

$$X = \begin{pmatrix} 1 & 2 & 0 & 1 \\ 0 & 1 & 3 & 1 \\ 2 & 1 & 0 & 2 \\ 1 & 0 & 1 & 3 \end{pmatrix}, \quad K = \begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix}.$$

Sliding K across every valid 3×3 window of X (stride 1, no padding) and summing the elementwise product at each position produces a 2×2 **feature map**:

$$X * K = \begin{pmatrix} 0 & 0 \\ -1 & -4 \end{pmatrix}.$$

For instance, the bottom-right entry comes from the window $\begin{pmatrix} 1 & 3 & 1 \\ 1 & 0 & 2 \\ 0 & 1 & 3 \end{pmatrix}$: $(1 \cdot 1 + 3 \cdot 0 + 1 \cdot -1) + (1 \cdot 1 + 0 \cdot 0 + 2 \cdot -1) + (0 \cdot 1 + 1 \cdot 0 + 3 \cdot -1) = 0 - 1 - 3 = -4$. The large negative values mark where brightness drops sharply left-to-right; this single fixed kernel is a hand-designed edge detector, but Section 7.3's backpropagation applies to convolutional weights exactly as it does to any other weight — in a trained CNN, kernels like this one are *learned*, not hand-specified.

Applying 2×2 **max-pooling** (stride 2) directly to X takes the maximum of each non-overlapping 2×2 block:

$$\text{maxpool}(X) = \begin{pmatrix} 2 & 3 \\ 2 & 3 \end{pmatrix}.$$

The output is half the resolution of X in each dimension, and each entry is unaffected by which exact pixel, within its 2×2 block, held the maximum value — a small, deliberate amount of positional slack.




Figure 8.1: The convolution of the worked example: a 3×3 kernel slides across the 4×4 input (the burnt-orange box marks the window that produces the highlighted -4 output), and each position produces one entry of the 2×2 feature map. Weight sharing means the exact same nine kernel weights are reused at all four window positions.

8.2.2 Applications: Image and Signal Processing in Industry

CNNs are the default architecture for any data with strong local, translation-relevant structure: visual inspection of welds, corrosion, or component wear from camera images; defect detection on a production line; and, less obviously, one-dimensional convolution over *time* applied to vibration or acoustic waveforms, where “locality” means a short time window and “translation invariance” means a characteristic fault signature should be detected whenever it occurs in the signal, not only at one fixed time offset.

8.3 Recurrent Neural Networks (RNNs)

8.3.1 Sequences, Hidden State, and Backpropagation Through Time

Convolutional weight sharing exploits structure in *space*. **Recurrent Neural Networks** exploit structure in *time*: they process a sequence $x_1, x_2, \dots, x_T$ one step at a time, while maintaining a **hidden state** h_t that summarises everything relevant seen so far.

The Recurrent Update

$$h_t = \phi(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$

The same weight matrices W_{hh} and W_{xh} are reused at *every* time step — weight sharing again, this time across time rather than space.




Figure 8.2: An RNN “unrolled” through time. Every timestep reuses the identical weight matrices W_{hh} and W_{xh} — the network in the box is the same network at every step, not a different one per timestep. **Backpropagation through time** (BPTT) is exactly Chapter 7’s backpropagation applied to this unrolled diagram, treated as one long feedforward computation graph: gradients flow backward through every timestep, and the shared weight’s total gradient is the *sum* of its gradient contributions at every timestep it was used.

BPTT inherits Chapter 7’s vanishing-gradient problem directly, and often more severely: a gradient flowing back through 50 or 100 timesteps multiplies through 50 or 100 copies of essentially the same Jacobian, and Section 7.5.1’s 0.25^L decay applies with L now equal to the sequence length, not the network’s depth.

8.3.2 LSTM and GRU: Gates and Memory

The **Long Short-Term Memory** (LSTM) network was designed specifically to fix this. It adds a separate **cell state** C_t , a kind of protected memory conveyor belt, and three learned **gates** (each a sigmoid-activated layer, producing values in $[0, 1]$) that

control how much information passes through) that regulate what enters, leaves, and is forgotten from that memory:

LSTM Gate Equations

Forget gate: $f_t = \sigma(W_f[h_{t-1}, x_t] + b_f)$ **Input gate:** $i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$
Candidate cell update: $\tilde{C}_t = \tanh(W_C[h_{t-1}, x_t] + b_C)$ **Cell state update:**
 $C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$ **Output gate:** $o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$ **Hidden state:** $h_t = o_t \odot \tanh(C_t)$

($\odot$ is elementwise multiplication; $[h_{t-1}, x_t]$ denotes concatenation.) The forget gate can learn to hold $f_t \approx 1$ for information that should persist across very many timesteps, letting the cell state's gradient flow backward through the $C_t = f_t \odot C_{t-1} + \dots$ update largely unimpeded by the multiplicative 0.25^L decay of Section 7.5.1 — this additive, gated pathway is the entire mechanism by which LSTMs learn long-range dependencies that plain RNNs cannot. The **GRU** (Gated Recurrent Unit) simplifies this to two gates (update z_t and reset r_t) and no separate cell state, trading a little representational flexibility for fewer parameters and faster training, with broadly comparable performance in practice.

8.3.3 Applications: Time-Series Prediction and Anomaly Detection

RNNs, LSTMs, and GRUs were, until the rise of the Transformer (Section 8.5), the default architecture for plant telemetry: multi-step-ahead prediction of a process variable, remaining-useful-life estimation (Chapter 15) from a sensor history, and sequence-based anomaly detection, where a model trained to predict “normal” next values flags a large prediction error as a candidate anomaly.

8.4 The Attention Mechanism

8.4.1 The Key-Value-Query Formulation

Both weight sharing tricks so far are architectural: they hard-code an assumption about structure into the network's wiring. **Attention** instead lets the network *learn, per input, which other parts of the input to focus on* — a content-based, dynamic form of weight sharing rather than a positional, fixed one. Every input position produces three learned projections of itself: a **query** (what this position is “looking for”), a **key** (what this position “offers” to others), and a **value** (what this position actually contributes if attended to).

Scaled Dot-Product Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V$$

$QK^\top$ measures how well every query matches every key (a large dot product means “highly relevant”); dividing by $\sqrt{d_k}$ (the key dimension) keeps these scores in a numerically stable range before the softmax; the softmax turns the scores into a proper probability distribution (Chapter 1) over positions to attend to; and the weighted sum of values is the output — each position's new representation is a learned, input-dependent blend of every other position's contribution.

Attention on Two Tokens, in Full

Two tokens with 2-dimensional embeddings $\mathbf{x}_1 = (1, 0)$, $\mathbf{x}_2 = (0, 1)$, and fixed (for this illustration) projection matrices

$$W_Q = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad W_K = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad W_V = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}.$$

Projecting: $Q = XW_Q = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$, $K = XW_K = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$, $V = XW_V = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$. With $d_k = 2$, the score matrix is

$$\frac{QK^\top}{\sqrt{2}} = \begin{pmatrix} 0 & 0.7071 \\ 0.7071 & 0 \end{pmatrix},$$

and row-wise softmax gives the attention weights

$$\begin{pmatrix} 0.330 & 0.670 \\ 0.670 & 0.330 \end{pmatrix},$$

so token 1's new representation is $0.330 V_1 + 0.670 V_2 = (1.000, 0.330)$, and token 2's is $0.670 V_1 + 0.330 V_2 = (1.000, 0.670)$. Each token's output is now a genuine *blend* of both tokens' values, weighted by how relevant the attention mechanism learned the other token to be — and, critically, this blending weight was computed from the tokens' *content* (Q and K), not from their fixed position in the sequence.

8.4.2 Self-Attention and Multi-Head Attention

When Q , K , and V are all projections of the *same* input sequence (as in the worked example above), the mechanism is called **self-attention**: every position attends to every other position in the same sequence, including itself. **Multi-head** attention runs several independent self-attention computations ("heads") in parallel, each with its own learned W_Q, W_K, W_V , and concatenates their outputs — allowing different heads to specialise in different kinds of relationship (one head might learn to track nearby words, another long-range dependencies, in a language model; one might track short-term correlation and another a slower periodic pattern, in a multivariate sensor sequence).

8.4.3 Positional Encoding

Unlike an RNN, self-attention has no inherent notion of order at all: every position attends to every other position symmetrically, so shuffling the input sequence would leave attention's output (reshuffled, but otherwise unchanged) structurally identical. **Positional encoding** restores order information by adding a fixed, deterministic pattern to each position's input embedding before it ever reaches the attention layers, most commonly

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{2i/d}}\right), \quad PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{2i/d}}\right),$$

a different sinusoid frequency for each embedding dimension, chosen so that the encoding for any fixed relative offset can be recovered as a linear function of the encoding at any position — giving the network an efficient way to learn relative-position relationships on top of the raw encoding.

8.5 The Transformer Architecture

8.5.1 Encoder-Decoder Structure




Figure 8.3: The Transformer’s encoder-decoder structure, simplified. The encoder (left) builds a contextual representation of the input sequence through N stacked self-attention/feed-forward blocks. The decoder (right) generates the output sequence one token at a time, using *masked* self-attention (so a position cannot attend to future positions it hasn’t generated yet) and a second attention layer that queries the encoder’s output directly — the mechanism by which the decoder consults the input sequence at every generation step.

8.5.2 Why Transformers Dominate Modern AI

Two properties, both direct consequences of replacing recurrence with attention, explain the Transformer’s dominance since 2017. **Parallelisation:** an RNN must process $x_1, \dots, x_T$ strictly in sequence, since h_t depends on h_{t-1} ; a Transformer’s self-attention computes every position’s output simultaneously, letting training exploit modern parallel hardware (GPUs, Chapter 11) far more efficiently. **Long-range dependencies:** in an RNN, information from position 1 reaches position T only after passing through $T - 1$ recurrent steps, each one a further opportunity for Section 7.5.1’s vanishing-gradient decay; in self-attention, position 1 attends to position T *directly*, in a single step, regardless of how far apart they are.

8.5.3 Applications: Language Models, Time-Series, and Multimodal Data

The Transformer, introduced for machine translation, is now the dominant architecture for large language models (the ILM of Chapter 10 is built on it), and has been adapted successfully to raw time-series forecasting (attention over recent timesteps in place of recurrence), to images (Vision Transformers, treating image patches as a sequence of “tokens”), and to multimodal systems that attend jointly across text, sensor, and image tokens within a single unified architecture.

8.6 Autoencoders and Variational Autoencoders (VAEs)




Figure 8.4: An autoencoder: an encoder compresses the input into a low-dimensional latent code z (Chapter 1’s rate-distortion theory, made architectural), and a decoder reconstructs the input from z alone. Trained to minimise reconstruction error, the bottleneck forces the network to discover a compact representation of whatever regularity the input data has.

A plain autoencoder’s latent space has no particular structure; a **Variational Autoencoder** (VAE) fixes this by making the encoder output a *distribution* over z (typically a Gaussian mean and variance) rather than a single point, and adding a penalty that pulls this distribution toward a simple prior:

VAE Loss

$$\mathcal{L}_{\text{VAE}} = \mathbb{E}[\|\mathbf{x} - \hat{\mathbf{x}}\|^2] + D_{\text{KL}}(q(z | \mathbf{x}) \| p(z))$$

(first term: reconstruction loss. second term: the regulariser.)

The second term is exactly the Kullback-Leibler divergence defined in Chapter 2, Section 2.9, here forcing the learned latent distribution $q(z | \mathbf{x})$ toward a simple prior $p(z)$ (typically a standard Gaussian) rather than letting it collapse to disconnected points anywhere in latent space. This is precisely the Information Bottleneck tradeoff of Chapter 1, Section 1.5.5, given an explicit architecture and a trainable objective: compress (the KL term, pulling toward a simple, low-information prior) against predict/reconstruct (the reconstruction term), with the balance between them a direct, literal instance of β in that earlier, more abstract formulation.

8.7 Generative Adversarial Networks (GANs)

A **GAN** trains two networks against each other: a **generator** G that maps random noise z to synthetic data $G(z)$, and a **discriminator** D that tries to distinguish real data from G ’s output.

The GAN Minimax Objective

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

D is trained to maximise this (correctly label real data as real, $D(\mathbf{x}) \rightarrow 1$, and generated data as fake, $D(G(z)) \rightarrow 0$); G is trained to *minimise* it, i.e. to fool D into predicting $D(G(z)) \rightarrow 1$. Training alternates between the two, each improving against the other, ideally converging to a generator whose output distribution is indistinguishable from the real data distribution. GANs have historically produced

some of the most visually convincing synthetic images of any generative approach, at the cost of a training procedure (a two-player game, rather than a single minimisation) that is notoriously less stable than the VAE's straightforward gradient descent on one objective.

8.8 Architecture Selection for Industrial Data

Data type	Structure exploited	Typical architecture
Time series (single/few sensors)	temporal locality	LSTM/GRU, 1D CNN
Time series (many sensors, long-range)	long-range dependence	Transformer
Images (visual inspection)	spatial locality	CNN
Text (manuals, logs, reports)	sequential/semantic	Transformer
Multimodal (sensors + text + images)	mixed structure	fused/multi-branch

This table is a starting point, not a rule: the honest answer to “which architecture” is almost always “try the simplest one that respects your data’s structure, and let cross-validation (Chapter 3) decide,” since the table’s entries interact (a long univariate series may do perfectly well in an LSTM; a short, low-dimensional image patch may not need a CNN at all).

Honest Boundary

Architecture is art as much as science. Every architecture in this chapter encodes an assumption about structure (locality for CNNs, sequential order for RNNs, learned relevance for attention), and matching that assumption to the actual structure of the data is a matter of judgment, experimentation, and domain knowledge, not a formula this chapter, or any other, can fully mechanise.

No architecture is universally best. This is the No Free Lunch theorem of Chapter 3, Section 3.7.2, applied one level up: just as no learning algorithm dominates on every possible problem, no architecture dominates on every possible data structure. A Transformer applied to a short, simple sequence and a CNN applied to data with no spatial structure at all can both underperform a far simpler model, not despite their sophistication but because of it — representational power spent modelling structure that was never there in the first place.

9. Supervised, Unsupervised, and Reinforcement Learning

9.1 What You Already Know

You already carry the vocabulary this chapter formalises. You know the difference between **labeled data** (a fault record that says not just what the sensors read, but which specific fault it was) and **unlabeled data** (a year of sensor logs with no such annotation at all, because no one has ever gone through and tagged it). And if you have ever tuned a PID controller by rewarding yourself mentally for a smooth response and wincing at an overshoot, you already have an intuition for a **reward signal**: a number that says how well things are going, without saying what the *correct action* would have been. These three distinctions — labeled, unlabeled, reward-driven — define the three learning paradigms this chapter covers, and, not coincidentally, every one of them will turn out to be a piece the ILM (Part III, starting next chapter) needs.

9.2 Supervised Learning

Supervised learning is every technique this book has built so far, from Chapter 3's linear regression to Chapter 8's Transformer: a model is trained on input-output pairs $(\mathbf{x}_i, y_i)$ and learns to predict y from $\mathbf{x}$. **Classification** predicts a discrete label (which of five known fault types is present); **regression** predicts a continuous quantity (the remaining useful life of a bearing, in hours). Every tool developed in Chapters 3-8 — loss functions, gradient descent, backpropagation, CNNs, Transformers — applies without modification; what changes chapter to chapter is only the data and the architecture, never the underlying estimation logic of Chapter 2's MLE.

→ ILM

Supervised learning trains the ILM's **task-specific heads**: a classification head that sorts an incoming maintenance report into one of a known set of fault categories, a regression head that predicts a numeric quantity (RUL, expected downtime) from a description or a sensor window. These heads sit on top of the shared Transformer backbone (Chapter 8, Chapter 10) and are exactly what Chapter 11's **fine-tuning** adapts to a specific industrial task.

9.3 Unsupervised Learning

Unsupervised learning finds structure in data with no labels at all. Three tasks recur throughout the rest of this book.

9.3.1 Clustering

Clustering groups similar points together without being told what the groups mean. The simplest and still most widely used algorithm, ***k*-means**, alternates two steps: assign each point to its nearest of k centroids, then move each centroid to the mean of the points assigned to it.

Clustering Equipment by Operating Regime, in 3D

Six pump units report three sensor readings each — a 3D vector $(\Delta T, \Delta P, V)$: temperature deviation ($^{\circ}\text{C}$), pressure deviation (bar), and vibration level (mm/s):

Unit	ΔT	ΔP	V
U1	1.0	0.5	0.8
U2	1.2	0.6	0.9
U3	0.8	0.4	0.7
U4	6.0	3.0	4.5
U5	6.5	3.2	4.8
U6	5.8	2.9	4.2

Initialising two centroids at $U1 = (1.0, 0.5, 0.8)$ and $U4 = (6.0, 3.0, 4.5)$, every point's Euclidean distance to each centroid already separates the data perfectly on the first assignment step (U1-U3 are close to centroid 1; U4-U6 are close to centroid 2), so *k*-means converges in a single iteration, to

$$\mathbf{c}_1 = (1.00, 0.50, 0.80), \quad \mathbf{c}_2 = (6.10, 3.03, 4.50).$$

The two clusters correspond to two distinct **operating regimes** — a tight, low-stress cluster and a visibly-elevated, higher-stress cluster — discovered with no fault labels anywhere in the process.




Figure 9.1: The six equipment units in their full 3-dimensional sensor space, coloured by the cluster k -means assigned them to. The two navy crosses mark the converged centroids. Real ILM sensor embeddings run to hundreds of dimensions; three is the smallest number that already shows the geometry is a genuine spatial clustering problem, not merely a 1D thresholding exercise.

9.3.2 Dimensionality Reduction

Dimensionality reduction compresses high-dimensional data into fewer dimensions while preserving as much structure as possible — exactly Chapter 1’s rate-distortion theory and Chapter 3’s curse-of-dimensionality worked example, now applied as a preprocessing step. **Principal Component Analysis** (PCA) finds the directions of greatest variance in the data (the eigenvectors of the covariance matrix from Chapter 2, Section 2.5) and projects onto the top few; a trained autoencoder’s bottleneck (Chapter 8, Section 8.6) is a non-linear generalisation of exactly the same idea.

9.3.3 Anomaly Detection

Anomaly detection identifies data points that do not fit the patterns the rest of the data exhibits, using exactly the tools Section 9.3’s other two tasks already introduced: a point far from every cluster centroid, or a point an autoencoder reconstructs badly (Chapter 8), is a natural anomaly candidate. Chapter 16 is devoted entirely to doing this well.

→ ILM

Unsupervised learning discovers structure in the ILM’s raw, *unlabeled* industrial corpus before any labels exist at all — clustering similar maintenance reports, reducing sensor telemetry to a compact embedding for retrieval, and flagging outlier documents or readings. This is precisely what happens *underneath* Chapter 11’s Retrieval-Augmented Generation: the documents RAG retrieves are found via nearest-neighbour search in exactly the kind of embedding space k -means just clustered.

9.4 Reinforcement Learning

9.4.1 Markov Decision Processes

Reinforcement Learning (RL) trains an agent to choose **actions** that maximise cumulative **reward** through trial and error, with no labeled “correct action” ever provided directly — only the consequences, as reward, of whatever action was actually tried. The formal setting is a **Markov Decision Process** (MDP), a direct extension of Chapter 2’s Markov chain (Section 2.8) with actions and rewards added:

Markov Decision Process

An MDP is a tuple $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$: a set of states $\mathcal{S}$, a set of actions $\mathcal{A}$, transition probabilities $P(s' | s, a)$, a reward function $R(s, a)$, and a discount factor $\gamma \in [0, 1)$ that weighs future reward against immediate reward.

9.4.2 Q-Learning and Policy Gradients

Q-learning learns the **action-value function** $Q(s, a)$: the expected total discounted future reward of taking action a in state s and acting optimally thereafter. It satisfies the **Bellman equation**,

$$Q(s, a) = R(s, a) + \gamma \max_{a'} Q(s', a'),$$

which is solved iteratively, exactly the way Chapter 5’s perceptron was solved iteratively: start with a guess (often $Q = 0$ everywhere) and repeatedly apply the equation as an update rule until it converges.

Q-Learning for a Simplified Pump Controller

A simplified pump controller has three states, Idle, Normal, Stressed, and two actions, Continue and ReduceLoad, with deterministic transitions and the following immediate rewards (positive for productive operation, sharply negative for continuing under stress):

State	Action	Next state	Reward
Idle	Continue	Normal	+1
Idle	ReduceLoad	Idle	0
Normal	Continue	Stressed	+2
Normal	ReduceLoad	Normal	+1
Stressed	Continue	Stressed	−5
Stressed	ReduceLoad	Normal	+1

With $\gamma = 0.9$ and Q initialised to zero everywhere, applying the Bellman update synchronously across all six (s, a) pairs converges (to six decimal places) after 45 iterations:

State	$Q(\cdot, \text{Continue})$	$Q(\cdot, \text{ReduceLoad})$	Learned action
Idle	14.706	13.234	Continue
Normal	15.234	14.706	Continue
Stressed	8.234	14.706	ReduceLoad

The learned policy is exactly what an experienced operator would choose: keep running under Idle and Normal conditions, but back off under Stress — and

this policy was discovered entirely from the reward numbers, with no rule ever stating “reduce load when stressed” directly. This is RL’s central strength and, as Section 9.4.3 makes concrete, its central danger in the same breath: the policy is only as good as the reward function that produced it, and small changes in that handful of reward numbers can quietly change the learned policy’s behaviour.

Policy gradient methods take a different route to the same goal: rather than learning $Q(s, a)$ and acting greedily with respect to it, they directly parameterise a policy $\pi_\theta(a | s)$ (often, exactly a softmax over a neural network’s output, Chapter 4) and adjust θ by gradient ascent (Chapter 4, Chapter 7) on expected reward. The REINFORCE estimator,

$$\nabla_\theta J(\theta) = \mathbb{E}[\nabla_\theta \log \pi_\theta(a | s) \cdot G],$$

where G is the observed cumulative reward following action a , is backpropagation (Chapter 7) applied to a policy network, with the observed reward standing in for a label that, in RL, simply does not exist.

9.4.3 Applications in Control: When RL Is Appropriate, and When It Is Dangerous

RL is a natural fit whenever the “correct” action is genuinely unknown in advance and can only be discovered through consequences — tuning a control policy for a novel process, optimising a complex multi-step scheduling problem. It is dangerous, and this book takes a firm position on this point, in exactly the settings this whole series is about: **RL trained (or worse, still learning) directly against a real, physical, critical process is learning by trial and error against equipment that can be damaged and people who can be hurt by the “error” half of that phrase.** Every credible industrial deployment of RL for critical control either trains exclusively in high-fidelity simulation before any real deployment, or wraps the learned policy’s output in hard, non-learned safety limits it cannot override — a direct preview of the **Deterministic Islands** of Chapter 12: an RL policy may *recommend* an action, but the boundary of what actions are physically permitted at all is exactly the kind of static, auditable, non-stochastic constraint a Deterministic Island enforces, sitting downstream of the learned policy and entirely outside its ability to learn its way around.

9.5 Semi-supervised and Self-supervised Learning

Labels are expensive; raw industrial data is comparatively cheap. **Semi-supervised learning** trains on a small labeled set together with a much larger unlabeled one, letting the unlabeled data’s structure (Section 9.3) inform decisions the small labeled set alone could not support reliably. **Self-supervised learning** goes further and manufactures its own labels directly from unlabeled data’s structure: hide part of the input and train the model to predict it. This is not a minor technique; it is, as Chapter 11 shows in full, *the* technique that pre-trains a large language model at all.

→ ILM

Self-supervised learning is exactly how the ILM’s backbone acquires its initial industrial language competence, entirely before any supervised fine-tuning: masking words in a maintenance manual and training the model to predict them (Masked Language Modeling), or training it to predict the next token in a fault report (Causal Language Modeling) — both examined formally in Chapter 11, Section 11.3, and both requiring not one single hand-written label.

9.6 Transfer Learning and Fine-Tuning

Transfer learning takes a model already trained on one task (often a general-purpose LLM, trained on a vast, generic text corpus) and adapts it to a new, related task, rather than training from scratch. **Fine-tuning** is the most direct form of transfer: continue training the pre-trained model’s weights (or, as Chapter 11’s Parameter-Efficient Fine-Tuning shows, only a small additional set of weights) on a smaller, task-specific labeled dataset.

→ ILM

This is the literal definition of how this book proposes the ILM itself should be built: not trained from nothing, but transferred from a general-purpose pre-trained language model and fine-tuned on industrial text — the entire subject of Chapter 11 — rather than reinvented from a blank slate. Chapter 10, Section 10.4.2’s “domain adaptation” is exactly this idea, named at the point the ILM’s architecture is first laid out.

9.7 Applications in Critical Infrastructure


Paradigm	Typical application	Caution
Supervised	fault classification, RUL regression	needs representative labeled failure data
Unsupervised	operating-regime clustering, embedding for retrieval	clusters found are not automatically meaningful
Reinforcement	control policy optimisation, scheduling	never learn directly against a live critical process
Self-supervised	pre-training the ILM backbone on raw industrial text	needs large, reasonably clean text corpora
Transfer/fine-tuning	adapting a general LLM to a specific plant	the base model’s biases transfer along with its competence

Honest Boundary

No learning paradigm is a silver bullet. Every paradigm in this chapter solves a different problem shape, and choosing among them is a direct function of what data is available (labeled? unlabeled? only a reward signal?), what the task actually requires (a decision on this input, a summary of this document, a control policy over time?), and, especially in the setting this book is written for,

how much risk a wrong answer carries. The ILM described starting next chapter combines several of these paradigms deliberately — self-supervised pre-training, supervised fine-tuning, and, wrapped carefully in the deterministic guardrails of Chapter 12, a bounded role for reinforcement learning — precisely because no single paradigm in this chapter, alone, is adequate to the whole problem.

THE INDUSTRY LANGUAGE MODEL (ILM)



THE INDUSTRY LANGUAGE MODEL (ILM)

10. Why an ILM?

10.1 What You Already Know

If you have used a general-purpose large language model — asked it to draft an email, summarise a document, or explain a concept — you already have direct, personal experience of both its strengths and its failure modes. You have likely also encountered the idea of a **domain-specific model**: a version of a general model narrowed toward one field (a legal-document model, a medical-coding model), trained or adapted so that its competence concentrates where a specific user actually needs it, at some cost to its competence everywhere else. Every argument in this chapter is an attempt to answer one question precisely, for the domain this entire book is about: exactly *which* of a general LLM's failure modes are merely inconvenient elsewhere, and which are disqualifying the moment the domain is a nuclear plant, a chemical reactor, or a power grid.

10.2 The Limits of General LLMs in Critical Applications

10.2.1 Hallucinations: The Model "Confabulates" Facts

A large language model, at its mathematical core, is exactly Chapter 8's Transformer: a system trained to predict a plausible *continuation* of text. Nothing in that training objective distinguishes a plausible-sounding continuation from a *true* one; the two coincide often enough, on general knowledge, to be genuinely useful, and diverge often enough, on specific or obscure facts, to be genuinely dangerous. A general LLM asked for the maximum permitted operating pressure of a specific pump model it has never seen documentation for will not, in general, refuse to answer; it will produce a fluent, confident, plausible-sounding number — a **hallucination**, or, more precisely, a **confabulation**: the model is not lying, because lying requires knowing the truth and saying otherwise; it is generating the statistically likely continuation of the prompt, with no internal distinction at all between "recalled from training data" and "the shape a plausible answer to this kind of question usually takes."

10.2.2 Lack of Domain Knowledge

Even a hallucination-free general LLM simply does not know *this plant's* specific procedures, *this equipment model's* exact maintenance history, or the content of a maintenance manual that was never part of its training corpus at all (most industrial SOPs, vendor manuals, and internal fault logs are not public internet text, and could not have been in any general model's training data). No amount of general

competence substitutes for access to the specific document that actually answers the question.

10.2.3 Non-Deterministic Behaviour

A general LLM generates text by repeatedly sampling from a probability distribution over the next token — exactly Chapter 8’s softmax, Chapter 1’s probability distribution, now used to *generate* rather than merely to classify. This sampling is controlled by a **temperature** parameter T that rescales the logits before the softmax: $P(\text{token}) = \text{softmax}(z/T)$. A low temperature concentrates probability on the single most likely token (nearly deterministic); a higher temperature flattens the distribution, increasing variety at the direct cost of consistency.

The Same Question, Twice, at Different Temperatures

A model asked “what is the maximum permitted operating pressure?” assigns these logits to four candidate numeric completions, in bar:

Candidate	120 (correct)	125	110	150
Logit	4.2	3.9	3.5	2.0

Applying temperature-scaled softmax at four settings:

T	$P(120)$	$P(125)$	$P(110)$	$P(150)$
0.1	0.952	0.047	0.001	0.000
0.7	0.485	0.316	0.178	0.021
1.0	0.426	0.316	0.212	0.047
1.5	0.374	0.306	0.234	0.086

At the low temperature typically used for “factual” queries in consumer chat products, the correct answer is heavily favoured but still not guaranteed (95.2%, not 100%). At the moderate temperatures ($T = 0.7$ to 1.0) common in general-purpose deployment for more natural-sounding text, the correct value 120 is barely ahead of the plausible-but-wrong 125, and the model will, over repeated queries, actually emit each of the four candidates some non-trivial fraction of the time. **Two identical questions, asked twice, can produce two different numbers** — not because anything went wrong, but because sampling from a probability distribution is exactly what a language model, used this way, is designed to do. For a safety-relevant numeric fact, this is not a minor quirk to be tuned away with a lower temperature; it is a structural property of the generation mechanism itself, and Section 10.3.3 states plainly why it cannot be accepted for facts of this kind.

10.2.4 Inability to Verify

A general LLM’s output arrives with no attached, trustworthy measure of its own correctness. The model’s own stated confidence (if asked) is itself just another generated, sampled piece of text, with no guaranteed relationship to the true probability the answer is correct — a model can be, and frequently is, fluently and confidently wrong. Chapter 13 examines uncertainty quantification in depth; the point here is narrower and more basic: nothing in a general LLM’s ordinary operation tells its user which category a given answer falls into.

10.2.5 Safety: One Mistake Can Be Catastrophic

Every failure mode above is a nuisance in a general consumer application (a wrong trivia fact, an invented citation) and a potential trigger for physical harm in the domain this book is written for. A hallucinated pressure limit, confidently delivered, is not merely wrong text; if acted on, it is a wrong operating parameter on real equipment. The asymmetry between these two settings, not any change in the underlying technology, is the entire reason this chapter, and this book, exist.

10.3 Definition of the ILM (Industry Language Model)

10.3.1 Data

An ILM is trained and grounded on data a general LLM typically never sees: equipment manuals, standard operating procedures, fault and maintenance records, operator logs, sensor telemetry, and internal technical reports — the specific, proprietary, often non-public documents that actually describe a given plant or fleet.

10.3.2 Goal

The ILM's purpose spans three related capabilities: **understanding** industrial language (parsing a fault report, interpreting a maintenance note written in an engineer's shorthand), **generation** (drafting a report, summarising a shift's events), and **decision support** (surfacing the relevant procedure, flagging an anomaly, recommending — never silently executing — a course of action).

10.3.3 Requirements: Reliability, Verifiability, Determinism Where Needed

The ILM's Three Requirements

Reliability: performance that holds up consistently across the actual distribution of industrial queries, not only on a curated benchmark. **Verifiability:** every output must be traceable to a source or a computation that can be checked, not merely trusted. **Determinism where needed:** *not everywhere* — open-ended reasoning and drafting can and should remain the flexible, probabilistic process Chapter 8's Transformer naturally is; but a narrow, identifiable class of queries (stated safety limits, regulatory thresholds, calibration constants) must be answered *identically every time*, with no dependence on temperature, sampling, or which server happened to handle the request.

This last requirement is the thread the rest of this book, especially Chapter 12, exists to make precise: the ILM is not proposed as a uniformly deterministic system (that would sacrifice everything generative language models are good at) nor as a uniformly stochastic one (that would repeat Section 10.2.3's failure on every safety-relevant fact). It is proposed as a system that *knows which of its own outputs need to be which*, and routes accordingly.

10.4 Architecture of the ILM

10.4.1 Transformer-Based

The ILM's backbone is the Transformer architecture of Chapter 8, Section 8.5, unmodified in its core mechanism: multi-head self-attention, positional encoding, an encoder-decoder or decoder-only structure depending on the specific task. Nothing about being an *industry* language model changes the underlying neural architecture; what changes is everything wrapped around it.

10.4.2 Adaptations

Four adaptations distinguish an ILM from a general-purpose LLM pointed at industrial text.

Time-series adaptation. Plant telemetry is not natural language, but the Transformer's attention mechanism (Chapter 8, Section 8.4) generalises to it directly: a **temporal embedding** encodes each timestep's sensor reading the way a token embedding encodes a word, and **multi-scale attention** lets the model attend over multiple time resolutions at once (recent fine-grained readings and longer-term trends) in a single architecture.

Multivariate adaptation. A real plant reports many **heterogeneous channels** simultaneously (temperature, pressure, vibration, current, each with its own units, sampling rate, and noise characteristics, exactly Chapter 2's multivariate distributions made architectural); the ILM's input layer must embed each channel appropriately before attention can relate them to each other and to text.

Domain adaptation. Pre-training on general text gives a model broad language competence; continued pre-training on industrial text (Chapter 11, Section 11.2) and **Retrieval-Augmented Generation** (RAG, Chapter 11, Section 11.6) ground its answers in the plant's own documents at query time, directly addressing Section 10.2.2's lack-of-domain-knowledge failure.

Deterministic anchors. This is the direct architectural response to Section 10.2.3's non-determinism failure, and this book's central engineering contribution: certain facts, computations, and rules are not generated by the stochastic Transformer at all, but resolved through a separate, deterministic mechanism the Transformer consults, the way a calculator consults a fixed arithmetic rule rather than guessing at a plausible-sounding sum. Chapter 12 is devoted entirely to this mechanism; this section introduces its central data structure for the first time.

A First Look at the ILMVector

Every embedding that flows through the ILM — whether it represents a token, a retrieved document chunk, or a sensor reading — is, in ordinary Transformer terms, just a vector of numbers, and the overwhelming majority of them, at any given moment, are exactly that and nothing more. Only a small subset ever participates in a Deterministic Island (Chapter 12) at all. The ILM's vector type reflects this directly: identity and content are universal, present on every vector; determinism bookkeeping is a separate, optional block, present only for the vectors that need it.

ILMVector.cs — a first, preliminary sketch

```

1 public sealed record ILMVector
2 {
3     // This vector's own identity -- what another vector's
4     // PointerTo (below) actually refers to.
5     public required Guid VectorId { get; init; }
6
7     // The ordinary Transformer embedding (Chapter 8). Every
8     // vector in the ILM has one; nothing else on this record
9     // is guaranteed to be present.
10    public required float[] Embedding { get; init; }
11
12    // Present only if this vector has ever taken part in a
13    // Deterministic Island (Chapter 12). Null for the ordinary
14    // token, sensor, and document embeddings that make up most
15    // of the ILM's traffic -- they were never asked to be
16    // deterministic, so they carry no opinion on the subject.
17    public DeterminismMetadata? Determinism { get; init; }
18 }
19
20 public sealed record DeterminismMetadata
21 {
22     // Is this vector's content currently resolved and authoritative
23     // (Chapter 12's Static Vault, Frozen Snapshot\index{Frozen
24     // Snapshot}, ...), or is it
25     // only a candidate awaiting resolution (see PointerTo below)?
26     public required bool IsDeterministic { get; init; }
27
28     // Versioning: when is this vector's content valid?
29     public required string Version { get; init; }
30     public required DateTime ValidFrom { get; init; }
31     public DateTime? ValidTo { get; init; }
32
33     // If set, this vector defers to another vector's VectorId --
34     // e.g. a parameter that "points to" the current governing
35     // standard. Chapter 12 fixes the routing rule this pointer
36     // obeys.
37     public Guid? PointerTo { get; init; }
38 }

```

Two records, not one, and the split is the point. `ILMVector` carries only what *every* vector in the system has: *identity* (`VectorId` — without it, no other vector could ever name this one as the target of a `PointerTo`) and *content* (`Embedding`, the only field an ordinary Transformer would ever have). Naming this type “DeterministicVector” was an earlier, tempting mistake this book corrects here on sight: it would have asserted, in the type’s very name, a property (determinism) that is false for almost every instance of it. `DeterminismMetadata` — nullable, attached only when relevant — carries the four ideas that actually concern determinism: *whether this vector is currently resolved and authoritative* (`IsDeterministic`), *when it is valid* (`Version`, `ValidFrom`, `ValidTo` — a stated safety limit that changed last year must not be silently used for a query about last year’s operation), and *whether it defers to another vector* (`PointerTo`, matched against some other vector’s `VectorId`). This last field matters most:

if vector *A*'s metadata points to vector *B*'s VectorId, and *B* is deterministic, the correct behaviour is to resolve *A*'s query through *B*'s deterministic content, not through whatever stochastic value *A* might otherwise have generated on its own. This routing behaviour, and the full schema it requires (including how conflicting pointers and expired versions are resolved), is developed completely in Chapter 12; this sketch is intentionally incomplete, introduced here only so that Chapter 11's training discussion and Chapter 12's full treatment have a shared object to refer back to.

10.5 Training Data

10.5.1 Sources

Manuals and vendor documentation, standard operating procedures, historical maintenance records, fault and incident reports, operator shift logs, and raw sensor telemetry — the same list as Section 10.3.1, now considered specifically as *training material* rather than as the knowledge the model is meant to draw on at query time (the distinction matters: some of this data trains the model's weights directly, Chapter 11's pre-training and fine-tuning; some of it is better left out of training entirely and instead retrieved at query time via RAG, so that updating a procedure means editing a document, not retraining a model).

10.5.2 Preprocessing

Raw industrial text needs the same tokenization Chapter 8's Transformer requires for any text, with domain-specific care: industrial **tokenization** must handle unit-laden numeric expressions ("120 bar", "45°C"), equipment codes, and abbreviations that a general-purpose tokenizer, trained on generic web text, will often fragment badly. **Normalization** maps equivalent expressions (different date formats, unit systems, or equipment-naming conventions across sites) onto a consistent internal representation. **Data augmentation**, used cautiously, can expand a small labeled fault-category dataset by paraphrasing or back-translation, subject to Chapter 3's warning that augmented data is not a substitute for genuinely representative new examples.

10.5.3 Handling Sparse Data and Class Imbalance

Industrial fault data is, almost by definition, **imbalanced**: catastrophic failures are (fortunately) rare, meaning "failure" examples are vastly outnumbered by "normal operation" examples in any real historical dataset — the same base-rate problem Chapter 2, Section 2.3's Bayes'-theorem worked example already warned about, now facing the ILM's training pipeline directly rather than a single sensor's alarm logic. Chapter 16 addresses this in full for anomaly detection specifically; here it is enough to flag that naively training on an imbalanced industrial dataset without correction (reweighting, resampling, or a loss function designed for imbalance) reliably produces a model that has learned to predict "normal" and be right most of the time, while being useless for the rare cases that actually matter.

10.5.4 Data Ethics and Security

Industrial training data frequently contains information a plant has strong reasons to protect: proprietary process parameters, safety incident details that carry legal sensitivity, and sometimes personal data about the operators involved in a logged incident. Training an ILM on this data inherits every obligation the plant already has around that data's handling, storage, and access control, and adds a new one specific to language models: a sufficiently capable model can sometimes be induced to reproduce specific training examples verbatim, meaning sensitive training data is not automatically safe merely because it was "only" used for training rather than published directly.

Honest Boundary

The ILM is not a panacea; it is a specialised tool. Every adaptation in Section 10.4 narrows the model toward industrial competence at some cost to its general-purpose fluency, and no amount of architectural cleverness substitutes for what Section 10.5 already implied: **data quality determines model quality**, completely and without exception. An ILM trained on incomplete, outdated, or poorly curated plant documentation will confidently reproduce that documentation's gaps and errors, not correct for them — Section 10.2.1's hallucination risk, in a different, equally dangerous disguise: not inventing facts from nothing, but faithfully learning wrong facts that were already sitting in the training data.

11. Training and Adaptation of the ILM

11.1 What You Already Know

You already have the two-stage mental model this chapter formalises, even if you have only encountered it as a user: a large language model is first **pre-trained**, once, at enormous cost, on a vast and general corpus, acquiring broad linguistic competence with no specific task in mind; it is then **fine-tuned**, far more cheaply, to specialise that general competence toward a specific use. This chapter is the complete, technical version of that two-stage story, applied specifically to building the ILM introduced in Chapter 10.

11.2 Pre-training

Pre-training is Chapter 9, Section 9.5’s self-supervised learning, at scale: the ILM’s Transformer backbone (Chapter 8) is trained on a very large corpus of industrial text — manuals, procedures, historical fault reports, technical documentation — with labels manufactured automatically from the text itself, exactly as Section 9.5 previewed. No human ever hand-labels this corpus; the supervision signal is the text’s own structure.

11.3 Masked Language Modeling and Causal Language Modeling

Two self-supervised objectives dominate in practice, and they produce models with different strengths.

MLM vs. CLM

Masked Language Modeling (MLM, the BERT family): randomly mask a fraction of input tokens and train the model to predict the masked tokens from the *surrounding context on both sides*. Suited to *understanding* tasks (classification, extraction), because the model learns a representation informed by full bidirectional context. **Causal Language Modeling** (CLM, the GPT family): train the model to predict each token from only the tokens *before* it, left to right. Suited to *generation* tasks (drafting, summarising), because this is exactly how text is produced one token at a time at inference.

A Single Masked-Token Prediction, in Full

Training on the sentence "Pump P-101 tripped due to high [MASK] pressure," with the true masked word "discharge," the model's output layer (Chapter 4's softmax, applied over a candidate vocabulary) assigns these logits to four plausible completions:

Candidate	discharge (true)	suction	differential	inlet
Logit	5.1	4.3	3.8	2.9

Softmax converts these to probabilities: $P(\text{discharge}) = 0.546$, $P(\text{suction}) = 0.245$, $P(\text{differential}) = 0.149$, $P(\text{inlet}) = 0.061$. The cross-entropy training loss (Chapter 4, Section 4.3.2) for this one masked position is $-\log(0.546) \approx 0.606$ nats. Backpropagation (Chapter 7) uses this single number to nudge every weight that contributed to it, across potentially billions of parameters, very slightly toward assigning more probability to "discharge" the next time a similar context appears. Multiplied across millions of masked positions in the corpus, this one mechanism is the entirety of how the model acquires industrial vocabulary and typical phrasing, with no engineer ever writing a rule that "discharge pressure" is a common collocation.

11.4 Fine-Tuning

Fine-tuning continues training a pre-trained model's weights on a smaller, labeled, task-specific dataset — Chapter 9's supervised learning, now starting from a competent initialisation rather than from random weights. Three ILM tasks recur: **classification** (sort an incoming report into one of a known set of fault categories, exactly Chapter 9's "→ ILM" box), **question answering** (extract or generate a direct answer to an operator's question from the relevant procedure), and **summarisation** (condense a long shift log or incident report into a short brief). Each is fine-tuned with the appropriate loss from earlier chapters: cross-entropy over fault categories for classification, and, for QA and summarisation, the same CLM next-token cross-entropy Section 11.3 already introduced, now applied to generating the answer or summary text itself.

11.5 Parameter-Efficient Fine-Tuning: LoRA and Adapters

Fine-tuning every one of a large model's weights is expensive in memory and storage, and risks **catastrophic forgetting** — degrading the general competence pre-training built, in the course of specialising it. **Parameter-Efficient Fine-Tuning** (PEFT) methods freeze the pre-trained weights entirely and train only a small number of additional parameters.

LoRA — Low-Rank Adaptation

Rather than fine-tuning a weight matrix $W \in \mathbb{R}^{d \times d}$ directly, LoRA freezes W and learns a low-rank update

$$W' = W + BA, \quad B \in \mathbb{R}^{d \times r}, \quad A \in \mathbb{R}^{r \times d}, \quad r \ll d.$$

Only B and A are trained; W never changes.

How Much Does Low-Rank Adaptation Actually Save?

A single weight matrix in a mid-sized Transformer layer might have $d = 4,096$, so W itself has $4,096^2 = 16,777,216$ parameters. Fine-tuning W directly means updating (and storing a separate copy of) all 16.8 million of them. LoRA with rank $r = 8$ instead trains B ($4,096 \times 8$) and A ($8 \times 4,096$):

$$|B| + |A| = 4,096 \times 8 + 8 \times 4,096 = 65,536 \text{ parameters,}$$

a **256-fold reduction**, at 0.39% of the original parameter count, for this one matrix. **Adapters** achieve a similar saving by a different route: small bottleneck layers (down-project, non-linearity, up-project) inserted between the frozen Transformer's existing layers, with only the adapter's own weights trained. Both techniques mean a single pre-trained ILM backbone can support many different task-specific or even site-specific fine-tunes, each stored as a lightweight add-on measured in megabytes rather than a full duplicate copy of a multi-gigabyte model.

11.6 Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) addresses Chapter 10, Section 10.2.2's lack-of-domain-knowledge failure directly, without retraining anything: at query time, the system first **retrieves** the most relevant document chunks from the plant's own corpus (via nearest-neighbour search in an embedding space, Chapter 9, Section 9.3.1's clustering geometry, run in reverse), then **generates** its answer with those chunks inserted directly into the model's context.




Figure 11.1: Retrieval-Augmented Generation. (The generator also receives the original query alongside the retrieved chunks; omitted from the diagram for clarity.) The document store holds ILMVectors (Chapter 10, Section 10.4.2): most carry no determinism metadata at all, but a chunk that corresponds to a formally verified safety limit carries a non-null Determinism block with `IsDeterministic=true`, and the generator’s role for that chunk narrows to *reporting* the retrieved value faithfully, not generating a fresh number for it — RAG, combined with the ILMVector’s determinism metadata, is the mechanism by which Section 10.2.3’s non-determinism failure is addressed for exactly the class of query where it matters most.

11.7 Prompt Engineering for Industrial Contexts

How a query is phrased measurably changes an LLM’s output, even with identical underlying knowledge available to it. A vague prompt (“what’s the pressure limit?”) leaves the model to guess which equipment, which limit, and which units are meant; a well-constructed industrial prompt states the equipment model, the specific parameter, and the expected answer format explicitly (“For pump model P-101, state the maximum permitted discharge pressure in bar, citing the source document”), which both improves accuracy and, critically, makes the answer’s provenance checkable. **Few-shot prompting** (including two or three worked examples of the desired question-and-answer format directly in the prompt) further steers the model’s output format without any weight update at all — a zero-training-cost complement to Section 11.4’s fine-tuning, useful exactly when a task is too narrow or too rare to justify a dedicated fine-tune.

11.8 Evaluation: Perplexity, BLEU, ROUGE, and Human Evaluation

Perplexity

$$\text{Perplexity} = \exp(\bar{L})$$

where $\bar{L}$ is the average per-token cross-entropy loss (Section 11.3) over a held-out set, in nats.

Perplexity is Chapter 1’s entropy, re-expressed: it measures, on average, how many equally-likely options the model was effectively choosing among at each position — a perplexity of 1 means perfect, certain prediction; a perplexity of 10 means the model’s uncertainty at each token was, on average, comparable to choosing uniformly among 10 options.

From Loss to Perplexity

A model achieving an average cross-entropy loss of $\bar{L} = 0.083$ nats per token on held-out industrial text has perplexity $e^{0.083} \approx 1.087$ — very low, indicating the industrial vocabulary and phrasing is highly predictable to this model (likely reflecting a narrow, repetitive domain corpus, such as standardised procedure text). A general-purpose model evaluated on the same narrow corpus, without domain adaptation, might show $\bar{L} = 2.3$ nats, giving perplexity $e^{2.3} \approx 9.97$ — a roughly ninefold difference in effective uncertainty per token, and a direct, quantitative measurement of exactly the domain-adaptation benefit Chapter 10, Section 10.4.2 argued for qualitatively.

BLEU and **ROUGE** evaluate generated text (a summary, a translated procedure) against one or more human-written references by measuring n -gram overlap: BLEU emphasises precision (how much of what the model generated appears in the reference), ROUGE emphasises recall (how much of the reference appears in what the model generated) — complementary, and both are surface-level text-overlap metrics that can be fooled by a fluent paraphrase that says the same thing in different words, or, more dangerously, by text that overlaps heavily in wording while being factually wrong in a single critical number. **Human evaluation** — domain experts directly reviewing model outputs against defined criteria — remains, despite its cost, the only method that reliably catches this specific, high-consequence failure mode, and no automatic metric in this section is proposed as a substitute for it in a safety-relevant deployment.

11.9 Ethical Issues: Bias, Fairness, and Transparency

A model pre-trained partly on general text and fine-tuned on a specific plant's history inherits both layers' biases: general societal biases present in web-scale pre-training data, and site-specific biases present in whatever historical maintenance and incident data it was fine-tuned on (if certain shifts, equipment, or sites were historically under-reported or over-scrutinised, the model learns that pattern as if it were a property of the equipment itself, not an artifact of reporting history). **Transparency** — being able to state which data a given answer draws on, and being honest with operators about the model's known limitations — is not a nicety layered on afterward; it is a precondition for the "verifiability" requirement Chapter 10, Section 10.3.3 already established as non-negotiable. Chapter 14 returns to all three of these issues in full.

Honest Boundary

Fine-tuning is not a guarantee of correctness. A fine-tuned model that performs well on its validation set has learned to fit that validation set's particular distribution of questions; nothing guarantees it generalises correctly to the full range of questions it will actually be asked in deployment, and Chapter 3's bias-variance tradeoff applies to a fine-tuned ILM exactly as it does to any other trained model.

Evaluation is hard, and often subjective. Perplexity, BLEU, and ROUGE are cheap to compute and easy to report, which is exactly why they are so often over-relied upon; none of them directly measures the property that actually matters for this book's purpose — whether a specific, safety-relevant answer

was correct. Rigorous evaluation of an ILM intended for critical infrastructure requires the harder, more expensive combination of automatic metrics, structured human expert review, and, for the specific class of queries Chapter 10, Section 10.3.3 flagged as needing determinism, Chapter 13's formal verification — not any single number this chapter has introduced, taken alone.

12. Deterministic Islands

12.1 What You Already Know

You already know what **reproducibility** means in an engineering context: run the same test under the same conditions and get the same result, every time, or the result cannot be trusted as a measurement at all. You know the difference between a **deterministic** system (a calculator: the same inputs always produce the same output) and a **stochastic** one (a dice roll, or, as Chapter 4 through Chapter 8 established in full mathematical detail, a neural network). And you know, from the very first pages of this book, that critical infrastructure imposes a safety requirement ordinary software does not always have to meet: for a defined, safety-relevant class of questions, “usually correct” is not an acceptable standard. This chapter is where these three familiar ideas are combined into the specific engineering mechanism this entire book has been building toward.

12.2 The Problem

12.2.1 Neural Networks Are Stochastic

Chapter 4’s random weight initialisation, Chapter 7’s dropout, and Chapter 10’s temperature-scaled sampling (Section 10.2.3) are not three unrelated technical details; they are three separate sources of the same underlying fact: **every stage of a neural network’s life, from training to inference, involves a source of randomness**, deliberately introduced because Chapter 7 showed it helps training (Section 7.5.2’s escape from saddle points) and Chapter 10 showed it helps generation (natural-sounding, varied text). The randomness is not a bug anywhere else in this book. Here, for a specific class of query, it is exactly the problem.

12.2.2 Two Identical Inputs, Two Different Outputs

Chapter 10, Section 10.2.3 already made this concrete and quantitative: the same question, asked twice, at a realistic deployment temperature, can return two different numeric answers, purely as a consequence of sampling from a probability distribution correctly and as designed. Nothing has malfunctioned. This is, simply, what the system in Chapters 4 through 8 does.

12.2.3 In Critical Applications, This Is Unacceptable

An operator asking twice for the maximum permitted operating pressure of a piece of equipment, on two different shifts, and receiving two different numbers, has not encountered a minor inconsistency. They have encountered a system that cannot be

trusted to state a fact reliably — and Chapter 10, Section 10.2.5 already established what a wrong operating parameter, acted on, can cost. **Deterministic Islands** are this book's answer: a small, carefully bounded set of mechanisms that guarantee identical output for identical input, for exactly the queries that need it, without requiring the entire ILM to sacrifice the generative flexibility that makes it useful everywhere else.

12.3 The Deterministic Islands

Three distinct mechanisms, increasing in flexibility and decreasing in the strength of guarantee they offer, make up the Deterministic Islands.

12.3.1 Static Vault: Hash-Based Routing

The strongest guarantee is the simplest mechanism: for a known, finite set of safety-critical facts, bypass the neural network entirely. A **Static Vault** maps a normalised query directly to a stored, pre-verified answer via a cryptographic hash, exactly the way a hash table looks up any other key.

The Static Vault Guarantee

`answer = Vault[SHA-256(normalised query)]`

Because SHA-256 is deterministic (the same input byte sequence always produces the same 256-bit output) and the vault is an ordinary lookup table, identical input guarantees identical output **absolutely**, not "with high probability": nothing about this path touches the stochastic Transformer at all.

StaticVault.cs

```

1 using System.Security.Cryptography;
2 using System.Text;
3
4 public sealed class StaticVault
5 {
6     private readonly IReadOnlyDictionary<string, ILMVector> _table;
7
8     public StaticVault(IReadOnlyDictionary<string, ILMVector> table)
9         => _table = table;
10
11     // Deterministic, one-way key derivation from a normalised query.
12     public static string HashQuery(string normalizedQuery)
13     {
14         var bytes = SHA256.HashData(Encoding.UTF8.GetBytes(
15             normalizedQuery));
16         return Convert.ToHexString(bytes).ToLowerInvariant();
17     }
18
19     public ILMVector? Lookup(string normalizedQuery)
20     {
21         var key = HashQuery(normalizedQuery);
22         return _table.TryGetValue(key, out var vector) ? vector : null;
23     }
24 }
```

Hashing a Real Query

The normalised query "what is the maximum allowed discharge pressure for pump p-101?" (lower-cased, whitespace-trimmed — normalisation must be fixed and documented, since even a single differing character produces a completely unrelated hash) hashes, via SHA-256, to

```
f20ab319ace115290fd196a1c1d5b574e60ed0b39692add217c4ff0bcf03e574.
```

This exact 64-character hex string is the vault's lookup key. Any future query that normalises to the identical string — regardless of server, timestamp, or model version — produces this identical key, and therefore this identical lookup. The corresponding ILMVector carries a Determinism block with `IsDeterministic = true`, and its Embedding encodes the verified answer ("120 bar"), not a value the Transformer ever generated. Every other field on this particular vector's Determinism block is populated too (`Version`, `ValidFrom`); it is precisely because this vector *does* participate in a Deterministic Island that none of `DeterminismMetadata`'s fields are left unset here — contrast this with the ordinary token and sensor embeddings discussed in Chapter 10, whose `Determinism` property is simply `null`.

12.3.2 Frozen Snapshot: Locked Software, Seed, and CPU Inference

The Static Vault handles facts that can be enumerated in advance. Some deterministic behaviour instead needs to come from the model *itself*, run in a mode that produces bitwise-identical output every time. This requires locking down every remaining source of non-determinism in the inference pipeline, not just the sampling temperature:

Sources of Non-Determinism a Frozen Snapshot Locks

Software version: different releases of the model weights, or even the inference library, can produce different numerical results from the same nominal computation. **Random seed:** any remaining sampling step must be seeded identically, every time. **CPU, not GPU, inference:** floating-point addition is not associative ($(a + b) + c$ need not exactly equal $a + (b + c)$ in finite precision); GPU inference parallelises operations like the sum in Chapter 4's neuron across many threads, whose completion order (and therefore the exact rounding accumulated) can vary run to run. A fixed, single-threaded CPU execution order removes this specific source of run-to-run bitwise drift.

FrozenSnapshotConfig.cs

```
1 public sealed record FrozenSnapshotConfig
2 {
3     public required string ModelWeightsHash { get; init; }
4     public required string RuntimeVersion { get; init; }
5     public required int RandomSeed { get; init; }
6     public required bool CpuOnlyInference { get; init; }
7 }
```

A Frozen Snapshot gives **bitwise reproducibility across time**: the same query, run against the same locked configuration next year, returns the identical output,

even though nothing in this mechanism is a simple lookup table — the full Transformer runs every time, just under conditions engineered to remove every remaining degree of freedom from the computation.

12.3.3 Intra-Vector Routing: Islands the Model Creates for Itself

The Static Vault and Frozen Snapshot are both, in different ways, set up in advance by a human engineer. **Intra-Vector Routing** is the mechanism Chapter 10, Section 10.4.2 already previewed: when an ILMVector's `Determinism.PointerTo` field references another vector's `VectorId`, resolution must follow that pointer to the target and defer to *its* determinism, not whatever value the original vector might otherwise have produced.

VectorResolver.cs

```

1 public interface IVectorStore
2 {
3     ILMVector? Get(Guid vectorId);
4 }
5
6 public static class VectorResolver
7 {
8     public static ILMVector Resolve(
9         ILMVector start, IVectorStore store, int maxHops = 8)
10    {
11        var current = start;
12        var visited = new HashSet<Guid> { current.VectorId };
13
14        // Follow the pointer chain only while the current vector is
15        // (a) tracked for determinism at all, (b) not yet resolved,
16        // and (c) actually defers somewhere. An ordinary vector with
17        // Determinism == null falls straight through and is
18        // returned as-is: it was never part of an island to begin
19        // with, so there is nothing to resolve.
20        while (current.Determinism is { IsDeterministic: false,
21            PointerTo: Guid next })
22        {
23            if (!visited.Add(next))
24                throw new InvalidOperationException(
25                    $"Cycle detected: vector {next} was already visited
26                    .");
27
28            if (visited.Count > maxHops)
29                throw new InvalidOperationException(
30                    "Pointer chain exceeded the maximum permitted hops.
31                    ");
32
33            current = store.Get(next)
34                ?? throw new InvalidOperationException(
35                    $"Dangling pointer: vector {next} does not exist.");
36        }
37        return current;
38    }
39 }

```

Two failure modes are guarded against explicitly, deliberately, and are not optional hardening: a **cycle** (vector *A* points to *B*, which points back to *A*) would otherwise loop forever, and a **dangling pointer** (a referenced `VectorId` that no longer exists, perhaps deleted by an update, Section 12.4.4) must fail loudly rather than silently falling back to an ungrounded, stochastic answer.

Arbitration Mechanisms

Three components govern how Intra-Vector Routing's result is actually used.

Arbiter, Causality Lock, Human-in-the-Loop

The **Arbiter** decides, for a given query, which mechanism governs the answer: Static Vault first (the strongest guarantee), then a resolved deterministic vector, then Frozen Snapshot, and only then the ordinary stochastic generator (Section 12.4.3 gives the full precedence order). The **Causality Lock** is engaged whenever a user explicitly requests a deterministic island (this book's own standing requirement, stated plainly: once engaged, *every* vector consulted for the remainder of that response must resolve to a Determinism block with `IsDeterministic = true`, with no silent fallback to stochastic generation anywhere in the path to the final answer). **Human-in-the-Loop** is the escalation path when the Causality Lock is engaged but no deterministic resolution exists for some part of the answer: the system reports that it cannot answer deterministically, rather than quietly relaxing the guarantee it was asked for.

CausalityLock.cs

```

1 public sealed class CausalityLock
2 {
3     public bool IsEngaged { get; private set; }
4     public void Engage() => IsEngaged = true;
5     public void Release() => IsEngaged = false;
6 }
7
8 public sealed class DeterminismViolationException : Exception
9 {
10     public DeterminismViolationException(string message) : base(message) { }
11 }
12
13 public static class GuardedResolver
14 {
15     // When the lock is engaged, every vector on the answer path
16     // must resolve to Determinism.IsDeterministic == true, or the
17     // request fails loudly -- it never silently falls back to the
18     // stochastic generator.
19     public static ILMVector ResolveUnderLock(
20         ILMVector start, IVectorStore store, CausalityLock
21         causalityLock)
22     {
23         var resolved = VectorResolver.Resolve(start, store);
24         var isResolved = resolved.Determinism is { IsDeterministic:
25             true };
26
27         if (causalityLock.IsEngaged && !isResolved)

```

```

26     {
27         throw new DeterminismViolationException(
28             $"Vector {resolved.VectorId} has no deterministic
29             resolution, " +
30             "but a Causality Lock is engaged for this response.");
31     }
32     return resolved;
33 }

```

12.4 Architecture of the Islands

12.4.1 How They Are Integrated into the ILM




Figure 12.1: The Deterministic Islands, integrated into the ILM. Every query passes through the Arbiter, which tries the mechanisms in order of guarantee strength (Section 12.4.3): Static Vault, Frozen Snapshot, Intra-Vector Routing, and only then the ordinary stochastic generator of Chapter 8. If a Causality Lock is engaged and none of the deterministic mechanisms resolve the query, the system escalates to Human-in-the-Loop rather than silently answering stochastically.

12.4.2 Code Examples: The Arbiter

```

Arbiter.cs
1 public sealed class Arbiter
2 {
3     private readonly StaticVault _vault;
4     private readonly IVectorStore _store;
5     private readonly CausalityLock _lock;
6
7     public Arbiter(StaticVault vault, IVectorStore store, CausalityLock
      causalityLock)
8     {
9         _vault = vault;
10        _store = store;

```

```

11     _lock = causalityLock;
12 }
13
14 public ILMVector Answer(string normalizedQuery, ILMVector?
15     candidate)
16 {
17     // 1. Static Vault: the strongest guarantee, tried first.
18     var vaultHit = _vault.Lookup(normalizedQuery);
19     if (vaultHit is not null) return vaultHit;
20
21     // 2. If the caller has a candidate vector (e.g. from RAG
22     //    retrieval, Chapter 11), resolve any pointer chain.
23     if (candidate is not null)
24         return GuardedResolver.ResolveUnderLock(candidate, _store,
25             _lock);
26
27     // 3. No deterministic path found.
28     if (_lock.IsEngaged)
29         throw new DeterminismViolationException(
30             "No deterministic mechanism resolved this query.");
31
32     throw new InvalidOperationException(
33         "Fall through to the stochastic generator (Chapter 8).");
34 }

```

12.4.3 Conflict Resolution Between Islands

Two mechanisms can, in principle, disagree — a Static Vault entry and an Intra-Vector-routed answer for what is nominally the same fact, perhaps because one was updated and the other was not yet. The Arbiter resolves this with a fixed precedence order, never a runtime heuristic: **Static Vault always wins**, because it is the only mechanism with an absolute, non-probabilistic guarantee; next, a resolved Intra-Vector-routed vector (`Determinism.IsDeterministic = true`); then Frozen Snapshot output. Any detected disagreement between two deterministic sources is *itself* logged as an error requiring human review (Section 12.4.4) — two sources that are both supposed to be authoritative and disagree indicates a data maintenance failure, not a case for the Arbiter to quietly adjudicate on its own.

12.4.4 Updating Islands: When and How

An `ILMVector`'s `Determinism` block is never edited in place. Updating a safety limit means: (1) the existing vector's `Determinism.ValidTo` is set to the effective date of the change (the vector remains queryable for historical questions about the period it was valid), and (2) a *new* `ILMVector`, with a new `VectorId` and a `Determinism` block whose `Version` is incremented and whose `ValidFrom` is set to the same effective date, is created to replace it. This append-only discipline is what makes `ValidFrom/ValidTo` meaningful at all (Chapter 10, Section 10.4.2) and gives every past answer a permanent, honest record of exactly which version of the truth produced it. Every such update requires human sign-off; this chapter's closing *Honest Boundary* states plainly why this is not automated.

12.5 Worked Example: An Island in Industrial Control

Two Questions, Two Very Different Guarantees

Question 1: "What is the maximum allowed pressure?" (clarified to pump P-101, discharge side). The Arbiter normalises the query, computes its SHA-256 hash (Section 12.3.1's worked example), and finds a Static Vault hit: an ILMVector whose Determinism block has `IsDeterministic=true`, encoding "120 bar". This answer is returned **identically**, every time this exact question is asked, by any user, on any server, indefinitely into the future — until a human-approved update (Section 12.4.4) supersedes it.

Question 2: "What do you recommend for the current situation?" No fixed fact answers this; it depends on the live operating context, and genuinely calls for the stochastic Transformer's generative reasoning (Chapter 8) over the current sensor state and retrieved procedures (Chapter 11's RAG). The Arbiter finds no Static Vault hit and routes to the stochastic generator — but with a boundary: if the generated recommendation references a specific numeric limit ("stay below the maximum discharge pressure"), *that specific reference* is resolved through the Static Vault or Intra-Vector Routing, not generated freely, exactly as Chapter 9, Section 9.4.3 already previewed for reinforcement-learned control policies: the open-ended reasoning stays stochastic and flexible; the hard numeric boundary it must respect does not.

Honest Boundary

Islands do not solve all problems. They guarantee reproducibility for the specific, enumerable class of facts placed inside them; they say nothing about whether the value stored was correct in the first place (Chapter 10's own closing Honest Boundary already made this point about data quality, and that warning applies to a Static Vault entry exactly as it applies to training data), and they do nothing at all for the vast majority of genuinely open-ended queries an ILM will field, which correctly remain the stochastic generator's responsibility.

Maintenance is complex. Section 12.4.4's append-only update discipline is deliberately conservative, and deliberately manual, for a reason worth stating plainly: an automated process that could edit a safety limit without human review would recreate exactly the unverifiable, opaque behaviour Deterministic Islands exist to eliminate, just one level removed. **Extension requires human intervention** by design, not by an oversight this book expects a future version to correct.

13. Evaluation and Verification of the ILM

13.1 What You Already Know

You already know the basic discipline of a held-out **test set**: never judge a model's quality using the same data it was trained on, or the judgement is meaningless (Chapter 3's bias-variance lesson, in its most practical form). What you may not yet have separated cleanly are two different questions that get asked in the same breath but demand very different evidence: **accuracy** (how often is the model right, on average, on data like its test set?) and **trust** (should a specific person, on a specific occasion, act on this specific answer?). A model can be highly accurate in aggregate and still be exactly the wrong thing to trust on the one query that matters most right now. This chapter is about closing that gap, not just measuring the first quantity more carefully.

13.2 Test Sets, Validation Sets, Cross-Validation

Chapter 3, Section 3.4.1 already introduced *k*-fold cross-validation for choosing a model's hyperparameters. A rigorous ILM evaluation keeps three data splits conceptually distinct: the **training set** (used to fit weights, Chapter 7), the **validation set** (used to choose hyperparameters and decide when to stop fine-tuning, Chapter 11), and the **test set** (touched exactly once, at the very end, to report a final number). Reusing the test set to make any decision, however small, quietly turns it into a second validation set and invalidates the final number's claim to be an honest estimate of performance on genuinely new data — a subtle, easy-to-commit version of Chapter 3's overfitting, one level up.

13.3 Metrics: Accuracy, Precision, Recall, F1, and AUC-ROC

The Confusion Matrix and Its Derived Metrics

For a binary classifier, every prediction falls into one of four categories: true positive (TP), false positive (FP), true negative (TN), false negative (FN).

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}, \quad \text{Precision} = \frac{TP}{TP + FP},$$

$$\text{Recall} = \frac{TP}{TP + FN}, \quad F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Why Accuracy Alone Misleads on Imbalanced Data

A fault classifier is evaluated on 200 held-out maintenance reports, 20 of which describe an actual fault (Chapter 10, Section 10.5.3's imbalance warning, now made concrete). The classifier produces:

	Predicted fault	Predicted normal
Actual fault (20)	TP = 15	FN = 5
Actual normal (180)	FP = 8	TN = 172

$$\text{Accuracy} = \frac{15 + 172}{200} = 93.5\%, \quad \text{Precision} = \frac{15}{23} \approx 65.2\%,$$

$$\text{Recall} = \frac{15}{20} = 75.0\%, \quad F_1 \approx 69.8\%$$

Now compare against the laziest possible classifier, one that predicts "normal" unconditionally: it scores $TP = 0$, $FN = 20$, $TN = 180$, for an accuracy of $\frac{180}{200} = 90.0\%$ — a suspiciously respectable-looking number from a model that has learned nothing whatsoever about faults, and would miss every single one in deployment. This is exactly Chapter 2, Section 2.3's base-rate trap, restated as a model-evaluation failure: whenever the positive class is rare, accuracy is dominated by how well the negative class is classified, almost regardless of what happens on the class that actually matters. Precision, recall, and F_1 do not have this blind spot, which is exactly why they, not raw accuracy, are the metrics this book recommends reporting for any imbalanced industrial classification task (Chapter 16 returns to this in full for anomaly detection specifically).




Figure 13.1: A representative ROC (Receiver Operating Characteristic) curve: true positive rate against false positive rate, traced out as the classification threshold sweeps from strict to lenient. The dashed diagonal is a coin flip ($AUC = 0.5$); this curve's Area Under the Curve, $AUC \approx 0.898$, summarises the classifier's ability to rank a random positive example above a random negative one, independent of any single chosen threshold — a useful complement to Precision/Recall/ F_1 , which all depend on wherever the threshold was actually set.

13.4 Robustness: Adversarial Examples and Distribution Shift

A model can score well on every metric in Section 13.3 and still fail in two specific, dangerous ways. An **adversarial example** is an input deliberately, often imperceptibly, perturbed to cause misclassification — a well-studied vulnerability of the CNNs of Chapter 8, and a genuine concern for any ILM component that processes external or partially untrusted input (Chapter 14, Section 14.4 returns to this as a security issue specifically). **Distribution shift** is a subtler, more common failure: the live data the deployed model sees gradually diverges from the training distribution (new equipment installed, a process change, a sensor recalibrated), degrading performance with no adversary involved at all. Chapter 2, Section 2.11's KL-divergence worked example is precisely the standard tool for detecting this: tracking D_{KL} between a reference distribution and a rolling window of recent inputs flags exactly this kind of drift, in an ILM's input distribution as readily as in the single-sensor example that first introduced it.

13.5 Explainability: SHAP, LIME, and Attention Visualization

Three complementary techniques make a model's individual predictions more inspectable, without requiring the model itself to be simple. **SHAP** (SHapley Additive exPlanations) borrows a solution concept from cooperative game theory: treat each input feature as a "player" contributing to the prediction, and attribute credit according to that feature's average marginal contribution across every possible subset of the other features — a principled, if computationally expensive, way to answer "how much did this particular input value push the prediction, accounting for everything else present?" **LIME** (Local Interpretable Model-agnostic Explanations) takes a cheaper route to a similar goal: fit a simple, interpretable model (Chapter 3's linear regression) to approximate the complex model's behaviour in a small neighbourhood around one specific prediction, then read the explanation off the simple local model instead of the complex global one. **Attention visualization** is specific to the Transformer (Chapter 8, Section 8.4): plotting the attention weights themselves shows, directly and without any separate approximation, which tokens or timesteps the model weighted most heavily in producing a given output — the cheapest of the three techniques, precisely because the explanation is a genuine internal quantity the model already computed, not a separate model built to approximate it after the fact.

13.6 Uncertainty Quantification

Every prediction examined so far has been a single point estimate. **Uncertainty quantification** asks a model to report *how confident* it is, not only what it predicts.

Bayesian neural networks take Chapter 2, Section 2.10's MAP estimation one step further: rather than a single point estimate $\hat{\theta}_{\text{MAP}}$ for the network's weights, maintain a full posterior distribution $P(\theta \mid \text{data})$ over them, and average predictions over that distribution at inference time. This is principled but expensive, and rarely used at the full scale of a modern Transformer.

MC Dropout (Monte Carlo Dropout) is the practical alternative, and a genuinely elegant reuse of existing machinery: Chapter 7, Section 7.6.4 introduced dropout as a *training-time* regulariser, randomly disabled at inference for a normal, single, deterministic-looking prediction. MC Dropout instead *leaves dropout active at inference*, runs the same input through the network many times, and treats the resulting spread of outputs as an empirical uncertainty estimate.

Quantifying Uncertainty in a RUL Prediction

An ILM's remaining-useful-life head predicts a bearing's RUL, in days, across ten independent MC Dropout forward passes on the identical input:

42.1, 45.3, 41.8, 47.2, 44.0, 43.5, 46.8, 40.9, 44.7, 45.9

The sample mean is $\bar{x} = 44.22$ days, with sample standard deviation $s = 2.15$ days, giving an approximate 95% interval

$$\bar{x} \pm 1.96s \approx [40.0, 48.4] \text{ days.}$$

A single point prediction of "44 days" invites a maintenance scheduler to treat that number as exact. The interval makes the model's genuine uncertainty visible and actionable: schedule the inspection with an 8-day margin in mind, not

a false sense of precision Section 13.6 has just shown the model never actually had. This is also a direct, practical illustration of Chapter 10, Section 10.3.3's "determinism where needed" principle in its mirror image: an open-ended predictive estimate is exactly the kind of output that *should* carry an honest uncertainty band, precisely because Chapter 12's Deterministic Islands are reserved for the narrower class of facts that must not.

13.7 Formal Verification: Specifications and Model Checking

Formal verification proves, mathematically, that a system satisfies a stated specification for *every* possible input, not merely the inputs in a test set. **Model checking** is the standard technique: exhaustively (or, for infinite state spaces, symbolically) explore a system's reachable states and verify the specification holds in all of them.

This book has, in fact, already given a small, complete example of exactly this kind of reasoning, in Chapter 12.

Verifying VectorResolver Terminates — A Specification, Proven

The specification: `VectorResolver.Resolve` (Chapter 12, Section 12.3.3) always terminates, either by returning a vector or by throwing a diagnosed exception, never by looping forever.

Proof. Each loop iteration either exits (the guard condition fails) or calls `visited.Add(next)`. Because `visited` is a `HashSet`, a repeated `next` causes `Add` to return `false`, which the code checks explicitly and turns into a thrown exception — so the loop cannot revisit a state without terminating. On every iteration that does not throw, `visited` strictly grows by exactly one element, and the explicit check `visited.Count > maxHops` throws once it exceeds a fixed bound. A strictly growing quantity, bounded above, cannot grow forever: the loop must exit within `maxHops` iterations, by exception if not by ordinary return. ■

This is a genuine instance of formal verification, at a small enough scale to state and check by hand: a specification ("this terminates"), and a proof that holds for *every* possible input graph the function could ever be given, not merely the ones tested. It also illustrates formal verification's real cost honestly: the property proven is narrow (termination), and required inspecting the exact implementation; nothing here proves the *returned value* is correct, only that the function does not hang. Larger properties, over larger systems, require correspondingly larger specifications and substantially more automated tooling.

13.8 Certification: Standards

Deploying an ILM in critical infrastructure increasingly means satisfying external standards, not only internal engineering judgement. **IEC 61508** and its sector-specific derivatives govern functional safety for systems whose failure can cause harm, predating AI but increasingly interpreted to cover AI-based safety functions. **ISO/IEC 42001** is an AI management system standard, addressing organisational process (governance, risk management, continual improvement) rather than any single model's technical properties. **NIST's AI Risk Management Framework** provides a structured process for identifying, measuring, and managing AI risk, without prescribing a specific pass/fail technical bar. None of these standards can

certify that a specific model's specific weights are "safe" in the way Chapter 12's Static Vault guarantees a specific fact; they certify that a rigorous *process* was followed in developing, evaluating, and monitoring the system — a different, complementary claim to everything else in this chapter.

Honest Boundary

Verification is not proof, for the system as a whole. Section 13.7's worked example proved one narrow, well-defined property of one small function. Nothing in this chapter, or this book, proves that an entire multi-billion-parameter Transformer behaves correctly on every conceivable input; that specific claim is, with current techniques, not provable at the scale of a modern ILM, and any vendor claiming otherwise should be asked exactly which specification was verified, over exactly which component, by exactly which method.

Certification is a process, not a stamp. Passing an IEC 61508 assessment or aligning with NIST's AI RMF is evidence that a disciplined process was followed at a point in time; it is not a permanent guarantee that survives the model's next fine-tune, the next data-drift event Section 13.4 described, or the next change to the surrounding system. Treat every certification in this chapter as the start of an ongoing monitoring obligation, not the end of one.

14. Ethics, Safety, and Governance

14.1 What You Already Know

Every engineer working in critical infrastructure already carries a professional duty of care that predates any of the AI in this book: a design decision that could plausibly harm someone is not merely a technical choice, it is an ethical one, and it is the engineer's, not merely the organisation's. You have also likely encountered the broader public debate about the **ethics of automation**: who is responsible when an automated system causes harm — the person who built it, the person who deployed it, the person who approved its use, or the system itself? This book has no interest in resolving that debate philosophically. It takes a narrower, engineering position, stated plainly now and defended through the rest of this chapter: responsibility does not transfer to a machine, ever, and every mechanism in this book, from Chapter 12's Deterministic Islands to this chapter's governance structures, is designed to keep a specific, named human accountable for a specific class of decision.

14.2 Bias and Fairness

Chapter 11, Section 11.9 already introduced the basic mechanism by which an ILM inherits bias: both from general pre-training data and from site-specific fine-tuning data. This section makes **fairness** itself precise enough to measure, and shows why measuring it is harder than it first appears.

Two Common Fairness Criteria

Demographic parity: the model flags examples at the same rate across groups, $P(\hat{y} = 1 \mid G = A) = P(\hat{y} = 1 \mid G = B)$. **Equalized odds:** the model's true positive rate *and* false positive rate are equal across groups, $P(\hat{y} = 1 \mid y = 1, G = A) = P(\hat{y} = 1 \mid y = 1, G = B)$ and likewise for $y = 0$.

The Same Classifier, Two Crews, Two Different Outcomes

A fault-flagging classifier is applied, unmodified, to maintenance reports from two crews:

	TP	FN	FP	TN
Crew 1 (100 reports, 20 faults)	16	4	12	68
Crew 2 (100 reports, 10 faults)	7	3	4	86

Crew 1: TPR = 80.0%, FPR = 15.0%, flag rate = 28%

Crew 2: TPR = 70.0%, FPR = 4.4%, flag rate = 11%

The *same* classifier, with no crew-specific logic anywhere in it, flags Crew 1's reports at more than double Crew 2's rate, and catches genuine faults at a meaningfully higher rate for Crew 1 too. Neither demographic parity nor equalized odds holds. This is not necessarily evidence the model is malfunctioning: it may simply reflect a genuine difference in each crew's equipment, workload, or reporting style, reflected honestly in the base rates (20 versus 10 actual faults). It is also not necessarily fine: if Crew 1's higher flag rate means their reports face more scrutiny and delay for reasons that trace back to reporting culture rather than genuine equipment risk, the classifier is amplifying an existing inequity, not measuring one. **A known impossibility result** (Chouldechova, Kleinberg et al.) makes this tension structural, not a sign of insufficiently careful engineering: whenever base rates genuinely differ between groups, as they do here, demographic parity and equalized odds cannot both hold exactly, for any classifier, except in degenerate cases. Choosing which fairness criterion to prioritise is a judgement call about the specific decision's consequences, not a parameter a training algorithm can optimise its way out of.

14.3 Transparency and Explainability

Chapter 13, Section 13.5 already gave the technical toolkit (SHAP, LIME, attention visualization) for explaining an individual prediction. Transparency is the broader, organisational discipline this toolkit serves: a stated commitment to make explainable *which* data and *which mechanism* produced a given answer, as a matter of course, not only when specifically challenged. This is not an abstract aspiration in this book's architecture; it is already, concretely, half-built. Every ILMVector's Determinism block (Chapter 10, Section 10.4.2), when present, carries a Version and validity window that is itself an audit trail: a Static Vault answer (Chapter 12, Section 12.3.1) can always be traced to the exact fact, and the exact version of that fact, that produced it. Transparency, in this architecture, is not a separate reporting layer bolted onto the ILM after the fact; it falls directly out of the same data structure built for determinism in the first place.

14.4 Security: Adversarial Attacks and Data Poisoning

Chapter 13, Section 13.4 introduced adversarial examples as a robustness concern; here they are a security concern, with a deliberate adversary rather than an accident of distribution shift. **Evasion attacks** craft an input, at inference time, to fool an already-trained model. **Data poisoning** is the more insidious threat for an ILM specifically: an attacker corrupts *training* data so that the resulting model learns a hidden, attacker-chosen behaviour — a backdoor triggered by some rare, specific pattern the attacker controls, invisible in ordinary validation because it was never designed to be found by ordinary testing.

This is precisely why Chapter 12, Section 12.4.4's human-sign-off requirement for updating a Deterministic Island is not solely a correctness control; it is a **security control** in its own right. An attacker able to insert or silently modify a Static Vault entry could plant a false safety limit with exactly the same absolute authority Section 12.3.1 built that mechanism to guarantee for legitimate facts — the append-only, human-approved update discipline is what stands between the Static Vault's

strength as a guarantee and its potential as an attack surface, and it is worth re-reading Chapter 12’s closing Honest Boundary now with this specific threat in mind.

14.5 Privacy: Federated Learning and Differential Privacy

Federated learning trains a shared model across multiple sites without ever centralising their raw data: each site trains locally on its own data and shares only model updates (gradients or weight deltas, Chapter 7), which are aggregated centrally — directly useful when several plants or operators want to benefit from a shared ILM without pooling commercially or legally sensitive operational data in one place.

Differential Privacy

A randomised mechanism $\mathcal{M}$ is ϵ -differentially private if, for any two datasets D, D' differing in one record, and any outcome S ,

$$P[\mathcal{M}(D) \in S] \leq e^\epsilon \cdot P[\mathcal{M}(D') \in S].$$

Smaller ϵ means a stronger guarantee: the mechanism’s output is nearly indistinguishable whether or not any single individual’s record was included at all.

Adding Calibrated Noise to a Sensitive Count

A query asks how many high-severity safety incidents occurred at a site last quarter; the true count is 7. The **Laplace mechanism** achieves ϵ -differential privacy for a count query (sensitivity $\Delta f = 1$, since one record changes the count by at most 1) by adding noise drawn from a Laplace distribution with scale $b = \Delta f / \epsilon$:

ϵ	Noise scale b	Sampled noise	Reported count
0.5 (strong privacy)	2.0	−3.76	3.24
2.0 (weaker privacy)	0.5	+0.41	7.41

The trade-off is visible in the numbers themselves, not just in the abstract: strong privacy ($\epsilon = 0.5$) adds enough noise that a single query’s reported count is only loosely related to the true value (each fresh query would draw a different noise sample); weaker privacy ($\epsilon = 2.0$) reports a number close to the truth but offers a correspondingly weaker guarantee about what an observer could infer about any single incident’s record. Neither setting is “correct” in the abstract; the choice of ϵ is a policy decision balancing statistical utility against individual privacy, made deliberately and documented, not defaulted to.

14.6 Regulatory Framework

The EU AI Act (Regulation 2024/1689, in force since August 2024) takes a risk-tiered approach: a small set of practices are prohibited outright; **high-risk** systems (a category that plausibly includes many industrial safety-relevant ILM deployments) face the Act’s most demanding obligations — risk management, data gov-

ernance, technical documentation, human oversight, and conformity assessment; limited-risk systems face narrower transparency duties (such as disclosing that a user is interacting with an AI system); minimal-risk systems face none of the above. As of this writing, the Act’s high-risk obligations, originally due to apply from August 2026, have been deferred by an amending “Digital Omnibus” package to December 2027 (for standalone, use-based high-risk systems) and August 2028 (for high-risk systems embedded in already-regulated products, the category most industrial control equipment would fall under) — dates that, like any live regulatory timeline, may move again before this book reaches its reader; verify the current schedule directly before treating any date here as final. **NIST’s AI Risk Management Framework** (Chapter 13, Section 13.8) complements this with a voluntary, process oriented structure (Govern, Map, Measure, Manage) rather than a binding legal tier system, and is often used by organisations operating in both the EU and US regulatory environments as a common internal process satisfying both.

14.7 Human Oversight

Human-in-the-Loop vs. Human-on-the-Loop

Human-in-the-loop: a human must actively approve an action before it takes effect — exactly Chapter 12, Section 12.3.3’s Human-in-the-Loop escalation, required whenever a Causality Lock is engaged and no deterministic mechanism resolves the query. **Human-on-the-loop:** the system acts autonomously within pre-approved bounds, with a human monitoring and able to intervene or override, but not required to approve every individual action.

Chapter 9, Section 9.4.3 already argued that reinforcement-learned control recommendations should never act directly on critical equipment; human-on-the-loop, with the human retaining override authority and the system’s action bounded by Chapter 12’s Deterministic Islands, is the specific oversight model that argument implies. Human-in-the-loop is reserved for the narrower, highest-stakes case: a query the system has determined it cannot answer deterministically at all, under a lock that was engaged specifically because a deterministic answer was required.

14.8 Responsibility and Accountability


Every mechanism this chapter and Chapter 12 have described — the Arbiter’s fixed precedence order, the append-only update log, the Causality Lock’s escalation to a named human reviewer — exists to answer one question clearly, in advance of ever needing to ask it under pressure: *who, specifically, is accountable for this answer?* For a Static Vault fact, it is whoever signed off the entry (Chapter 12, Section 12.4.4). For a stochastic recommendation, it is whoever approved deploying that model for that task, informed by Chapter 13’s evaluation. For an escalated Human-in-the-Loop decision, it is the specific reviewer who made the call. Nothing in this architecture allows “the AI decided” to function as an answer to who is accountable; every path through it terminates at a named human, by design.

Honest Boundary

Ethics is not a checklist. Every technique in this chapter — a fairness metric, a differential-privacy guarantee, a human-oversight model — is a tool that supports an ethical judgement; none of them, individually or in combination, substitutes for the judgement itself. A system that satisfies every metric in this chapter and is still deployed somewhere it should not be remains a failure of ethics, not a false negative in a compliance spreadsheet.

Governance is an ongoing process, not a project with an end date. Chapter 13's closing Honest Boundary already made this point for certification specifically; it applies to every mechanism in this chapter equally. Bias drifts as data drifts (Chapter 13, Section 13.4); regulatory obligations shift, as Section 14.6's own dates demonstrate even within the writing of this book; and the accountable human named in Section 14.8 must be someone whose responsibility is actively maintained, not a name filled in once at deployment and never revisited. This chapter closes Part III not because these questions are settled, but because the engineering mechanisms — the Deterministic Islands of Chapter 12, the evaluation discipline of Chapter 13, and the governance structures here — are now all in place to keep asking them properly, for as long as the ILM remains in service.

APPLICATIONS IN CRITICAL INFRASTRUCTURE



15. Predictive Maintenance

15.1 What You Already Know

You already know the difference between **preventive maintenance** (service equipment on a fixed schedule — every 6 months, every 2,000 operating hours — regardless of its actual condition) and **predictive maintenance** (service equipment when its actual, measured condition indicates it needs it, not before and not after). You also already know, intuitively, that these two policies trade off differently against the **cost of downtime**: preventive maintenance risks servicing healthy equipment unnecessarily (wasted cost, wasted planned downtime) or missing a fault that develops between scheduled visits; predictive maintenance risks nothing structurally, but only works as well as its ability to actually detect degradation before it becomes failure. Every technique in this chapter, from Chapter 8’s architectures to Chapter 13’s uncertainty quantification, is now being pointed at one specific, concrete question: *when, and with how much confidence, should a specific piece of equipment be serviced?*

The Cost Asymmetry, in Real Numbers

A mid-sized turbine’s unplanned outage costs approximately \$45,000 per hour in lost production and emergency repair premiums, and averages 14 hours to resolve: \$630,000 per event. Historically, this turbine suffers 3 such unplanned events per year, for an expected annual cost of **\$1,890,000**. A predictive maintenance programme, costing \$150,000 per year to operate, catches the great majority of developing faults early enough to convert them into scheduled interventions, reducing unplanned events to an estimated 0.5 per year:

$$\text{Annual cost with PdM} = 0.5 \times \$630,000 + \$150,000 = \$465,000.$$

(the first term is the residual unplanned-outage risk; the second is the programme’s own operating cost.) Net annual saving: **\$1,425,000** — a figure large enough, on equipment of this scale, that the programme’s own predictive model does not need to be anywhere near perfect to be worth deploying; Chapter 13’s evaluation discipline (Section 13.3’s precision/recall framing) is exactly how “good enough” gets decided quantitatively rather than by intuition.

15.2 Data: Vibration, Temperature, Pressure, Current

Four sensor modalities dominate predictive maintenance for rotating and electrical equipment, each diagnostic for a different failure mechanism. **Vibration** (accelerometers, typically sampled at several kHz) is the single richest signal for me-

chanical faults — bearing wear, imbalance, misalignment, gear tooth damage — because each of these produces a distinct, physically derivable frequency signature (Section 15.3.2). **Temperature** (thermocouples, RTDs, or infrared thermography) reveals friction, lubrication failure, and electrical resistance faults, generally at much lower sampling rates since thermal processes evolve slowly relative to vibration. **Pressure** indicates flow restrictions, seal degradation, and cavitation in pumps and compressors. **Current** (motor current signature analysis) detects electrical and some mechanical faults indirectly, through the subtle way a developing fault modulates the current drawn by a motor — valuable precisely because it requires no new sensor at all on existing motor-driven equipment, only access to the electrical supply already being metered.

15.3 Feature Engineering: Time-Domain and Frequency-Domain

15.3.1 Time-Domain Features

Standard Time-Domain Vibration Features

$$\text{RMS} = \sqrt{\frac{1}{N} \sum_{i=1}^N x_i^2}, \quad \text{Crest Factor} = \frac{\max |x_i|}{\text{RMS}},$$

$$\text{Kurtosis} = \frac{\frac{1}{N} \sum_i (x_i - \bar{x})^4}{\left(\frac{1}{N} \sum_i (x_i - \bar{x})^2 \right)^2}$$

RMS tracks overall vibration energy; crest factor and kurtosis both respond specifically to **impulsiveness** — short, sharp spikes riding on an otherwise smooth signal, exactly the signature a bearing defect produces every time a damaged surface strikes another. A healthy, smoothly periodic signal has kurtosis near 3 (a Gaussian’s kurtosis) or below; a signal dominated by sharp, rare impulses has kurtosis far above 3.

Healthy vs. Faulty Bearing, by the Numbers

A 0.2-second vibration sample (5 kHz sampling) from a healthy bearing and, separately, from a bearing with a developing outer-race defect, both containing the same underlying 50 Hz rotational component:

	RMS	Crest Factor	Kurtosis
Healthy	0.357	1.70	1.55
Faulty (outer-race defect)	0.542	5.80	10.02

RMS alone rises only modestly (0.357 to 0.542, a 52% increase) — easy to miss against normal operating variation. Kurtosis rises by a factor of **6.5**, from below the Gaussian baseline of 3 to well above it, precisely because the fault adds sharp impacts, not just more overall energy. This is exactly why kurtosis and crest factor, not RMS alone, are the standard first-line features for detecting the *onset* of impulsive bearing faults, well before the fault is severe enough to raise overall vibration energy substantially.

15.3.2 Frequency-Domain Features

Chapter 1's Fourier-adjacent intuition (a signal decomposed into frequency content) becomes directly diagnostic here: a bearing's geometry predicts, exactly, the frequency at which a specific defect location will produce impacts.

Ball Pass Frequency, Outer Race (BPFO)

$$\text{BPFO} = \frac{n}{2} f_r \left(1 - \frac{d}{D} \cos \theta \right)$$

where n is the number of rolling elements, f_r the shaft rotation frequency, d the ball diameter, D the pitch diameter, and θ the contact angle. (BPFI, the inner-race equivalent, replaces the minus sign with a plus.)

Diagnosing a Specific Bearing from Its Spectrum

A bearing with $n = 9$ balls, $d = 7.5$ mm, $D = 39$ mm, $\theta = 0^\circ$, on a shaft running at 1,780 RPM ($f_r = 1780/60 \approx 29.67$ Hz), has a predicted

$$\text{BPFO} = 4.5 \times 29.67 \times \left(1 - \frac{7.5}{39} \right) \approx \mathbf{107.8 \text{ Hz.}}$$

If a measured vibration spectrum shows a distinct peak at 107.8 Hz (and its harmonics, 215.6 Hz, 323.4 Hz, ...), this is not merely correlated with an outer-race defect; it is the specific, predicted physical signature of one, computed in advance from bearing geometry alone, before any machine-learning model ever sees the data. Feature engineering for vibration-based predictive maintenance routinely includes the energy present at these physically-derived frequencies as explicit input features, giving a downstream model (Section 15.4) domain-informed features no amount of raw-waveform training data alone reliably re-discovers on its own at practical dataset sizes.

15.4 Models: RNN, LSTM, and Transformer

Chapter 8 already built every architecture this section needs; predictive maintenance mainly determines which one fits the data at hand, following Chapter 8, Section 8.8's selection table directly. Short, regularly-sampled univariate or few-channel vibration windows often do well with a 1D CNN (Section 8.2) or LSTM (Section 8.3.2). Long-horizon degradation trends spanning weeks or months, combined with sparse maintenance-log text, favour a Transformer (Section 8.5), whose attention mechanism relates a current reading to relevant history at any distance without the vanishing-gradient cost Chapter 7, Section 7.5.1 identified for long RNN sequences.

15.5 Remaining Useful Life (RUL) Estimation

RUL estimation predicts how much operating time remains before a component crosses a failure threshold, framed either as direct regression (Chapter 3, Chapter 4) on time-to-failure, or as forecasting a continuous **health index** forward until it crosses a defined threshold.




Figure 15.1: A remaining-useful-life forecast: the health index (a composite of the time- and frequency-domain features of Section 15.3) declines with operating time, and RUL is the forecast distance to the failure threshold. This is the identical prediction task as Chapter 13, Section 13.6’s MC Dropout worked example: the same bearing, evaluated “today” at day 155, produced a mean RUL of 44.22 days with a 95% confidence interval of [40.0, 48.4] days (shaded band, days 195–203.4) — not a single point on this curve, but a genuine, quantified spread of plausible threshold-crossing times.

15.6 Examples: Turbines, Pumps, and Transformers

A **gas or steam turbine**’s dominant failure modes are blade fatigue and bearing wear, diagnosed primarily through vibration spectra (Section 15.3.2) and, for combustion turbines, exhaust temperature spread across measurement points. A **centrifugal pump**’s dominant modes are cavitation (diagnosed through a characteristic broadband high-frequency vibration signature and pressure fluctuation), seal wear, and bearing degradation (Section 15.3.1’s kurtosis-based features apply directly). A **power transformer**’s dominant degradation mode is insulation ageing, diagnosed not through vibration at all but through **dissolved gas analysis** (DGA): specific combinations of hydrogen, methane, ethylene, and other gases dissolved in the transformer’s insulating oil are a well-established, decades-old chemical fingerprint of specific fault types (partial discharge, thermal fault, arcing), and are a clean example of a domain-specific feature set no generic time-series architecture would discover without a domain expert first identifying which chemical measurements matter.

15.7 Predictive Maintenance Inside the ILM

Nothing in this chapter’s RUL estimate, health index, or cost calculation should be a Deterministic Island: each is a genuine prediction, correctly stochastic, and Section 15.5’s uncertainty band is precisely the honest way to report one. What *should* be Island-backed, inside a real ILM deployment, are the fixed facts those predictions are built from and compared against: the manufacturer’s rated bearing life, a maintenance-interval threshold set by regulation rather than by any model, or the \$630,000 unplanned-outage figure Section 15.1’s worked example treated

as a known input. Resolving *those* facts through Chapter 12's Static Vault, rather than letting a language model regenerate them on request, is what keeps a predictive maintenance report from quietly drifting: the prediction can improve from one model version to the next; the manufacturer's rated limit it is compared against must not silently change underneath it.

Honest Boundary

Prediction is not certainty. Every RUL estimate in this chapter is exactly what Chapter 13, Section 13.6 already demonstrated for this same worked example: a distribution, not a fact. Reporting "44 days" without the surrounding uncertainty band invites exactly the false precision Chapter 13 warned against, and nothing in this chapter's models changes that — a more sophisticated architecture produces a better-calibrated distribution, not a disappearing one.

Maintenance decisions remain human. A predictive model's output is, in every case in this chapter, a recommendation (Chapter 9, Section 9.4.3's framing, restated for this domain): the decision to pull a turbine from service, at real cost, on the strength of a probabilistic forecast, is exactly the kind of consequential judgement Chapter 14, Section 14.7's human-on-the-loop oversight model reserves for a named, accountable person, informed by the model rather than replaced by it.

16. Anomaly Detection

16.1 What You Already Know

You already have, from Chapter 1, the vocabulary this chapter needs most: the distinction between **noise** (random, expected variation, carrying no information worth acting on) and **signal** (a pattern that means something). An **anomaly**, in this chapter's sense, is exactly a signal that looks like noise until you know to look for it: a data point, or a pattern of data points, that does not fit the distribution the rest of the data was drawn from. You may also already sense the difference between an **outlier** (a point statistically unlikely under a *known* distribution) and **novelty** (a genuinely new regime the model has never seen at all, for which "unlikely under the known distribution" may not even be the right question). Every method in this chapter is a different mathematically precise way of asking the same question Chapter 9, Section 9.3.3 first posed: does this belong?

16.2 What Is an Anomaly? Point, Contextual, Collective

Three Kinds of Anomaly

Point anomaly: a single data point far from the rest of the data, considered in isolation — a pressure reading of 400 bar in a system rated to 150. **Contextual anomaly:** a value unremarkable in general, but anomalous given its context — a cooling-water temperature of 28°C is unremarkable in August and highly anomalous in January; the *same* value is normal or anomalous depending on Chapter 2, Section 2.3's conditioning variable. **Collective anomaly:** no single point in a sequence is individually unusual, but the *pattern* across the sequence is — a slow, steady upward creep in bearing temperature, each individual reading within normal range, but the trend itself signalling a developing fault no single-point method would ever flag.

This taxonomy matters operationally, not just conceptually: a method tuned to catch point anomalies (Section 16.3.1) will systematically miss collective anomalies, and no amount of additional point-anomaly tuning fixes that, because the failure is a mismatch of method to anomaly *type*, not a matter of sensitivity.

16.3 Methods: Statistical, Distance-Based, Density-Based, Reconstruction-Based

16.3.1 Statistical Methods

The simplest method is also the one every engineer already uses informally: flag a reading whose **z-score** $z = (x - \mu)/\sigma$ (Chapter 2, Section 2.4) exceeds a threshold, conventionally $|z| > 3$.

The Masking Effect: When the Anomaly Hides Itself

Twenty temperature readings, nineteen normal (mean 74.67, standard deviation 2.08) and one clear anomaly at 95.0. Computed against the *clean* statistics (the 19 normal readings alone), the anomalous reading's z-score is

$$z = \frac{95.0 - 74.67}{2.08} \approx \mathbf{9.77}$$

— unmistakable. But a real-time monitor typically estimates μ and σ *including* the anomaly itself, since it does not know in advance which point is anomalous. Recomputed over all twenty readings, the contaminated mean and standard deviation shift to 75.69 and 4.98, and the same reading's z-score drops to

$$z = \frac{95.0 - 75.69}{4.98} \approx \mathbf{3.88}$$

— still over threshold, but much less extreme. This is the **masking effect**: a single extreme point inflates the very statistics used to judge it, working to hide itself. With two or more simultaneous anomalies, or a threshold set only slightly above 3, masking can push a genuine anomaly's contaminated z-score below the detection threshold entirely. Robust statistical anomaly detection uses estimators (the median and median absolute deviation, rather than the mean and standard deviation) that are far less sensitive to exactly this kind of contamination.

16.3.2 Distance-Based and Density-Based Methods

Distance-based methods score a point by its distance to its k -th nearest neighbour (Chapter 9, Section 9.3.1's clustering geometry, read as a score rather than a cluster assignment): a large distance to the nearest cluster of normal points is itself the anomaly signal, with no assumption of Gaussian-distributed data required at all. **Density-based** methods refine this further: the **Local Outlier Factor** (LOF) compares a point's local density to the local density of its neighbours, rather than to the dataset as a whole, which lets it correctly flag a point as anomalous even when it sits in a moderately dense region, provided that region is still much sparser than its immediate neighbours' own neighbourhoods — catching anomalies that a single global density or distance threshold would miss entirely.

16.3.3 Reconstruction-Based Methods

Reconstruction-based methods, covered in full in Section 16.4, take an entirely different approach: train a model to compress and reconstruct *normal* data well, then flag any input it reconstructs *badly* as anomalous, on the logic that a model which has only ever learned to compress normal patterns has no efficient representation for anything else.

16.4 Autoencoders for Anomaly Detection

Chapter 8, Section 8.6 introduced the autoencoder as an architecture; here is its most direct industrial application. Trained *exclusively* on data from confirmed-normal operation, an autoencoder’s bottleneck (Chapter 1’s rate-distortion theory, architectural again) is forced to learn a compact representation of *normal* patterns specifically. At inference, the reconstruction error

$$e(\mathbf{x}) = \|\mathbf{x} - \hat{\mathbf{x}}\|^2$$

serves directly as an anomaly score: normal data, resembling the training distribution, reconstructs well (low e); anomalous data, unlike anything the autoencoder learned to compress, reconstructs poorly (high e). The threshold on $e(\mathbf{x})$ above which a point is flagged is a direct precision/recall trade-off (Chapter 13, Section 13.3), chosen deliberately rather than defaulted to.

16.5 One-Class SVM and Isolation Forest

One-Class SVM learns a boundary that encloses the bulk of the normal training data as tightly as possible in feature space, using the same kernel machinery that makes a support vector machine able to draw non-linear boundaries; any point falling outside the learned boundary at inference is flagged anomalous. **Isolation Forest** takes a structurally different, and computationally cheaper, approach:

Isolation Forest

Build many random trees, each splitting the data on a randomly chosen feature at a randomly chosen threshold. An anomaly, being “few and different,” is isolated (separated into its own leaf) by *far fewer* random splits, on average, than a normal point embedded in a dense cluster of similar points. The anomaly score is a function of the average path length $E[h(x)]$ across all trees: shorter average path $\Rightarrow$ higher anomaly score.

Isolation Forest, on Real Data

100 points sampled from a 2D Gaussian centred at the origin (representing normal multivariate sensor readings), plus two deliberately placed outliers at (6, 6) and (5.5, -5). Training an Isolation Forest (100 trees) and scoring every point (higher score = more anomalous):

	Mean anomaly score
Normal points (100)	-0.199
Planted anomalies (2)	+0.192

The separation is clean and large — roughly 0.39 units on a score scale where the entire normal population clusters tightly around -0.2 — achieved with no assumption about the data’s distribution shape at all, unlike Section 16.3.1’s z-score method, and without needing to train on the enclosing-boundary concept One-Class SVM relies on. This combination of simplicity, speed (each tree needs only $O(\log n)$ splits to isolate a typical point), and distribution-free operation is why Isolation Forest is often the first method tried on a new industrial anomaly-detection problem, before reaching for anything more elaborate.

16.6 Deep Learning: GANs and VAEs

Chapter 8’s generative architectures extend reconstruction-based detection further. A **VAE** (Section 8.6) trained on normal data scores anomalies using both reconstruction error *and* the KL-divergence term of its own training loss: an input that the encoder maps to a region of latent space far from where normal data concentrates produces a large KL term even before reconstruction is considered, a second, complementary signal beyond Section 16.4’s plain autoencoder. A **GAN**’s discriminator (Section 8.7), trained to distinguish real normal data from the generator’s output, can be repurposed directly as an anomaly scorer: an input the discriminator confidently rejects as “not like the real distribution it learned” is, by that same judgement, a candidate anomaly.

16.7 Evaluation: Precision@k and AUPRC

Chapter 13’s accuracy, precision, recall, and AUC-ROC all apply to anomaly detection, but Chapter 13, Section 13.3’s own worked example already showed accuracy misleading on imbalanced data — and anomaly detection is, essentially by definition, the most imbalanced classification problem this book examines. Two metrics are particularly suited to it.

Precision@k and AUPRC

Precision@k: rank all instances by anomaly score, and measure precision among only the top k — directly answering “of the k alarms an operator will actually investigate, how many are real?” **AUPRC** (Area Under the Precision-Recall Curve): analogous to Chapter 13’s AUC-ROC, but plotting precision against recall instead of true-positive rate against false-positive rate — far less forgiving than AUC-ROC when negatives vastly outnumber positives, because it does not let an overwhelming number of easy true negatives flatter the score.

Why AUPRC, Not AUC-ROC, for Rare Anomalies

A ranked list of 20 events (5 true anomalies) from an anomaly detector gives

$$\text{Precision@3} = \frac{2}{3} \approx 66.7\%, \quad \text{Precision@5} = \frac{3}{5} = 60.0\%,$$

$$\text{Precision@10} = \frac{4}{10} = 40.0\%,$$

with AUC-ROC = 0.880 and AUPRC = 0.774 — here, the two tell a broadly similar story. They diverge sharply, though, at realistic industrial imbalance ratios. A second, larger evaluation (1,000 events, only 10 true anomalies — a 1% positive rate, typical for rare industrial failures) run against a mediocre detector gives

$$\text{AUC-ROC} = 0.876 \quad (\text{looks good}), \quad \text{AUPRC} = 0.083 \quad (\text{reveals the truth}).$$

The same detector, on more realistically imbalanced data, looks strong on AUC-ROC and reveals itself as barely better than random ranking on AUPRC. AUC-ROC’s false comfort here comes from the same source as Chapter 13’s accuracy trap: with 990 easy true negatives to rack up correct calls against, the ROC curve stays close to the top-left corner almost regardless of how well the 10 genuine anomalies are actually ranked. AUPRC, driven entirely by precision among positive predictions, has nowhere to hide this. **Report AUPRC, not**

AUC-ROC alone, for any anomaly-detection task with realistic industrial imbalance.

16.8 Examples: Leak Detection, Sensor Failure, and Cyberattack

Leak detection in a pipeline is a textbook collective anomaly (Section 16.2): a small, gradual pressure or flow discrepancy between two measurement points, individually within normal sensor noise at any single instant, but persistent and one-directional over time in a way pure noise would not be. **Sensor failure** itself produces recognisable, specific anomaly signatures distinct from the process anomalies a sensor is meant to detect: a “stuck-at” failure (the reading freezes at a constant value while the true process continues to vary), a slow calibration drift (a collective anomaly, structurally identical to the leak-detection case above but in the sensor itself rather than the process), or intermittent dropout. **Cyberattack** detection (Chapter 14, Section 14.4) applies these same methods to network traffic and control-system command patterns rather than physical sensor readings: an Isolation Forest or autoencoder trained on normal control-network traffic flags an anomalous command sequence by the same logic Section 16.4 and Section 16.5 already established, with the “normal data” being legitimate operator and automated control traffic rather than a physical process signal.

16.9 Anomaly Detection Inside the ILM

An anomaly score itself — Section 16.3 through 16.6’s distance, reconstruction error, or isolation-path length — is correctly stochastic and model-dependent: it should change as a model is retrained on more data, and reporting it with the AUPRC discipline Section 16.7 established is the honest way to convey how much to trust it. The **alarm threshold** applied to that score, however, is frequently not a modelling choice at all: a regulatory reporting deadline once an anomaly is confirmed, or an engineering alarm setpoint fixed by a safety case (Chapter 19, Section 19.4), is exactly the kind of stable fact Chapter 12’s Static Vault exists to protect. An ILM built on this chapter’s methods should resolve that threshold through a Deterministic Island precisely so that two detectors, trained months apart on different data, still agree on the one number — where the line between “notable” and “actionable” actually falls — that no amount of model improvement is entitled to move quietly on its own.

Honest Boundary

Anomaly detection is inherently imbalanced. Every method in this chapter is being asked to find a small number of needles in an overwhelming haystack of normal operation, and Section 16.7’s worked example showed exactly how badly the wrong evaluation metric can hide poor real-world performance under this imbalance. Evaluate every anomaly detector the way Section 16.7 recommends, not the way that produces the most flattering single number.

False alarms have a cost, and it compounds. A detector tuned for high recall (catch every real anomaly) at the expense of precision floods operators with false alarms; the well-documented consequence, **alarm fatigue**, is that

operators begin discounting or ignoring alerts from a system that cries wolf too often — including, eventually, the alert that was real. The choice of detection threshold is therefore not a purely statistical decision to be optimised against a metric in isolation; it is an operational and human-factors decision, made in view of how many false alarms a specific operating team can absorb before the system's warnings, real ones included, stop being trusted at all.

17. Control and Optimization

17.1 What You Already Know

You already know the **feedback loop** that underlies every control system in this chapter: measure the process, compare the measurement to a setpoint, compute a control action from the difference, apply it, and measure again. You already know the **PID controller**, the workhorse implementation of that loop:

PID Control Law

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de(t)}{dt}$$

where $e(t) = r(t) - y(t)$ is the error between setpoint r and measured output y .

And you likely already know the **state-space** representation of a dynamic system, $\mathbf{x}_{k+1} = A\mathbf{x}_k + B\mathbf{u}_k$, $\mathbf{y}_k = C\mathbf{x}_k$ — a direct generalisation of Chapter 2, Section 2.8's Markov chain, with a control input $\mathbf{u}_k$ now added as a second lever, alongside the system's own dynamics, that shapes the next state.

A PID Controller, Computed

A furnace temperature loop has setpoint 350°C, current reading 342°C (error $e = 8.0$), accumulated integral error 45.0, and a derivative (rate of change of error) of -1.2 (the error is currently shrinking). With gains $K_p = 2.5$, $K_i = 0.15$, $K_d = 0.8$:

$$u = (2.5)(8.0) + (0.15)(45.0) + (0.8)(-1.2) = 20.0 + 6.75 - 0.96 = \mathbf{25.79}.$$

Each term contributes a distinct, interpretable share: the proportional term (20.0) reacts to the current error size; the integral term (6.75) pushes harder because error has persisted; the derivative term (-0.96) pulls back slightly because the error is already improving, damping overshoot. This decomposability, one number per physically distinct cause, is exactly what the neural control methods later in this chapter trade away in exchange for handling non-linear, high-dimensional dynamics no fixed three-term law can capture.

17.2 Model Predictive Control (MPC) with Neural Networks

Model Predictive Control replaces PID's reactive, single-instant correction with an explicit look-ahead: at every timestep, predict the system's trajectory over a horizon of N future steps under a candidate control sequence, choose the sequence that

minimises a cost function (tracking error plus, typically, a penalty on control effort), apply only the *first* action of that sequence, then repeat the entire optimisation at the next timestep with newly measured data — the **receding horizon** principle.

One Step of MPC, by Hand

A tank level follows $x_{k+1} = x_k + 0.5u_k$, currently at $x_0 = 6.0$, target $x_{\text{ref}} = 10.0$, horizon $N = 3$. Evaluating four candidate constant control actions against the cost $\sum_{k=1}^N (x_k - x_{\text{ref}})^2 + 0.1u^2$:

u	Predicted trajectory x_1, x_2, x_3	Cost
2	7.0, 8.0, 9.0	15.20
4	8.0, 10.0, 12.0	12.80
6	9.0, 12.0, 15.0	40.80
8	10.0, 14.0, 18.0	99.20

$u = 4$ wins: it reaches the target exactly at $k = 2$ and only mildly overshoots by $k = 3$, beating both the sluggish $u = 2$ (never reaches target within the horizon) and the aggressive $u = 6$ and $u = 8$ (large overshoot, heavily penalised). Only $u = 4$ is actually applied; at the next timestep, with a freshly measured x , this entire four-way comparison runs again. **Neural MPC** replaces the simple linear prediction model $x_{k+1} = x_k + 0.5u_k$ used here with a learned model — an LSTM or Transformer (Chapter 8), trained exactly as Chapter 15 trains a predictive maintenance model — when the true plant dynamics are too non-linear for a simple physics-based model to predict accurately, while keeping the receding-horizon optimisation loop itself unchanged.

17.3 Reinforcement Learning for Control

Chapter 9, Section 9.4 already built RL's foundations, including a worked Q-learning example for a simplified pump controller. For continuous control (a valve position, not a discrete Continue/ReduceLoad choice), policy-gradient methods (Chapter 9, Section 9.4.2) extend naturally to continuous action spaces, with algorithms such as DDPG and PPO handling the added complexity of an infinite, continuous action set rather than a small discrete one — the underlying REINFORCE logic (gradient ascent on expected reward via backpropagation, Chapter 7) is unchanged.

17.4 Safe RL: Constrained RL and Shielding

Chapter 9, Section 9.4.3 stated the position plainly: RL should never act directly on critical equipment without a boundary it cannot learn its way around. This section makes that boundary a formal, named technique, not just a warning.

Constrained RL and Shielding

Constrained RL reframes the optimisation itself: instead of maximising expected reward alone, maximise expected reward *subject to* an expected safety cost staying below a threshold, $\max_{\pi} \mathbb{E}[R]$ s.t. $\mathbb{E}[C] \leq \delta$, typically solved via a Lagrangian relaxation that penalises constraint violation during training. **Shielding** takes a different, stronger approach: rather than training the policy to *learn* to respect safety limits, wrap the trained policy's proposed action in an independent, non-learned safety layer that overrides any proposed action falling outside a hard-coded safe range, *before* it ever reaches the actuator.

Shielding, read carefully, is nothing other than Chapter 12's Deterministic Islands applied to a control action instead of a factual query: the shield's safe range is exactly the kind of static, auditable, non-stochastic constraint Chapter 12's Arbiter enforces, and the architecture composes directly.

Shield.cs — an RL policy's proposal, checked against a Deterministic Island

```

1 public sealed class Shield
2 {
3     private readonly Arbiter _arbiter;
4
5     public Shield(Arbiter arbiter) => _arbiter = arbiter;
6
7     // proposedAction: e.g. a valve opening percentage the RL
8     // policy recommends. The safe upper bound is resolved through
9     // exactly the same Arbiter (Chapter 12) that answers any other
10    // safety-critical query -- Static Vault first, then a routed
11    // ILMVector, never the stochastic generator.
12    public double Enforce(double proposedAction, string limitQuery)
13    {
14        ILMVector limitVector = _arbiter.Answer(limitQuery, candidate:
15            null);
16        double maxSafe = limitVector.Embedding[0]; // decoded safety
17            limit
18        return Math.Min(proposedAction, maxSafe);
19    }
20 }
```

A Shielded Action, in Practice

An RL policy, optimising for throughput, proposes opening a valve to 85%. The Shield resolves the query "maximum safe valve opening under current conditions" through the Arbiter, which returns a Static Vault hit (Chapter 12, Section 12.3.1) of **75%**. The enforced action is

$$\text{Enforce}(85, 75) = \min(85, 75) = \mathbf{75\%}.$$

The RL policy is free to *propose* 85% again on the next timestep, and to learn, over time, that proposals above 75% are never actually applied — but the plant is never exposed to an 85% opening while that learning happens, because the shield's bound comes from a deterministic, pre-verified source entirely outside the policy's own learned parameters.

17.5 Digital Twins

A **digital twin** is a continuously updated virtual replica of a physical system: a simulation model, structurally similar to Chapter 15's predictive models but run forward in time rather than fit to historical data, kept synchronised with the real system's current state via ongoing sensor measurement (a state-estimation problem, philosophically the same recursive update as Chapter 2, Section 2.10's Bayesian inference: prior belief about the system's state, updated by each new measurement). Digital twins serve two distinct roles in this chapter's methods: as the **training and testing ground** for RL policies (Chapter 9, Section 9.4.3's requirement that RL train in simulation before any real deployment, made concrete), and as the **prediction model** inside MPC's receding-horizon loop when the physical system is too complex or too risky to probe directly with exploratory control inputs.

17.6 Examples: Reactor Control, Combustion, and Energy Management

Reactor control (temperature and pressure regulation in a chemical or nuclear process) is the paradigm case for this chapter's entire safety argument: MPC or RL may optimise for efficiency and throughput, but the hard temperature and pressure limits are non-negotiable Deterministic Islands, shielded exactly as Section 17.4 describes, with certified, non-AI safety systems retaining final physical authority regardless of what any model recommends. **Combustion optimisation** balances fuel efficiency against emissions, a genuinely multi-objective problem MPC's cost function (Section 17.2) naturally accommodates by weighting both terms. **Energy management** (balancing generation, storage, and demand across a grid or micro-grid) operates at a longer timescale than the other two examples, and increasingly combines RL-based scheduling with hard constraints (line capacity limits, reserve margins) enforced exactly as shielding enforces a valve limit.

Honest Boundary

AI in the control loop must be certified. Every technique in this chapter that touches an actuator, not merely a display or a report, falls squarely within Chapter 13, Section 13.8's certification standards (IEC 61508 above all) and Chapter 14's governance requirements, not as an optional extra but as a precondition for deployment. A model that performs impressively in simulation (Section 17.5) has demonstrated nothing yet about its behaviour under a certification process's adversarial scrutiny, and should not be represented as ready for the control loop until it has been.

Safety is paramount, and it is not the model's job to provide it. Every method in this chapter, from PID through shielded RL, is more capable and more sophisticated than the one before it, and none of that growing sophistication changes where this book places the actual safety guarantee: not in the increasingly capable controller, but in the Deterministic Island the controller's output must pass through before it ever reaches physical equipment. That ordering — capability inside, hard limits outside, and never the reverse — is the single idea this entire book has been building toward since Chapter 1, and this chapter is where it finally meets an actuator.

18. Decision Support

18.1 What You Already Know

You already draw a sharp line, in ordinary professional life, between a **recommendation** and a **decision**. A colleague who recommends a course of action has offered you information and their own judgement; you remain free to weigh it, combine it with context they did not have, and choose differently. A **decision**, once made, carries **accountability** (Chapter 14, Section 14.8) in a way a mere recommendation does not. Every chapter in Part IV has, in fact, already respected this line without necessarily naming it: Chapter 15's RUL estimate, Chapter 16's anomaly flag, and Chapter 17's shielded control proposal were all, precisely, recommendations — numbers and flags a human or a downstream system acts on, never actions the ILM took by itself. This chapter names that pattern explicitly and gives it a formal structure.

18.2 Decision Support Systems

A **Decision Support System** (DSS) is the architecture that turns Part IV's individual predictive components into something a human can actually act on: it gathers evidence (a RUL estimate here, an anomaly score there, a shielded control recommendation elsewhere), synthesises it into a small number of clearly stated options, and presents those options, with their trade-offs made explicit, to an accountable person. This is not a new architecture layered on top of everything else in this book; it is the ILM itself, viewed from the output side, with Chapter 12's Deterministic Islands guaranteeing that any hard fact feeding into the recommendation is trustworthy and Chapter 13's evaluation discipline guaranteeing that any probability or prediction feeding into it is honestly calibrated.

When Does the Expected Cost Favour Acting Now?

Chapter 15's predictive maintenance model estimates a 65% probability that a bearing fails within the next 30 days if left in service. Two options are available: schedule immediate maintenance, at a known cost of \$80,000, or wait and monitor, risking Chapter 15's own \$630,000 unplanned-outage cost if the failure does occur.

$$E[\text{cost of waiting}] = (0.65)(\$630,000) + (0.35)(\$0) = \$409,500.$$

Since $\$409,500 \gg \$80,000$, immediate maintenance is the expected-cost-minimising choice. Solving for the failure probability at which the two options

break even,

$$p_{\text{breakeven}} = \frac{\$80,000}{\$630,000} \approx 12.7\%,$$

reveals something a single expected-cost number alone would not: even if the model's confidence were far lower than 65%, expected-cost logic would *still* favour immediate action down to a failure probability of about one in eight. This is exactly the kind of reasoning a Decision Support System should surface explicitly, not bury inside a single recommended action — the breakeven point tells the accountable human how sensitive the recommendation actually is to the model's uncertainty, which the recommendation alone does not.

18.3 Human-AI Collaboration

Chapter 14, Section 14.7 already named the two governing oversight models, human-in-the-loop and human-on-the-loop; decision support is where the distinction is felt most directly by the person doing the work. Effective human-AI collaboration in this setting means the system states not just its recommendation but its confidence (Chapter 13, Section 13.6's uncertainty quantification), its reasoning (Section 18.4), and, as the worked example above showed, how sensitive the recommendation is to being wrong — giving the accountable human what they need to override the recommendation on grounds the model could not have had access to (a planned outage window, a spare-parts shortage, a regulatory inspection already scheduled), rather than forcing a choice between blind acceptance and blind rejection.

18.4 Explainable AI for Decisions

Chapter 13, Section 13.5 already built the technical toolkit; for a decision support system specifically, explainability is not an optional diagnostic but part of the recommendation itself. A recommendation to schedule maintenance is far more useful, and far more checkable by exactly the person who must be accountable for acting on it, when it arrives already showing *which* features drove it (SHAP or LIME, Chapter 13, Section 13.5) or *which* part of an incoming maintenance report the model attended to most strongly (attention visualization, Chapter 8, Section 8.4) — the explanation is not decoration on top of the number; it is the part of the output that lets the human actually exercise the judgement Section 18.1 said they retain.

18.5 Scenario Analysis

A single expected-cost figure, as Section 18.2's worked example already hinted, can understate the risk a decision-maker actually cares about. **Scenario analysis** (typically implemented as a Monte Carlo simulation, Chapter 9's clustering geometry and Chapter 13's uncertainty methods both feeding into it) runs many simulated futures under a candidate policy and reports the *distribution* of outcomes, not just its mean.

The Tail Risk a Single Expected Value Hides

Simulating 10,000 scenarios of the “wait and monitor” policy from Section 18.2’s worked example, with the repair cost itself drawn from a realistic distribution around \$630,000 (rather than treated as a fixed number) when failure occurs:

Statistic	Value
Mean cost	\$410,578
Median cost	\$484,349
90th percentile	\$823,726
99th percentile	\$1,165,458
$P(\text{cost} > \$1,000,000)$	2.98%

The mean, \$410,578, closely matches Section 18.2’s simpler expected-value calculation (\$409,500) — a useful cross-check that the two methods agree where they should. But the mean alone hides that nearly **3% of simulated futures** under the “wait” policy cost over a million dollars, nearly triple the mean. A decision-maker who is risk-averse, or who is accountable to a board that will ask about worst-case exposure rather than only expected exposure, needs exactly this tail information, and no single expected-cost number, however correctly computed, delivers it. Scenario analysis is how a Decision Support System delivers it instead.

18.6 Examples: Fault Diagnosis, Maintenance Planning, and Crisis Management

Fault diagnosis decision support combines Chapter 9’s classification heads with Chapter 13’s explainability to present not just “which fault” but the evidence for it, letting a technician confirm or override the diagnosis before acting on it. **Maintenance planning** decision support is precisely Section 18.2’s worked example, generalised across an entire fleet: ranking many pieces of equipment by expected cost of inaction, letting a limited maintenance budget be allocated where the expected-cost argument is strongest, with Section 18.5’s scenario analysis available wherever the ranking is close enough that tail risk should break the tie. **Crisis management** decision support operates under the tightest time pressure of the three, synthesising live sensor data, Chapter 16’s anomaly flags, and Chapter 17’s control-system state into a small number of clearly ranked response options for an incident commander, with Chapter 12’s Deterministic Islands ensuring that any hard fact (a stated safety limit, a required response time) presented under this pressure is exactly right, not merely plausible.

Honest Boundary

AI advises; humans decide. This is not a slogan appended to close the chapter; it is the literal architecture Part IV has built, component by component, since Chapter 15: every predictive number, every anomaly flag, every control recommendation, and every decision support output in this chapter has been engineered specifically to inform a human’s judgement, never to substitute for it. The Deterministic Islands of Chapter 12 protect the facts a recommendation relies on; Chapter 13’s evaluation protects the honesty of its uncertainty; Chapter 14’s governance names who is accountable for acting on it. None of these

mechanisms, individually or together, make the decision. That step, by design, is the one place in this entire architecture where this book has never proposed that the ILM belongs.

19. Safety and Compliance

19.1 What You Already Know

You already work inside a **regulatory environment**: every plant this book has been written for operates under external safety oversight, inspection regimes, and documented compliance obligations that predate any AI in this book by decades. You may also already know the concept of a **safety case**: a structured, evidence-based argument that a specific system is acceptably safe for a specific application in a specific environment — not a claim that a system is safe in the abstract, but a documented case, open to scrutiny, that it is safe *for this use*. This chapter asks a narrow, practical question the rest of the book has been building toward: what does it actually take to bring an ILM through this kind of scrutiny, in the specific regulatory environments this book's readers work within?

19.2 Regulations: Nuclear (IAEA), Aerospace (FAA, EASA), Medical (FDA)

Three sectors most relevant to this book's readers are each, as of this writing, actively building AI-specific regulatory guidance on top of their existing safety frameworks, at different stages of maturity.

Nuclear (IAEA). The IAEA's 2025 publication *Considerations for Deploying Artificial Intelligence Applications in the Nuclear Power Industry* (Nuclear Energy Series No. NR-T-1.26) is explicit that AI adoption in nuclear power plants "remains in the nascent stages, characterized by pilot projects and limited deployment," and the guidance itself is framed as general considerations, not a binding certification standard — it complements, rather than replaces, the existing IAEA nuclear safety standards this book's readers already operate under. The IAEA's February 2026 International Symposium on AI and Nuclear Energy signalled continued, active development of this guidance rather than a settled framework.

Aerospace (FAA, EASA). EASA's AI Roadmap 2.0 and its associated Concept Paper progression are the most developed sector-specific AI certification framework among the three discussed here: guidance already exists for Level 1 AI ("human augmentation/assistance") and Level 2 AI ("human-AI cooperation and collaboration"), explicitly built around a human-oversight principle tied to the EU AI Act's Article 14. Level 3 AI ("advanced automation," where a human may not be present during operation) remained under active public consultation as of 2026, with EASA describing this as requiring "a strong limitation for safety-related systems" until further guidance matures. The FAA operates on a parallel track, coordinating internationally on AI certification approaches without, at the time of writing, having published a directly comparable staged framework of its own.

Medical (FDA). The FDA's framework is the most procedurally concrete of the three, built around three linked mechanisms: **Good Machine Learning Practice** (GMLP, developed jointly with Health Canada and the UK's MHRA) sets quality expectations for data, training, and validation; the **Predetermined Change Control Plan** (PCCP, finalised in FDA guidance through 2025) lets a manufacturer pre-specify exactly which future model updates are already authorised, avoiding a full new submission for every retraining cycle; and **Total Product Lifecycle** (TPLC) oversight replaces a one-time approval with continuous real-world performance monitoring. By early 2026, the FDA had authorised over 1,350 AI-enabled medical devices, roughly double the 2022 figure — concrete evidence that a lifecycle-based, rather than one-time, certification model can scale.

19.3 Certification of AI Systems

A common pattern emerges across all three sectors, despite their different domains and different maturity levels, and it lines up directly with mechanisms this book has already built. Each framework moves away from a single, one-time approval (the traditional model for static, non-learning systems) toward **lifecycle-based oversight**: the FDA's PCCP and TPLC name this explicitly; EASA's staged Level 1/2/3 framework and IAEA's "nascent stage" framing reflect the same underlying recognition that a learning system's behaviour is not fully fixed at the point of initial approval. Each framework also requires **documented human oversight** (EASA's Article 14 tie-in, this book's Chapter 14, Section 14.7) and **data and process quality** (the FDA's GMLP, this book's Chapter 10, Section 10.5). Chapter 13, Section 13.8's IEC 61508, ISO/IEC 42001, and NIST AI RMF sit underneath all three sector-specific frameworks as the general-purpose functional-safety and AI-governance infrastructure each of them draws on or references.

19.4 Safety Cases

A safety case is conventionally structured as a set of **claims**, supported by **argument**, supported in turn by **evidence** (the CAE pattern; Goal Structuring Notation, GSN, is the most common graphical notation for laying this structure out). This book's own architecture, read end to end, is already substantially a safety case for the ILM, waiting to be written up formally.

A Skeleton Safety Case for the ILM

Top claim: The ILM's outputs are acceptably safe for use in the decision-support role described in Chapter 18.

Sub-claim 1 — safety-critical facts are correct and reproducible. *Argument:* such facts are resolved exclusively through Deterministic Islands, never generated stochastically. *Evidence:* Chapter 12's Static Vault hash-lookup mechanism (Section 12.3.1) and its append-only, human-approved update discipline (Section 12.4.4).

Sub-claim 2 — predictions carry honest uncertainty. *Argument:* every probabilistic output is evaluated against held-out data using metrics appropriate to its class imbalance. *Evidence:* Chapter 13's precision/recall/AUPRC discipline (Sections 13.3, 16.7) and MC Dropout uncertainty quantification (Section 13.6).

Sub-claim 3 — no output acts on the physical system unsupervised. *Argument:* every control-relevant recommendation passes through a non-learned safety layer before reaching an actuator. *Evidence:* Chapter 17’s Shield mechanism (Section 17.4), composed directly from the same Arbiter used for Sub-claim 1.

Sub-claim 4 — accountability is traceable. *Argument:* every decision path terminates at a named, accountable human, and every fact’s provenance is recorded. *Evidence:* Chapter 14, Section 14.8’s accountability structure and Section 19.5’s traceability discipline, below.

A real safety case for a specific deployment would expand each of these four sub-claims with the specific evaluation results, audit logs, and sign-off records for that installation — but the *shape* of the argument, resting on exactly four pillars this book has already built chapter by chapter, does not change.

19.5 Auditing and Traceability

Every one of Section 19.4’s four sub-claims depends on being able to answer, after the fact, exactly which mechanism produced a given historical answer and why. This book’s architecture already carries the data needed: an ILMVector’s Determinism block (Chapter 10, Section 10.4.2) records Version, ValidFrom, and ValidTo for every fact it ever represents, and the Arbiter (Chapter 12, Section 12.4.2) determines, for every query, which mechanism — Static Vault, a routed vector, Frozen Snapshot, or the stochastic generator — answered it. A complete audit trail is simply the discipline of logging that Arbiter decision itself, every time, alongside the resolved vector’s own version metadata: not a new system bolted onto the ILM for compliance purposes, but the direct, textual record of decisions the architecture was already making internally.


19.6 Examples: AI Certification for Medical Devices and Autonomous Vehicles

Medical device AI certification is, as Section 19.2 detailed, the most mature of this chapter’s examples: a PCCP submitted alongside initial marketing authorisation lets a manufacturer specify in advance exactly how a diagnostic model may be retrained and redeployed without triggering a fresh full review each time, provided the retraining stays within the pre-agreed bounds — a direct, regulator-sanctioned parallel to Chapter 12, Section 12.4.4’s own append-only update discipline for a Deterministic Island. **Autonomous vehicles** sit in a less settled regulatory position than medical devices but draw on a comparable functional-safety base: ISO 26262 (automotive functional safety) and ISO/PAS 21448 (SOTIF, Safety of the Intended Functionality, specifically addressing hazards arising from a system’s performance limitations rather than from faults) play the role IEC 61508 plays more generally in Chapter 13, Section 13.8, with SOTIF in particular addressing exactly the gap this book has repeatedly named: a system can be free of traditional faults and still behave unsafely because of the inherent limitations of its learned, statistical nature.

Honest Boundary

Certification is necessary but not sufficient for trust. Chapter 13's own closing Honest Boundary already established that certification is a process, not a stamp; this chapter's tour through three live regulatory frameworks makes that abstract warning concrete. The IAEA's own guidance calls nuclear AI deployment "nascent"; EASA's own roadmap still limits Level 3 autonomy pending further work; even the FDA's comparatively mature PCCP framework governs *how* a model may change, not a guarantee that any given change is wise. None of the three frameworks in this chapter claims to certify that an AI system's specific judgement, on a specific case, was correct — they certify that a disciplined process stood behind it. Passing every certification gate this chapter describes is the beginning of a deployment's compliance obligations, not their conclusion, and no regulatory approval, however rigorously obtained, substitutes for the continuous evaluation, monitoring, and accountable human judgement this entire book has argued for from Chapter 13 onward.

SOFTWARE MODELING - DDD IN C#



20. Introduction to Software Modeling for Engineers

20.1 What You Already Know

You already practice **modularity** every time you design a system as separable sub-systems rather than one undifferentiated whole — a cooling loop, a control system, and an instrumentation package, each independently specifiable, testable, and replaceable. You already use **interfaces** in the general engineering sense: a pump’s datasheet specifies its inlet and outlet connections, its power requirements, and its performance envelope, without saying anything about the impeller geometry inside — exactly **separation of concerns**, the discipline of exposing only what a consumer needs to know and hiding everything else. This part of the book gives these three already-familiar engineering habits their precise software counterparts, in C#, so that every architecture this book has described — most importantly Chapter 12’s Deterministic Islands — can be read, from this point on, as actual, well-structured code rather than only as diagrams and prose.

20.2 Why Model Physics and AI with Code?

20.2.1 Simulations and Digital Twins

Chapter 17, Section 17.5’s digital twin is, concretely, a piece of running software that represents a physical system’s state and evolves it forward in time; how well that software is structured determines directly whether the twin stays comprehensible as its scope grows, or collapses into the kind of unmaintainable tangle this chapter’s discipline exists to prevent.

20.2.2 Educational Tools and Interactive Learning

A well-modelled simulation is not only a production tool; it is how an engineer builds intuition for a system too large, too slow, or too dangerous to experiment with directly — precisely the role this book’s own companion simulators, referenced throughout the NEXUS-1 series, are built to serve.

20.2.3 Data Analysis and Experimental Validation

Chapter 13’s evaluation discipline needs code that faithfully represents the domain being measured: a bug in how a “fault” is represented in code is a bug in every metric computed from it, no matter how correct the statistics.

20.2.4 Control Automation and Safety

Chapter 17, Section 17.4's Shield class was not a simplification for the page; production control software is held to exactly this standard of clarity, because Chapter 19's certification processes (Section 19.3) require reviewers to be able to read and verify exactly what a safety-relevant code path does.

20.3 Basic Principles of Object-Oriented Programming (OOP) in C#

The Four OOP Pillars

Classes and objects: a **class** is a blueprint (the datasheet); an **object** is a specific instance built from it (the physical pump on the floor with that datasheet). **Encapsulation:** an object exposes only the operations and data a consumer needs, hiding its internal state and implementation. **Inheritance:** a specialised class (Pump) can extend a more general one (Equipment), inheriting its shared structure while adding or overriding specifics. **Polymorphism:** code written against the general class (Equipment) works correctly, without modification, on any specialised class that extends it.

Equipment.cs — the four pillars, in one small hierarchy

```

1 public abstract class Equipment
2 {
3     public string Tag { get; }           // encapsulated identity
4     protected Equipment(string tag) => Tag = tag;
5     public abstract string Inspect();    // overridden per equipment
6     type
7 }
8 public sealed class Pump : Equipment   // inheritance
9 {
10     public double FlowRateM3PerHour { get; set; }
11     public Pump(string tag, double flowRate) : base(tag)
12         => FlowRateM3PerHour = flowRate;
13     public override string Inspect()    // polymorphism
14         => $"[Tag]: flow = {FlowRateM3PerHour} m3/h";
15 }
16
17 public sealed class Valve : Equipment
18 {
19     public double OpeningPercent { get; set; }
20     public Valve(string tag, double opening) : base(tag)
21         => OpeningPercent = opening;
22     public override string Inspect()
23         => $"[Tag]: opening = {OpeningPercent}%";
24 }
25
26 // A caller holding only List<Equipment> never needs to know which
27 // concrete type each item is -- this is polymorphism doing real work.
28 List<Equipment> plant = new() { new Pump("P-101", 120.0), new Valve("V
29     -12", 40.0) };
30 foreach (var item in plant) Console.WriteLine(item.Inspect());

```

Three further relationships describe how objects connect to one another: **composition** (a Reactor owns its FuelRods; the rods have no independent existence outside the reactor that contains them), **aggregation** (a MaintenanceTeam references Pumps it services, but the pumps exist independently of any one team), and **dependency** (a ReportGenerator briefly uses an ISensorReader to do its job, without owning or referencing it afterward). Modern C# adds several features this book has already used extensively without naming them: **properties** (the { get; init; } syntax throughout Chapters 10-12), **generics** (IReadOnlyDictionary<string, ILMVector> in Chapter 12's StaticVault), **LINQ** (declarative collection queries), and **async/await** (for I/O-bound operations like a database or network call, not used in this book's synchronous worked examples but standard in production inference services).

20.4 Introduction to Domain-Driven Design (DDD)

Domain-Driven Design is a discipline for keeping code and domain expertise in permanent correspondence, so that a plant engineer and a software engineer, reading the same class names, mean the same thing.

20.4.1 Ubiquitous Language

A **Ubiquitous Language** is a vocabulary, agreed between domain experts and developers, used consistently in conversation, documentation, and code alike — if operators say “trip,” the code says Trip, never a generic Event or Flag that requires a mental translation step every time it's read.

20.4.2 Bounded Contexts

A **Bounded Context** is the boundary within which a Ubiquitous Language, and the model built from it, applies consistently. “Flow” means something specific and consistent inside a HydraulicsContext; the same word inside a FinancialReportingContext might mean cash flow entirely. Chapter 21 and Chapter 22 are each, formally, a separate Bounded Context — the quantum-mechanical domain and the AI/ILM domain, respectively, each with its own internally consistent vocabulary that need not, and should not, be forced to share terms across the boundary.

20.4.3 Entities vs. Value Objects

Entities vs. Value Objects

An **Entity** has an **identity** that persists through change: the same physical pump, identified by its serial number, remains “the same pump” even after its flow rate, its impeller, or its owner changes. A **Value Object** has no identity at all; it is defined entirely by its attributes, and two Value Objects with identical attributes *are* the same value, interchangeably.

Entity vs. Value Object, contrasted directly

```
1 // Value Object: immutable, defined by its value alone.
2 public readonly record struct Pressure
3 {
```

```

4     public required double Bar { get; init; }
5 }
6 // Two Pressure values with Bar=120.0 are simply equal -- there is
7 // no sense in which they are "different pressures that happen to
8 // have the same reading."
9
10 // Entity: defined by identity, which survives attribute changes.
11 public sealed class Pump
12 {
13     public Guid Id { get; } // identity: never changes
14     public string SerialNumber { get; }
15     public double FlowRateM3PerHour { get; set; } // attributes: can
        change
16
17     public Pump(Guid id, string serialNumber, double flowRate)
18     {
19         Id = id;
20         SerialNumber = serialNumber;
21         FlowRateM3PerHour = flowRate;
22     }
23 }
24 // Two Pump objects with the same Id are the same pump, even if one
25 // has a stale FlowRateM3PerHour -- identity, not current attribute
26 // values, is what equality means for an Entity.

```

20.4.4 Aggregates and Aggregate Roots

An **Aggregate** is a cluster of related Entities and Value Objects treated as a single consistency boundary; the **Aggregate Root** is the one Entity through which all changes to the cluster must pass, so that the aggregate's **invariants** — rules that must always hold — can never be violated by a change made to one part in isolation.

CoolantLoop.cs — an Aggregate Root enforcing an invariant

```

1 public sealed class CoolantLoop
2 {
3     private readonly List<Pump> _pumps = new();
4     public IReadOnlyList<Pump> Pumps => _pumps;
5     public double MinimumTotalFlowM3PerHour { get; }
6
7     public CoolantLoop(double minimumTotalFlow)
8         => MinimumTotalFlowM3PerHour = minimumTotalFlow;
9
10    public void AddPump(Pump pump)
11    {
12        _pumps.Add(pump);
13        EnforceInvariant();
14    }
15
16    public void UpdateFlow(Guid pumpId, double newFlow)
17    {
18        var pump = _pumps.Single(p => p.Id == pumpId);
19        pump.FlowRateM3PerHour = newFlow;
20        EnforceInvariant(); // re-checked on every change, always
21    }

```

```

22
23     private void EnforceInvariant()
24     {
25         double total = _pumps.Sum(p => p.FlowRateM3PerHour);
26         if (total < MinimumTotalFlowM3PerHour)
27             throw new InvalidOperationException(
28                 "CoolantLoop invariant violated: total flow below
29                 minimum.");
30     }

```

No code anywhere else in the system is permitted to reach into `CoolantLoop`'s internal `_pumps` list directly and change a flow rate; every change must go through `CoolantLoop` itself, which is exactly what keeps the invariant unbreakable by construction, not merely by convention.

20.4.5 Repositories

A **Repository** is an abstraction for retrieving and storing Aggregates, hiding whatever actual storage mechanism (a database, a file, an in-memory cache) sits behind it: `IVectorStore` in Chapter 12, Section 12.3.3 is exactly a repository for `ILMVector`, already written, already in use, three chapters before this one named the pattern.

20.4.6 Domain Services

A **Domain Service** holds domain logic that does not belong naturally to any single Entity or Value Object — Chapter 12's `Arbiter` and `VectorResolver` are both Domain Services: the logic of *which* mechanism should answer a query is not a property of any one `ILMVector`, but a coordination rule that spans several of them.

20.4.7 Domain Events

A **Domain Event** records that something significant happened in the domain, for other parts of the system to react to asynchronously:

A Domain Event

```

1 public sealed record PumpFlowDroppedBelowThreshold(
2     Guid PumpId, double CurrentFlow, DateTime OccurredAt);

```

Chapter 22, Section 22.7 formalises exactly this pattern for the ILM itself: events such as `ModelTrained` or `IslandTriggered` are Domain Events in precisely this sense, letting, for instance, an audit log (Chapter 19, Section 19.5) subscribe to every `IslandTriggered` event without the `Arbiter` needing to know that any auditing is happening at all.

20.5 The Translation Table: Physics → DDD

Physics concept	DDD concept	Example
Domain of validity (e.g. Newtonian mechanics holds only at speeds $\ll c$)	Bounded Context	Classical Mechanics Context vs. Relativistic Context
Conservation law (energy, mass, momentum)	Aggregate invariant	CoolantLoop's minimum-flow check, above
Dimensioned quantity (120 bar, 350°C)	Value Object	Pressure, above

This is not a loose metaphor; it is the working design principle behind every code example the rest of Part V builds. A bounded context's boundary is exactly where a model's assumptions (Newtonian, non-relativistic; or, in Chapter 22, "this vector has been resolved deterministically") stop being guaranteed to hold. An aggregate's invariant is exactly the software-enforced version of a law the physical system itself obeys unconditionally. A value object's immutability mirrors a physical quantity's own nature: a pressure reading of 120 bar does not "become" a different pressure by being copied or passed around; it simply *is* 120 bar, wherever it appears.

Honest Boundary

DDD is a discipline, not a dogma. Not every class needs to be classified as strictly Entity or Value Object, not every cluster of objects benefits from being forced into an Aggregate, and applying this chapter's vocabulary mechanically, everywhere, produces over-engineered code with no more clarity than the tangle it replaced.

The physics analogy is a teaching bridge, not an identity. Chapter 1 through Chapter 19 used physical and statistical metaphors throughout to make software and mathematical concepts concrete for an audience of engineers; this chapter has done the same in reverse, using DDD to make software structure concrete by analogy to physics these readers already know cold. A bounded context is not literally a domain of validity, and an aggregate invariant is not literally a conservation law — the analogy is a bridge for understanding, sturdy enough to carry real design decisions, as Chapter 22 will show in full, but a bridge is still not the two things it connects.

21. Bounding the Quantum Context (Revisited)

21.1 What You Already Know

If you have read Volume II of this series, *From Core to Quantum*, you already carry the physics this chapter needs: a quantum state as a vector in a complex Hilbert space, superposition, the Born rule connecting amplitudes to measurement probabilities, unitary evolution under a Hamiltonian, and entanglement as a joint state that cannot be factored into independent single-qubit states. This chapter does not re-derive any of that physics.

If you have not read Volume II, four facts are enough to follow everything below: a quantum state is a vector of *complex* numbers, not a single number, whose squared magnitudes are probabilities that must sum to one; a measurement returns exactly one of these possibilities, chosen at random according to those probabilities, and irreversibly discards the rest; a quantum gate is a length-preserving (**unitary**) linear transformation of that vector; and two systems are **entangled** when their combined state cannot be split into two separate, independent per-system descriptions. Volume II derives all four from first principles; this chapter does not, but every one of them is also directly visible in the code that follows — reading this chapter as an *operational* description of what these four facts imply for software, checkable against Section 21.5's fully worked example, is a legitimate way to read it even before the underlying physical intuition feels fully secure.

This chapter does something narrower and, for this book's purpose, more valuable: it takes Chapter 20's Domain-Driven Design vocabulary and applies it, in full, to exactly one Bounded Context — quantum computing — producing real, runnable C# that a software engineer with no quantum background and a physicist with no software background can both read and agree describes the same thing.

21.2 Value Objects: StateVector, Operator, Basis

A quantum state, a gate, and a choice of basis are all, in Chapter 20, Section 20.4.3's exact sense, **Value Objects**: none of them has identity; two StateVectors with identical amplitudes are not "two states that happen to agree," they *are* the same state, interchangeably, exactly as two Pressure values of 120 bar were the same value in Chapter 20.

Quantum Value Objects

```

1 using System.Numerics;
2
3 public readonly record struct StateVector
4 {
5     public required Complex[] Amplitudes { get; init; }
6
7     public double Probability(int basisIndex)
8         => Amplitudes[basisIndex].Magnitude * Amplitudes[basisIndex].
           Magnitude;
9 }
10
11 public readonly record struct Operator
12 {
13     public required Complex[,] Matrix { get; init; }
14
15     public StateVector Apply(StateVector input)
16     {
17         int n = input.Amplitudes.Length;
18         var result = new Complex[n];
19         for (int i = 0; i < n; i++)
20         {
21             Complex sum = Complex.Zero;
22             for (int j = 0; j < n; j++)
23                 sum += Matrix[i, j] * input.Amplitudes[j];
24             result[i] = sum;
25         }
26         return new StateVector { Amplitudes = result };
27     }
28 }
29
30 public readonly record struct Basis
31 {
32     public required string[] Labels { get; init; } // e.g. {"|0>",
33         "|1>"}
34 }

```

Every gate this chapter uses — Hadamard, CNOT — is simply a specific Operator value: a fixed matrix, with no behaviour beyond Apply, exactly matching Chapter 20’s principle that a Value Object is defined entirely by its content.

21.3 Entities: Qubit and QuantumSystem

A Qubit, by contrast, has **identity** (Chapter 20, Section 20.4.3): “qubit 0” of a register remains “qubit 0” throughout a computation even as gates change what its measurement would reveal. QuantumSystem is the **Aggregate Root** (Chapter 20, Section 20.4.4) that owns the joint StateVector across every qubit in the register and enforces the one invariant quantum mechanics itself never violates: total probability sums to exactly 1.

Qubit.cs and QuantumSystem.cs

```

1 public sealed class Qubit
2 {
3     public Guid Id { get; }
4     public int Index { get; } // position within the register
5
6     public Qubit(Guid id, int index) { Id = id; Index = index; }
7 }
8
9 public sealed class QuantumSystem
10 {
11     private readonly List<Qubit> _qubits;
12     public StateVector State { get; private set; }
13
14     public QuantumSystem(int numQubits)
15     {
16         _qubits = Enumerable.Range(0, numQubits)
17             .Select(i => new Qubit(Guid.NewGuid(), i)).ToList();
18         var amplitudes = new Complex[1 << numQubits];
19         amplitudes[0] = Complex.One; // initial state
20         // |00...0>
21         State = new StateVector { Amplitudes = amplitudes };
22     }
23
24     public void ApplyGate(Operator gate)
25     {
26         State = gate.Apply(State);
27         EnforceNormalization();
28     }
29
30     private void EnforceNormalization()
31     {
32         double total = State.Amplitudes.Sum(a => a.Magnitude * a.
33             Magnitude);
34         if (Math.Abs(total - 1.0) > 1e-9)
35             throw new InvalidOperationException(
36                 "QuantumSystem invariant violated: state is not
37                 normalised.");
38     }
39 }

```

21.4 Modeling Quantum Collapse and Entanglement

Measurement is modelled as a method that samples a basis outcome according to the Born rule and then **collapses** the state — an operation with no classical software analogue, since it is irreversible and destroys the superposition Chapter 1’s information theory would describe as having just delivered exactly $-\log_2 P(\text{outcome})$ bits about the system, at the cost of every other possibility.

Measurement and Collapse

```

1 public int Measure(QuantumSystem system, Random rng)
2 {

```

```

3   var amps = system.State.Amplitudes;
4   double r = rng.NextDouble();
5   double cumulative = 0.0;
6   for (int i = 0; i < amps.Length; i++)
7   {
8       cumulative += system.State.Probability(i);
9       if (r < cumulative)
10      {
11          CollapseTo(system, i);
12          return i;
13      }
14  }
15  CollapseTo(system, amps.Length - 1); // floating-point fallback
16  return amps.Length - 1;
17 }
18
19 private void CollapseTo(QuantumSystem system, int basisIndex)
20 {
21     var amps = new Complex[system.State.Amplitudes.Length];
22     amps[basisIndex] = Complex.One;
23     // (System exposes an internal setter for State; omitted here.)
24 }

```

Entanglement needs no separate class or flag at all, which is itself the point worth naming explicitly: a `QuantumSystem` with n qubits already has a single joint `StateVector` of 2^n amplitudes, not n separate single-qubit states, so a state that cannot be written as a product of single-qubit states (an entangled one) is simply an ordinary `StateVector` value like any other. The domain model does not need to represent entanglement because it never represented qubits as separately-stateful in the first place — exactly the kind of modelling decision Chapter 20’s closing Honest Boundary warned would matter: forcing a `Qubit` Entity to carry its own `StateVec`tor would have made entanglement inexpressible, a real design error a DDD-literate but physics-naive developer could easily make.

21.5 Code Example: Gate Application, Measurement, and Time Evolution

A Verified Bell State, End to End

Starting from $|00\rangle$, applying a Hadamard gate to qubit 0 (implemented as $H \otimes I$ on the joint state) followed by a CNOT gate (control qubit 0, target qubit 1):

$$|00\rangle \xrightarrow{H \otimes I} \frac{|00\rangle + |10\rangle}{\sqrt{2}} \xrightarrow{\text{CNOT}} \frac{|00\rangle + |11\rangle}{\sqrt{2}}$$

the `QuantumSystem`’s `State.Amplitudes` become exactly (0.7071, 0, 0, 0.7071) in the $\{|00\rangle, |01\rangle, |10\rangle, |11\rangle\}$ basis — the canonical Bell state. Running `Measure` 100,000 times on independently prepared copies of this state (a Monte Carlo verification, exactly Chapter 13’s evaluation discipline applied to a quantum simulator instead of a classifier) gives:

Outcome	00⟩	01⟩	11⟩
Predicted probability	50%	0%	50%
Empirical frequency (100,000 trials)	50.01%	0.00%	49.99%

(|10⟩ is omitted; it was likewise measured 0.00% of the time.) The two qubits are never observed to disagree: measuring qubit 0 as 0 means qubit 1 is measured as 0 every single time, across all 100,000 trials, and likewise for outcome 1 — the defining, verifiable signature of entanglement, produced here by nothing more exotic than `QuantumSystem.ApplyGate` called twice.

Honest Boundary

This models finite-dimensional quantum mechanics only. Every class in this chapter represents a system as a finite complex vector in $\mathbb{C}^{2^n}$ — exactly right for qubits, and for any other genuinely discrete, finite-level quantum system, but not applicable as written to continuous degrees of freedom (a particle’s position or momentum, a harmonic oscillator’s infinite ladder of energy levels), which require infinite-dimensional Hilbert spaces and a substantially different computational representation entirely. Volume II’s treatment of the hydrogen atom and of continuous-variable systems relies on exactly this distinction, and this chapter’s `StateVector` class should never be reached for outside the finite-dimensional, discrete-level regime it was built for.

22. Bounding the AI Context

S,

Deterministic Islands

The AI Bounded Context (Chapter 20, Section 20.4.2) is the software model of everything Chapters 8 through 14 built: the Transformer architecture and its trained weights, the pre-training and fine-tuning processes that produce them, the act of answering a query, the query itself, and — the one piece of this context with no counterpart in an ordinary machine learning system — Chapter 12's Deterministic Islands. This chapter gives every one of these a precise place in Chapter 20's Entity, Value Object, and Aggregate vocabulary, and shows that the code already written in Chapters 10 through 12 was, without ever using these words, already correctly structured this way.

22.2 Domain Terminology

The AI Context's Ubiquitous Language

Model: a trained Transformer artifact — an architecture plus a specific set of weights. **Layer:** one building block of a Model's architecture (dense, convolutional, attention; Chapters 4, 8). **Tensor:** a multidimensional array of numbers, with a shape and a data type — the universal currency weights, activations, and embeddings are all expressed in. **Prompt:** the input to a single inference request. **Inference:** the act of producing an output from a Prompt. **TrainingSession:** one run of the training process (Chapters 7, 11), with its hyperparameters and resulting metrics. **Island:** a deterministic anchor — a hash and its associated, pre-verified answer (Chapter 12, Section 12.3.1).

22.3 Value Objects

Tensor, Layer, Prompt, and Island share no identity at all: two Tensors holding the same numbers in the same shape *are* the same tensor, exactly as Chapter 21's StateVector had no identity beyond its amplitudes.

AI Context Value Objects

```

1 public readonly record struct Tensor
2 {
3     public required float[] Data { get; init; }

```

```

4     public required int[] Shape { get; init; }
5 }
6
7 public readonly record struct Layer
8 {
9     public required string Kind { get; init; }    // "Dense", "
           Convolutional", "Attention"
10    public required int[] Shape { get; init; }
11 }
12
13 public readonly record struct Prompt
14 {
15     public required string Text { get; init; }
16
17     // The exact normalisation StaticVault.HashQuery (Chapter 12,
18     // Section 12.3.1) expects -- defined once, here, and reused.
19     public string Normalized => Text.Trim().ToLowerInvariant();
20 }
21
22 // A deterministic anchor: a hash and its pre-verified answer.
23 // Answer is an ILMVector (Chapter 10) whose Determinism block
24 // carries IsDeterministic = true -- Island does not duplicate that
25 // metadata, it names the (hash, answer) pair that Chapter 12's
26 // StaticVault has been storing since it was first introduced.
27 public readonly record struct Island
28 {
29     public required string Hash { get; init; }
30     public required ILMVector Answer { get; init; }
31 }

```

22.4 Entities

Model, TrainingSession, and InferenceSession each have identity that survives change: retraining a Model produces new Weights for the *same* model identity; appending a loss value to a TrainingSession in progress does not make it a different session.

AI Context Entities

```

1 public sealed class Model
2 {
3     public Guid Id { get; }
4     public IReadOnlyList<Layer> Architecture { get; }
5     public IReadOnlyList<Tensor> Weights { get; private set; }
6
7     public Model(Guid id, IReadOnlyList<Layer> architecture,
8         IReadOnlyList<Tensor> weights)
9     {
10         Id = id; Architecture = architecture; Weights = weights;
11     }
12
13     public void UpdateWeights(IReadOnlyList<Tensor> newWeights) =>
14         Weights = newWeights;
15 }

```

```

14 public sealed class TrainingSession
15 {
16     public Guid Id { get; } = Guid.NewGuid();
17     public IReadOnlyDictionary<string, double> Hyperparameters { get; }
18     public List<double> LossHistory { get; } = new(); // Chapter 7's
19         loss curve, tracked
20
21     public TrainingSession(IReadOnlyDictionary<string, double>
22         hyperparameters)
23         => Hyperparameters = hyperparameters;
24
25     public void RecordEpoch(double loss) => LossHistory.Add(loss);
26 }
27 public sealed class InferenceSession
28 {
29     public Guid Id { get; } = Guid.NewGuid();
30     public Prompt Prompt { get; }
31     public ILMVector? Result { get; private set; }
32     public DateTime OccurredAt { get; } = DateTime.UtcNow;
33
34     public InferenceSession(Prompt prompt) => Prompt = prompt;
35     public void Complete(ILMVector result) => Result = result;
36 }

```

22.5 Aggregate and the ILM Aggregate Root

ILM is the **Aggregate Root** (Chapter 20, Section 20.4.4) for this entire Bounded Context: every Model, every Island, and every Inference passes through it, exactly as CoolantLoop was the single point of entry for every change to its pumps.

ILM.cs — the Aggregate Root

```

1 public sealed class ILM
2 {
3     private readonly List<Model> _models = new();
4
5     // The single, shared backing store for every Island: constructing
6     // the StaticVault and the Arbiter from this same dictionary is
7     // what guarantees AddIsland is immediately visible to Infer --
8     // two separately-constructed vaults would silently never agree.
9     private readonly Dictionary<string, ILMVector> _vaultEntries = new
10     ();
11     private readonly Arbiter _arbiter;
12
13     private readonly InferenceService _inferenceService;
14     private readonly TrainingService _trainingService;
15     private readonly IslandService _islandService;
16
17     public IReadOnlyList<Model> Models => _models;
18     public IEnumerable<Island> Islands =>
19         _vaultEntries.Select(kv => new Island { Hash = kv.Key, Answer =
20             kv.Value });

```

```

19 public IReadOnlyDictionary<string, double> Configuration { get; }
20
21 public ILM(ILMReadOnlyDictionary<string, double> configuration,
22     IVectorStore vectorStore)
23 {
24     Configuration = configuration;
25     var vault = new StaticVault(_vaultEntries);
26     _arbiter = new Arbiter(vault, vectorStore, new CausalityLock())
27     ;
28     _inferenceService = new InferenceService(_arbiter);
29     _trainingService = new TrainingService();
30     _islandService = new IslandService();
31 }
32
33 public ILMVector Infer(Prompt prompt) => _inferenceService.Infer(
34     prompt);
35
36 public void AddIsland(Island island) => _islandService.Register(
37     island, _vaultEntries);
38
39 public TrainingSession Train(ILMReadOnlyDictionary<string, double>
40     hyperparameters, Model model)
41     => _trainingService.Run(hyperparameters, model);
42 }

```

22.6 Domain Services

TrainingService, InferenceService, and IslandService hold the coordination logic that spans multiple Entities and Value Objects at once — exactly Chapter 20, Section 20.4.6’s definition, and, for two of the three, not new code at all.

AI Context Domain Services

```

1 public sealed class InferenceService
2 {
3     private readonly Arbiter _arbiter; // Chapter 12, Section 12.4.2
4     public InferenceService(Arbiter arbiter) => _arbiter = arbiter;
5
6     public ILMVector Infer(Prompt prompt)
7     {
8         var session = new InferenceSession(prompt);
9         // The Arbiter (unchanged since Chapter 12) tries Static Vault
10        // first, then a routed ILMVector, before ever reaching a
11        // stochastic generator.
12        var result = _arbiter.Answer(prompt.Normalized, candidate: null
13        );
14        session.Complete(result);
15        return result;
16    }
17 }
18
19 public sealed class TrainingService
20 {

```

```

20 public TrainingSession Run(IReadOnlyDictionary<string, double>
    hyperparameters, Model model)
21 {
22     var session = new TrainingSession(hyperparameters);
23     // Chapter 7's backpropagation loop runs here, epoch by epoch;
24     // each iteration calls session.RecordEpoch(loss), exactly as
25     // Chapter 7's XorBackprop.cs tracked loss, one call site
26     // formalised as a TrainingSession instead of a local variable.
27     return session;
28 }
29 }
30
31 public sealed class IslandService
32 {
33     public void Register(Island island, Dictionary<string, ILMVector>
        vaultEntries)
34     {
35         if (island.Answer.Determinism is not { IsDeterministic: true })
36             throw new InvalidOperationException(
37                 "Only a resolved deterministic ILMVector may back an
                    Island.");
38         vaultEntries[island.Hash] = island.Answer; // Chapter 12,
            Section 12.4.4's
39     } // update
40     discipline governs replacement.
}

```

InferenceService does not reimplement Chapter 12's Arbiter; it is a thin wrapper around it, which is exactly the design Chapter 20's Repository and Domain Service patterns are meant to produce: the AI Context's Ubiquitous Language (Infer, Prompt) sits on top of machinery that has not changed since Chapter 12, not a parallel reimplementaion of it.

22.7 Domain Events

AI Context Domain Events

```

1 public sealed record ModelTrained(Guid ModelId, double FinalLoss,
    DateTime OccurredAt);
2 public sealed record InferencePerformed(Guid SessionId, string
    PromptText, DateTime OccurredAt);
3 public sealed record IslandAdded(string Hash, DateTime OccurredAt);
4 public sealed record IslandTriggered(string Hash, DateTime OccurredAt);

```

IslandTriggered matters most: raised every time the Arbiter resolves a query through a Static Vault hit, it is exactly the audit record Chapter 19, Section 19.5 described — a subscriber (a logging service, a compliance dashboard) can listen for this one event type and build a complete, append-only audit trail of every deterministic answer the ILM has ever given, without the Arbiter itself needing to know that auditing exists.

22.8 Code Example: Building a Minimal ILM in C#

The Same Query, End to End, Through the Full DDD Model

Registering an Island for the exact query Chapter 12, Section 12.3.1 first hashed, then querying the ILM twice — once with that exact question, once with an unrelated one:

Program.cs

```

1 // A minimal in-memory IVectorStore (Chapter 12, Section 12.3.3) --
2 // enough to satisfy the ILM's dependency for this walkthrough.
3 public sealed class InMemoryVectorStore : IVectorStore
4 {
5     private readonly Dictionary<Guid, ILMVector> _byId = new();
6     public ILMVector? Get(Guid vectorId)
7         => _byId.TryGetValue(vectorId, out var v) ? v : null;
8 }
9
10 var ilm = new ILM(
11     configuration: new Dictionary<string, double>(),
12     vectorStore: new InMemoryVectorStore());
13
14 var pressureLimit = new ILMVector
15 {
16     VectorId = Guid.NewGuid(),
17     Embedding = new float[] { 120f }, // decodes to "120 bar"
18     Determinism = new DeterminismMetadata
19     {
20         IsDeterministic = true,
21         Version = "rev-C",
22         ValidFrom = new DateTime(2024, 1, 1)
23     }
24 };
25
26 var island = new Island
27 {
28     Hash = StaticVault.HashQuery(
29         "what is the maximum allowed discharge pressure for pump p
30         -101?"),
31     Answer = pressureLimit
32 };
33 ilm.AddIsland(island); // raises IslandAdded
34
35 var hit = ilm.Infer(new Prompt {
36     Text = "What is the maximum allowed discharge pressure for pump
37     P-101?" });
38 // hit.Determinism!.IsDeterministic == true; raises IslandTriggered
39
40 var miss = ilm.Infer(new Prompt { Text = "What do you recommend
41     right now?" });
42 // no Static Vault entry matches; the Arbiter falls through toward
43 // the stochastic generator, exactly as Chapter 12's own worked
44 // example described for this second kind of question.

```

The first Infer call resolves through the identical 64-character hash Chap-

ter 12 computed by hand (f20ab319ace...bcf03e574); the second does not, and is handled by the open-ended path Chapter 12, Section 12.5 already walked through. Nothing about the DDD structure in this chapter changed either outcome — it gave the existing behaviour names (ILM, Prompt, Island) a plant engineer and a software engineer can both read the same way.

Honest Boundary

This is a teaching model; production systems require much more. The classes in this chapter omit concurrency control (many simultaneous Infer calls against one ILM), persistence (every Repository here is implicitly in-memory), authentication and authorisation on who may call AddIsland, and the operational tooling (metrics, retries, circuit breakers) any production inference service needs. The domain model this chapter built is the correct *starting point* for that production system — Chapter 20's entire argument was that getting the Entities, Value Objects, and Aggregate boundaries right first makes everything built on top of them easier to get right too — but the model alone is not the system, and no reader should deploy the code in this chapter directly against a critical process without everything this paragraph just named.

23. Integration and Application Architecture

Chapter 20, Section 20.4.2 established that a Bounded Context's vocabulary should not leak across its own boundary. A real application, though, needs its contexts to *cooperate* without merging their languages, and this book's own three contexts form a natural chain: the **Quantum Context** (Chapter 21) provides the theoretical foundations — nuclear structure and decay are, at bottom, quantum-mechanical phenomena — that the **Nuclear Context** (isotopes, decay chains, fission fragment yields; built in full in Volume I, *From Grid to Core*) builds on; the Nuclear Context in turn supplies the **Reactor Control Context** with exactly the physical quantities — isotopic inventories, decay heat, delayed-neutron fractions — a control system must track; and the Reactor Control Context supplies the **AI Context** (Chapter 22) with the raw material, Chapter 10 through 19's entire subject, that learning and decision support are built from. Each arrow is a one-way data dependency, not a shared vocabulary: the AI Context never needs to know what a delayed-neutron fraction *is* physically, only that a `DelayedNeutronDataRepository` can supply it as training data, in exactly the AI Context's own Tensor terms.

23.2 Use of Repositories

Section 23.1's chain of dependencies is implemented entirely through **Repositories** (Chapter 20, Section 20.4.5): each context exposes its data to its neighbour through a narrow, read-oriented interface, never through a shared database schema or a leaked internal class.

Repositories, one per context boundary

```

1 // Nuclear Context repositories -- consumed by, but not owned by,
2 // the AI Context.
3 public interface IIsotopeRepository
4 {
5     IsotopeData? GetByNuclide(string nuclide);
6     IEnumerable<IsotopeData> GetDecayChain(string rootNuclide);
7 }
8
9 public interface IDelayedNeutronDataRepository
10 {
11     DelayedNeutronGroup[] GetGroupsFor(string fissileIsotope);
12 }
13
14 // AI Context repositories (Chapter 20, Section 20.4.5 already named
```

```

15 // IVectorStore as exactly this pattern; ModelRepository and
16 // IslandRepository complete the set for the other two Entities and
17 // Value Objects Chapter 22 introduced).
18 public interface IModelRepository
19 {
20     Model? GetById(Guid modelId);
21     void Save(Model model);
22 }
23
24 public interface IIslandRepository
25 {
26     Island? GetByHash(string hash);
27     void Save(Island island);
28     IEnumerable<Island> GetAll();
29 }

```

IIslandRepository is a direct refinement of Chapter 22, Section 22.5's `_vaultEntries` dictionary: the Aggregate Root there used a plain in-memory dictionary because Chapter 22 was teaching the domain model in isolation; a real application injects an IIslandRepository implementation instead, so persistence can change (memory, file, database) without ILM or any Domain Service noticing.

23.3 Application Creation

23.3.1 Application Layer: Coordination of Domain Services

The **Application Layer** orchestrates a use case by calling Domain Services and Repositories in sequence; it contains no domain logic of its own, only coordination.

TrainIlmApplicationService.cs

```

1 public sealed class TrainIlmApplicationService
2 {
3     private readonly IDelayedNeutronDataRepository _neutronData;
4     private readonly IModelRepository _models;
5     private readonly TrainingService _trainingService;
6
7     public TrainIlmApplicationService(
8         IDelayedNeutronDataRepository neutronData,
9         IModelRepository models,
10        TrainingService trainingService)
11     {
12         _neutronData = neutronData;
13         _models = models;
14         _trainingService = trainingService;
15     }
16
17     public Guid RetrainOnLatestReactorData(Guid modelId, string
18        fissileIsotope)
19     {
20         var model = _models.GetById(modelId)
21             ?? throw new InvalidOperationException("Unknown model.");
22         var groups = _neutronData.GetGroupsFor(fissileIsotope);
23         // Reactor Control Context data (Section 23.1) becomes training

```

```

23      // input here -- the Application Layer is the only place that
24      // is allowed to know about both the Nuclear and AI contexts
25      // at once.
26      var hyperparameters = new Dictionary<string, double> { [ "
27          learningRate" ] = 0.001 };
28      var session = _trainingService.Run(hyperparameters, model);
29      _models.Save(model);
30      return session.Id;
31  }

```

23.3.2 Infrastructure Layer: Repository Implementation, Logging, Monitoring

The **Infrastructure Layer** provides concrete implementations of the interfaces Section 23.2 defined — a `SqlIsotopeRepository`, a file-backed `FileModelRepository` — plus cross-cutting concerns (logging, metrics, retries) that no Domain Service or Aggregate should need to reference directly. Chapter 19, Section 19.5's audit trail lives here: a concrete `IDomainEventListener` implementation (Section 23.4) that writes every `IslandTriggered` event to a durable log is Infrastructure, not domain logic.

23.3.3 Presentation Layer: CLI or Web API

The **Presentation Layer** is the thinnest layer of all: it translates an external request (an HTTP call, a console command) into a call against the Application Layer, and translates the result back — Section 23.5 gives this concretely.

23.4 Design Patterns Used

Four Patterns, Already at Work in This Book

Repository (Section 23.2): abstracts storage. **Factory**: encapsulates the logic of constructing a complex object correctly. **Observer / Event Listener**: lets independent parts of the system react to a Domain Event (Chapter 22, Section 22.7) without the event's source knowing who is listening. **Strategy**: makes an algorithm swappable behind a common interface.

Factory, Observer, and Strategy, each in one example

```

1  // Factory: constructs a Model with correctly-shaped, zero-initialised
2  // weights for a given architecture -- the one place that logic lives.
3  public static class ModelFactory
4  {
5      public static Model CreateUntrained(IReadOnlyList<Layer>
6          architecture)
7      {
8          var weights = architecture
9              .Select(l => new Tensor { Data = new float[l.Shape.
10                  Aggregate(1, (a, b) => a * b)], Shape = l.Shape })
11              .ToList();

```

```

10     return new Model(Guid.NewGuid(), architecture, weights);
11 }
12 }
13
14 // Observer: any number of listeners can react to a domain event
15 // without IslandService knowing they exist.
16 public interface IDomainEventListener<TEvent> { void Handle(TEvent
    domainEvent); }
17
18 public sealed class AuditLogListener : IDomainEventListener<
    IslandTriggered>
19 {
20     public void Handle(IslandTriggered domainEvent)
21         => Console.WriteLine($"[AUDIT] Island {domainEvent.Hash}
            triggered at {domainEvent.OccurredAt:0}");
22 }
23
24 // Strategy: TrainingService (Chapter 22) can be configured with any
25 // optimiser (Chapter 7, Section 7.6) behind one shared interface.
26 public interface IOptimizationStrategy
27 {
28     void Step(Model model, Tensor gradient);
29 }
30
31 public sealed class AdamStrategy : IOptimizationStrategy
32 {
33     public void Step(Model model, Tensor gradient) { /* Chapter 7,
        Section 7.6.2 */ }
34 }

```

23.5 Creating a Console or API for Interacting with the Model

A Minimal Web API, Presentation Layer

```

1 var builder = WebApplication.CreateBuilder(args);
2 builder.Services.AddSingleton<ILM>(sp => new ILM(
3     configuration: new Dictionary<string, double>(),
4     vectorStore: new InMemoryVectorStore()));
5 var app = builder.Build();
6
7 app.MapPost("/infer", (PromptRequest request, ILM ilm) =>
8 {
9     var result = ilm.Infer(new Prompt { Text = request.Text });
10    var isDeterministic = result.Determinism is { IsDeterministic: true
11        };
12    return Results.Ok(new { isDeterministic, embedding = result.
13        Embedding });
14 });
15
16 app.Run();
17
18 public sealed record PromptRequest(string Text);

```

Every safety property Chapter 12 built is still enforced at this outermost layer, unchanged: the endpoint has no way to bypass the Arbiter, because ILM.Infer is the only path to an answer this API exposes at all.

23.6 Testing

Three Levels of Test

Unit tests exercise one class in isolation. **Integration tests** exercise several classes wired together, with real (often in-memory) dependencies. **Acceptance tests** exercise a complete use case from the outside, in the Ubiquitous Language (Chapter 20, Section 20.4.1) a domain expert would use to describe it.

Three Tests, One per Level, in xUnit

```

1 public class IslandServiceTests // Unit
2 {
3     [Fact]
4     public void Register_Throws_WhenAnswerIsNotDeterministic()
5     {
6         var service = new IslandService();
7         var nonDeterministic = new ILMVector
8         {
9             VectorId = Guid.NewGuid(), Embedding = new float[] { 1f },
10            Determinism = null
11        };
12         var island = new Island { Hash = "abc", Answer =
13             nonDeterministic };
14
15         Assert.Throws<InvalidOperationException>(
16             () => service.Register(island, new Dictionary<string,
17                 ILMVector>()));
18     }
19 }

```

```

16 }
17
18 public class IlmInferenceIntegrationTests // Integration
19 {
20     [Fact]
21     public void Infer_ReturnsRegisteredIsland_ForMatchingPrompt()
22     {
23         var ilm = new ILM(new Dictionary<string, double>(), new
24             InMemoryVectorStore());
25         var answer = new ILMVector
26         {
27             VectorId = Guid.NewGuid(), Embedding = new float[] { 120f
28             },
29             Determinism = new DeterminismMetadata
30             {
31                 IsDeterministic = true, Version = "rev-A", ValidFrom =
32                 DateTime.UtcNow
33             }
34         };
35         ilm.AddIsland(new Island { Hash = StaticVault.HashQuery("max
36             pressure?"), Answer = answer });
37
38         var result = ilm.Infer(new Prompt { Text = "max pressure?" });
39         Assert.True(result.Determinism is { IsDeterministic: true });
40     }
41 }
42
43 // Acceptance ("Given a registered safety limit, when an operator asks
44 // for it twice, then both answers are identical" -- Chapter 10,
45 // Section 10.2.3's non-determinism failure, verified as a test.)
46 public class RepeatedQueryAcceptanceTests
47 {
48     [Fact]
49     public void SameQuestionTwice_ReturnsIdenticalAnswer()
50     {
51         // BuildIlmWithPressureLimitIsland() mirrors the setup already
52         // shown in full in IlmInferenceIntegrationTests, above.
53         var ilm = BuildIlmWithPressureLimitIsland();
54         var first = ilm.Infer(new Prompt { Text = "max pressure?" });
55         var second = ilm.Infer(new Prompt { Text = "max pressure?" });
56         Assert.Equal(first.Embedding, second.Embedding);
57     }
58 }

```

23.7 Future Extensions: Digital Twin of the Reactor

Chapter 17, Section 17.5's digital twin is the natural next Application Layer service for this architecture: connecting to real sensor data (a new Repository implementation, nothing else), predictive analytics (Chapter 15's models, trained through exactly Section 23.3.1's `TrainIlmApplicationService` pattern), ML.NET-based training (an alternative `IOptimizationStrategy` implementation, Section 23.4), and real-time dashboards (a Presentation Layer consumer of the same Domain Events Section 23.4's `AuditLogListener` already showed how to subscribe to). None of these


extensions require touching a single class this chapter or Chapter 22 already built; every one of them is new code arriving at an existing seam, which is the entire practical payoff of getting the Bounded Context boundaries right in the first place.

Honest Boundary

This is a teaching-scale architecture; production requires much more.

Every pattern in this chapter is genuine and production-grade in kind, but the specific code shown omits transaction management across repositories, authentication and authorisation on every endpoint (Section 23.5's API accepts any request), distributed tracing across the Application and Infrastructure layers, and the deployment, scaling, and disaster-recovery concerns any real critical-infrastructure system carries. Chapter 20 through this chapter have built the right *shape* for that system — correctly bounded contexts, correctly assigned responsibilities — and getting that shape right first is what makes every one of these missing pieces addable later without a rewrite; but a correct shape, on its own, is not yet a deployable system, and no reader should treat it as one.

TOWARDS ARTIFICIAL CERTAINTY



TOWARDS ARTIFICIAL CERTAINTY

24. Deterministic Islands — Deep Dive

24.1 What You Already Know

You already have the complete practical picture: Chapter 12 built Static Vault, Frozen Snapshot, and Intra-Vector Routing from the ground up, with worked code and worked examples; Chapter 22 gave the same mechanism its formal DDD structure. This chapter does not repeat either. It returns to the same three mechanisms one level deeper, supplying the mathematical foundations — exactly how strong the hashing guarantee actually is, exactly what “reproducible” means at the level of individual floating-point bits, and a formal definition precise enough to state as a theorem — that Chapter 12 was entitled to use without proving.

24.2 Mathematical Foundations

24.2.1 Hashing: SHA-256 and Collision Resistance

Chapter 12, Section 12.3.1 relied on SHA-256 producing “the same input byte sequence always producing the same 256-bit output” without asking how safe it is to treat two different queries as certainly different whenever their hashes differ. A **hash collision** — two distinct inputs producing the identical hash — is the one event that would let a Static Vault return the wrong entry for a query, and the **birthday bound** quantifies exactly how unlikely that is.

How Safe Is a 256-bit Hash, Really?

For an n -bit hash function assumed to behave like a random function, the number of distinct inputs that must be hashed before a collision becomes as likely as not is approximately

$$N_{50\%} \approx 1.1774\sqrt{2^n}.$$

For SHA-256 ($n = 256$):

$$N_{50\%} \approx 1.1774\sqrt{2^{256}} \approx 4.01 \times 10^{38}.$$

Even at an absurdly fast, physically implausible rate of 10^{18} hashes per second (roughly a billion times faster than any real Static Vault would ever be queried), reaching this many hashes would take

$$\frac{4.01 \times 10^{38}}{10^{18} \text{ s}^{-1}} \approx 4.01 \times 10^{20} \text{ seconds} \approx 1.27 \times 10^{13} \text{ years},$$

roughly **920 times the current age of the universe**. This is not merely “very unlikely”; it is the specific, standard justification, used throughout cryptography, for treating a 256-bit hash’s output as an unforgeable, effectively-unique

identifier for its input — the Static Vault’s entire correctness guarantee rests on this number being as large as it is, not on an informal appeal to “hashes are usually fine.”

24.2.2 Reproducibility: Seeds, Determinism, and Bitwise Accuracy

Chapter 12, Section 12.3.2’s Frozen Snapshot claimed that single-threaded CPU execution removes floating-point non-determinism that GPU parallelism introduces. The underlying fact is associativity failure, made concrete:

The Same Sum, Two Different Answers

Let $a = 10^{16}$, $b = -10^{16}$, $c = 1.0$, all exactly representable in IEEE 754 double precision. Mathematically, $(a+b)+c = a+(b+c) = 1$. Computed in floating-point arithmetic:

$$(a + b) + c = (10^{16} - 10^{16}) + 1.0 = 0.0 + 1.0 = \mathbf{1.0},$$

$$a + (b + c) = 10^{16} + (-10^{16} + 1.0) = 10^{16} + (-10^{16}) = \mathbf{0.0}.$$

$b + c$ rounds to -10^{16} because 1.0 is far below the precision double-precision arithmetic can represent at that magnitude (machine epsilon at 10^{16} is larger than 1) — the addition’s *order* changed the final answer from 1.0 to 0.0, not by an approximation error, but exactly, reproducibly, and inevitably given IEEE 754 rounding rules. GPU inference sums many partial products in an order that depends on which thread finishes first, which can vary run to run; single-threaded CPU execution fixes one specific summation order, and this worked example is the entire, concrete reason why that fix is necessary for bitwise reproducibility, not merely sufficient for it.

24.2.3 Formal Definition of a Deterministic Island

Deterministic Island, Defined

A **Deterministic Island** is a pair (H, A) where $H : \mathcal{Q} \rightarrow \{0, 1\}^{256}$ is a fixed hash function over normalised queries, and A is a *partial function* $A : \{0, 1\}^{256} \rightarrow V$ (defined only on the finite set of hashes currently registered) mapping a hash to an ILMVector with `Determinism.IsDeterministic = true`. The composition $A \circ H$ is the Static Vault’s lookup as a whole.

Theorem. For any two queries $q_1, q_2 \in \mathcal{Q}$ with $H(q_1) = H(q_2)$ both in A ’s domain, $(A \circ H)(q_1) = (A \circ H)(q_2)$. **Proof.** A is a function, so equal inputs to A produce equal outputs; $H(q_1) = H(q_2)$ by hypothesis, so $A(H(q_1)) = A(H(q_2))$. ▀

This is a genuinely trivial proof, and that triviality is the point: Chapter 12’s entire determinism guarantee reduces, once stated this formally, to nothing more than the ordinary mathematical fact that *function composition is well-defined*. All of the engineering difficulty this book has spent three chapters on — making sure H really is a pure function of its input (Section 24.2.1) and that A really is unambiguous at query time (Section 24.4) — exists entirely to make this one-line proof’s hypotheses true in practice, not to strengthen the proof itself.

24.3 Implementation, Revisited Through This Lens

Chapter 12, Sections 12.3.1-12.3.3 already gave complete, working C# for Static Vault, Frozen Snapshot, and Intra-Vector Routing; nothing in this section changes that code. What Section 24.2 adds is a way to reason about its *cost*: a Static Vault lookup is $O(1)$ expected time (an ordinary hash-table lookup); VectorResolver’s pointer-chasing is $O(\text{maxHops})$, a small constant bound by construction (Chapter 12, Section 12.3.3); and a Frozen Snapshot’s single-threaded CPU execution, while asymptotically no different from GPU inference, typically runs one to two orders of magnitude slower in wall-clock time for a large Transformer, a concrete latency cost that is the direct price of Section 24.2.2’s bitwise reproducibility guarantee — paid, deliberately, only for the narrow class of queries where Chapter 10, Section 10.3.3 said the guarantee is actually required.

24.4 Management

24.4.1 Updating and Extending Islands

Chapter 12, Section 12.4.4 already fixed the append-only update discipline (close the old ValidTo, open a new VectorId). Formally, this is an invariant on the *version history* of a single hash: for a fixed $h \in \{0, 1\}^{256}$, the ValidFrom/ValidTo intervals of every ILMVector ever registered under h must partition time without gap or overlap, so that “what was the answer to this query on this date” always has exactly one well-defined answer.

Extending an island — this chapter’s genuinely new idea — means deriving a new deterministic answer as a pure function of *other* islands’ answers, rather than registering an independent fact.

A Derived Island, and Why It Stays Deterministic

Three redundant pressure sensors each have an individually registered maximum safe reading, as three separate Islands: 45.0, 42.5, and 44.0 bar. A fourth, *derived* Island — the system’s overall safe operating limit, defined as the minimum across all three — computes

$$\min(45.0, 42.5, 44.0) = \mathbf{42.5 \text{ bar.}}$$

This derived value is itself deterministic, by the same one-line argument as Section 24.2.3’s theorem: $\min$ is an ordinary mathematical function, and a function applied to deterministic inputs produces a deterministic output. The engineering discipline this adds is purely organisational, not mathematical: a derived Island’s registered answer must be *recomputed and re-registered* (Section 24.4.1’s update discipline, applied automatically rather than manually) whenever any of its inputs changes, or it silently becomes stale — a new, composite failure mode this chapter’s closing Honest Boundary returns to.

24.4.2 Conflict Resolution, Recalled

Chapter 12, Section 12.4.3’s fixed precedence order (Static Vault, then a routed vector, then Frozen Snapshot) and Chapter 12, Section 12.3.3’s Arbiter, Causality Lock, and Human-in-the-Loop mechanisms are unchanged here; this deep dive adds no new resolution rule, because Section 24.2.3’s formal definition already shows the

existing rule is the entire content required: exactly one mechanism’s answer is ever used, by construction, so “conflict” can only ever mean two mechanisms disagreeing about what the ground truth is, never an ambiguity in which one governs.

24.5 Code Examples in C#

A Derived (Composite) Island

```

1 public static class DerivedIslandFactory<T> where T : Island {
2     {
3         // Builds a new deterministic Island as a pure function of
4         // already-registered Islands -- Section 24.4.1's worked example,
5         // as code. Because Func<...> here is a pure numeric function,
6         // the result inherits IsDeterministic = true directly.
7         public static Island DeriveMinimum(string derivedHash, params
8             Island[] inputs)
9         {
10             if (inputs.Any(i => i.Answer.Determinism is not {
11                 IsDeterministic: true }))
12                 throw new InvalidOperationException(
13                     "A derived Island requires every input to already be
14                     deterministic.");
15
16             float minValue = inputs.Min(i => i.Answer.Embedding[0]);
17             var answer = new ILMVector
18             {
19                 VectorId = Guid.NewGuid(),
20                 Embedding = new float[] { minValue },
21                 Determinism = new DeterminismMetadata
22                 {
23                     IsDeterministic = true,
24                     Version = "derived-v1",
25                     ValidFrom = DateTime.UtcNow
26                 }
27             };
28             return new Island { Hash = derivedHash, Answer = answer };
29         }
30     }
31 }

```

24.6 Case Study: A Deterministic Island for Reactor Safety

A reactor’s **SCRAM setpoint** — the exact core temperature at which an automatic emergency shutdown is triggered — is among the highest-consequence single facts an ILM could ever be asked about, and a natural, concrete case for everything this chapter has formalised. Registered once as an Island, its hash and answer sit in the Static Vault exactly as Chapter 12, Section 12.5’s pressure-limit example did; every query about it, from any operator, on any shift, for as long as it remains registered, resolves through Section 24.2.3’s *A ◦ H* composition, with the birthday-bound guarantee of Section 24.2.1 standing behind the promise that no other query could ever accidentally collide with it. When the setpoint is revised — following exactly the regulatory change-control process Chapter 19, Section 19.4’s skeleton safety case already cited as evidence — Section 24.4.1’s versioning discipline ensures a

query about last year's setpoint, perhaps raised during an incident investigation, still resolves correctly to last year's value, while today's query resolves to today's. No part of this case study required anything beyond the mechanisms Chapter 12 already built; what this chapter's deep dive adds is the confidence, now proven rather than asserted, that those mechanisms are mathematically sound at exactly the point — a life-safety setpoint — where soundness matters most.

Honest Boundary

Islands are not a panacea; they require careful design and maintenance.

Chapter 12's own closing Honest Boundary already said this once; this deep dive adds two specific, previously unnamed failure modes. **Stale derived islands** (Section 24.4.1): a composite Island computed from other Islands is only as current as the last time it was recomputed, and nothing in Section 24.5's code detects, on its own, that an input has changed without a recomputation being triggered — a purely organisational discipline this chapter cannot make automatic. **Island sprawl**: as an ILM's registered fact base grows into the thousands, the human review burden Chapter 12, Section 12.4.4 requires for every update grows with it, and a static-vault registry too large for any one reviewer to hold a mental model of is a governance failure waiting to happen, however mathematically sound each individual entry remains in isolation. Neither failure mode is solved by better mathematics; both are solved, if at all, by the same disciplined, bounded, human process Chapter 14 and Chapter 19 already argued for.

25. Governance and Regulation

25.1 What You Already Know

You already carry this chapter’s raw material: Chapter 14 built the ethics and governance argument in full, and Chapter 19 grounded it in three live regulatory frameworks. This chapter does not restate either. It asks the one question a governance chapter placed this late in the book is finally positioned to answer well: *having built the whole architecture, chapter by chapter, does it actually satisfy the frameworks Chapter 19 described?* The answer, mapped out explicitly below, is mostly yes, by design rather than by retrofit — and the gaps that remain are named honestly, not smoothed over.

25.2 Regulatory Frameworks, Mapped to This Book's Architecture

25.2.1 EU AI Act

Chapter 14, Section 14.6 already covered the Act’s risk tiers and current (2026) phased timeline. What follows is new: a direct mapping from the Act’s substantive obligations to the specific mechanism in this book that satisfies each one.

EU AI Act obligation	Book mechanism	Where
Risk management system	Evaluation discipline, precision/recall/AUPRC by task risk	Chapter 13
Data governance	Training-data sourcing, preprocessing, imbalance handling	Chapter 10, Section 10.5
Technical documentation	The ILM’s own DDD model as living documentation	Chapter 22
Record-keeping	IslandTriggered/InferencePerformed events, logged	Chapter 19, §19.5; Chapter 22, §22.7
Human oversight (Article 14)	Human-in-the-loop / human-on-the-loop	Chapter 14, Section 14.7
Accuracy, robustness, cybersecurity	Adversarial robustness, distribution shift	Chapter 13, Section 13.4

25.2.2 NIST AI Risk Management Framework

The NIST AI RMF’s four functions — Govern, Map, Measure, Manage — map onto four separate parts of this book with a cleanliness that was not an accident of order-

ing: Chapter 14 **governs** (policy, accountability, oversight structure); Chapter 10 **maps** (defines exactly what the ILM is and is not scoped to do); Chapter 13 **measures** (the entire evaluation and verification discipline); and Chapter 12 **manages** risk directly, in running code, through the Deterministic Islands that bound where stochastic behaviour is and is not permitted to reach.

25.2.3 Sector-Specific Regulations

Chapter 19, Section 19.2's tour of IAEA, EASA/FAA, and FDA guidance stands as written; this chapter adds one synthesis point Chapter 19 was not yet positioned to make. All three regulators, independently, converge on **lifecycle-based oversight over one-time approval** (Chapter 19, Section 19.3) — and Chapter 12, Section 12.4.4's append-only update discipline, together with Chapter 19, Section 19.5's audit trail, is precisely a lifecycle oversight mechanism, already built, not a compliance feature that would need to be added afterward to satisfy any of the three.

25.3 Responsibility and Accountability, Made Queryable

Chapter 14, Section 14.8 stated the principle: every decision path terminates at a named human. Chapter 22's Domain Events make this more than a policy statement — they make it a property the system can answer questions about directly.

Accountability as a Query, Not Only a Policy

```

1 public sealed class AccountabilityQueryService<T>
2     where T : IIslandTriggered
3 {
4     private readonly IReadOnlyList<T> _triggers;
5     private readonly IReadOnlyList<IslandAdded> _additions;
6
7     public AccountabilityQueryService(
8         IReadOnlyList<T> triggers, IReadOnlyList<
9             IslandAdded> additions)
10    {
11        _triggers = triggers; _additions = additions;
12    }
13
14    // "Which Island answered this query, and who approved it?"
15    // -- Chapter 12, Section 12.4.4 requires every addition to carry
16    // a human sign-off; a real IslandAdded event would include the
17    // approver's identity alongside Hash and OccurredAt.
18    public IEnumerable<(string Hash, DateTime Triggered)>
19        TriggersInRange(DateTime from, DateTime to)
20        => _triggers.Where(t => t.OccurredAt >= from && t.OccurredAt <=
21            to)
22            .Select(t => (t.Hash, t.OccurredAt));
23 }

```

This is Chapter 19, Section 19.5's auditing discipline, made concrete enough to run: "who is accountable for every deterministic answer given last quarter" is not a question that requires digging through logs after the fact; it is an ordinary query over data the architecture was already producing.

25.4 Ethical Issues, Extended: Autonomous Decision-Making

Chapter 14, Sections 14.2 through 14.5 covered bias, transparency, security, and privacy in full technical depth; one further question those sections did not yet name directly, **autonomous decision-making**, deserves the synthesis this book is now positioned to give. Chapter 19, Section 19.2's EASA Level 1/2/3 framework (human augmentation, human-AI cooperation, advanced automation) turns out to classify this book's own architecture precisely:

Where This Book's Architecture Sits, Deliberately

For **Deterministic Island-governed facts** (Chapter 12): the system does not decide anything; it retrieves a pre-verified answer. This is Level 1, human augmentation, by construction — no higher level of autonomy is even applicable to a lookup. For **Shield-gated control recommendations** (Chapter 17, Section 17.4): the model proposes, a non-learned boundary constrains, and a human retains override authority. This is Level 2, human-AI cooperation — never Level 3. For **open-ended generation** (RAG-grounded answers, decision support, Chapter 18): the model reasons flexibly, but the output is explicitly a recommendation, never an executed action. This, too, is Level 2 at most.

This book has never proposed, anywhere in twenty-four preceding chapters, an architecture that reaches EASA's Level 3. That is not an oversight this chapter is now correcting; it is the position the whole book has argued for from Chapter 9, Section 9.4.3 onward, now named in the regulatory vocabulary Chapter 19 introduced.

Honest Boundary

Governance is an ongoing process, not a one-time certification. Chapter 13 and Chapter 19 each made a version of this point already; the concrete extension worth adding here is a cadence, not just a warning: bias metrics (Chapter 14, Section 14.2) checked on a recurring schedule, not once at deployment; certification renewal (Chapter 19, Section 19.3) tracked against each regulator's own lifecycle requirements, not assumed permanent; and the audit queries Section 25.3 just made concrete, run routinely, not only when an incident forces the question. None of Section 25.2's mapping table is a claim that this book's architecture is compliant, once and permanently, with any framework it lists; it is a claim that the architecture gives an organisation the right *starting* evidence for an ongoing compliance relationship it must still actively maintain.

26. The Future — Artificial Certainty

26.1 What You Already Know

You already know that engineering trust in a new technology has always followed the same arc: a period of impressive but unpredictable capability, followed by the slow, unglamorous work of building the specific guarantees that let a capability be relied upon in consequential settings — steam power before the boiler code, electricity before earthing standards, aviation before airworthiness certification. Nothing about AI exempts it from this arc, and nothing in the twenty-five chapters before this one has argued otherwise. What follows is not a prediction that this arc will complete quickly, or easily, or without setbacks; it is a statement of where this book believes the arc is headed, built from mechanisms already demonstrated working, in running code, throughout this book.

26.2 The Vision: AI That Is Not Merely Intelligent, but Reliable

Every chapter of Part II built intelligence: the capacity to learn patterns, generalise, and generate fluent, useful output. Intelligence alone, Chapter 10 argued at length, is not what a nuclear plant, a grid operator, or a hospital needs from a system whose answer might be acted on. They need **reliability**: the guarantee that a specific, safety-relevant class of question gets the same, correct answer every time, under a certifiable process, with a named human accountable for every path through the system. This book's entire argument, restated in one sentence, is that intelligence and reliability are not the same property, do not require the same architecture, and — this is the constructive claim, not merely the diagnosis — *can be combined in one system without sacrificing either*.

26.3 Deterministic Islands as the Foundation of Trustworthy AI

Chapter 12 built this combination; Chapter 22 gave it formal structure; Chapter 24 proved its mathematical soundness. This section adds nothing technical to any of the three. It states the claim plainly, once, at the scale this closing chapter earns: a **Deterministic Island** is not a workaround bolted onto a fundamentally unreliable system to patch its worst failure mode. It is the specific, minimal, mathematically sound mechanism by which a system built from stochastic components can still make an absolute guarantee about a bounded, identifiable subset of its behaviour

— and every chapter since Chapter 9’s reinforcement-learning warning has been, in one form or another, an application of exactly this one idea.

26.4 The Neural Constitution: Rules, Constraints, Values

Every Deterministic Island, every Shield, every Arbiter precedence rule this book has built is, individually, a specific technical mechanism solving a specific problem. Together, they form something worth naming as a single artifact: a **Neural Constitution** — the complete, explicit set of rules an ILM must never violate, regardless of what its stochastic components learn or how its weights change under retraining.

NeuralConstitution.cs — Naming What This Book Already Built

```

1 public sealed record ConstitutionalRule(string Name, string EnforcedBy)
2     ;
3 public static class NeuralConstitution
4 {
5     public static readonly IReadOnlyList<ConstitutionalRule> CoreRules
6         = new[]
7     {
8         new ConstitutionalRule(
9             "A resolved deterministic fact is never overridden by a
10              generated one.",
11             "Arbiter's fixed precedence order, Chapter 12, Section
12              12.4.3"),
13         new ConstitutionalRule(
14             "No control action reaches an actuator outside its
15              certified safe range.",
16             "Shield, Chapter 17, Section 17.4"),
17         new ConstitutionalRule(
18             "A Causality Lock, once engaged, never silently falls back
19              to a stochastic answer.",
20             "GuardedResolver, Chapter 12, Section 12.3.3"),
21         new ConstitutionalRule(
22             "Every registered fact's provenance and approval are
23              permanently recorded.",
24             "Append-only update discipline, Chapter 12, Section 12.4.4"),
25         new ConstitutionalRule(
26             "No fully autonomous action is taken without a human
27              retaining override authority.",
28             "Human-on-the-loop, Chapter 14, Section 14.7; Chapter 25,
29              Section 25.4"),
30     };
31 }

```

Nothing in this class is new code; every rule’s `EnforcedBy` field points to a mechanism already built, tested, and formally argued for in an earlier chapter. What the `NeuralConstitution` adds is not enforcement — the enforcement already exists — but **legibility**: a single artifact a reviewer, a regulator, or an incoming engineer can read to see the complete list of guarantees the system makes, each traceable to exactly the chapter that built it.

26.5 Ethical AI

Chapter 14 and Chapter 25 built this book's ethical position in full technical depth; this section states only its trajectory. An ethically deployed ILM is not one that has passed a one-time review (Chapter 13, Chapter 19's shared warning against exactly this framing) but one whose **Neural Constitution** (Section 26.4) is itself subject to the same governance discipline as everything else in this book: reviewed on a cadence (Chapter 25, Section 25.4's closing point), extended only through the same append-only, human-approved process Chapter 12 already established for any other fact, and never silently amended by a retraining run that was never told the Constitution existed as a boundary on what it was permitted to change.

26.6 The Future of Critical Infrastructure

The trends already visible as this book is written — grid modernisation under growing renewable and storage penetration, small modular reactors with substantially more automated operating concepts than the plants Volume I was written against, and a general industry-wide push toward predictive rather than reactive operations (Chapter 15 through 18, in full) — all increase, not decrease, the pressure toward more AI in the operating loop. This book's position is not that this pressure should be resisted; it is that every one of these systems should be built the way Part III through Part VI built this one: stochastic capability wrapped in deterministic guarantees exactly where guarantees are required, never the reverse, and never assumed rather than engineered.

26.7 Quantum Neural Networks with Deterministic Islands

26.7.1 Introduction to Quantum Neural Networks (QNNs)

A **Quantum Neural Network** replaces some or all of a classical network's layers (Chapter 8) with parameterised quantum circuits (Chapter 21) operating on qubits instead of classical activations. Its appeal is the exponential state-space capacity Chapter 21, Section 21.2 already built: n qubits represent a state in a 2^n -dimensional space, offering, in principle, a representational capacity no classical layer of comparable size can match. The crucial, widely misunderstood limit sits directly beside this appeal: **quantum parallelism** lets a circuit evaluate a function across a superposition of every input simultaneously, but Chapter 21, Section 21.4's measurement collapse means only *one* outcome can ever be read out, chosen probabilistically according to the Born rule. A QNN gains nothing automatically from evaluating "all inputs at once" unless its circuit is specifically designed (interference between amplitudes, not brute enumeration) to concentrate measurement probability on the answer that matters — exactly what makes designing genuinely useful quantum algorithms hard, and why "quantum computers try every answer simultaneously" is a popular but seriously misleading simplification.

26.7.2 The Challenge of Noise and Decoherence

A qubit's superposition survives only as long as it is isolated from uncontrolled interaction with its environment; any such interaction — **decoherence** — degrades

the superposition toward an ordinary classical mixture, destroying exactly the interference effects Section 26.7.1 said a useful QNN depends on. As of 2026, current hardware (superconducting, trapped-ion, and photonic platforms alike) remains in what the field calls the **NISQ** era (Noisy Intermediate-Scale Quantum): roughly 50 to 1,000-plus physical qubits, two-qubit gate error rates around 0.1-1%, and no full, always-on error correction. **Quantum error correction** addresses this by encoding one robust **logical qubit** across many physical qubits (a distance-7 surface code, for example, uses 49 physical qubits per logical qubit); the field's most significant recent milestone, Google's Willow processor demonstrating in December 2024 that logical error rates *decrease* as more physical qubits are added — "below threshold" operation — is precisely what turns error correction from a theoretical requirement into a practical engineering path, though multiple independent teams (QuEra reporting 96 logical qubits from 448 physical qubits by April 2026 among them) had, by the time of this writing, still only reached the tens-to-hundreds-of-logical-qubits scale, well short of the millions of logical qubits current estimates (Gidney and Ekera, 2021) suggest breaking RSA-2048 via Shor's algorithm would require.

26.7.3 Deterministic Islands in the Quantum Regime

Applying Chapter 12's principles to a quantum system runs immediately into a fact Chapter 21 already established: **measurement collapse** means a quantum state cannot simply be read out and cached the way an `ILMVector` is, and the **no-cloning theorem** means an unknown quantum state cannot even be copied to check it non-destructively. Chapter 12's mechanisms adapt, but not naively.

A **Frozen Snapshot of a quantum circuit** cannot mean freezing a cached state vector; it means exactly what Chapter 12, Section 12.3.2 meant for a classical model, applied to the classical description of a quantum computation: fix the gate sequence, the hardware calibration data, and the classical pseudorandom seed used to decide how many measurement shots to take and how to aggregate them, so that re-running the identical circuit on the identical calibrated hardware reproduces the identical measurement *statistics* — reproducible at the level of the experiment's classical description and its output distribution, not at the level of any single run's individual, irreducibly probabilistic outcome.

Hash-based routing for quantum states follows the same resolution: Chapter 24, Section 24.2.1's SHA-256 hash is computed over the *circuit description* — which gates, in which order, on which qubits, exactly as classical and exactly as hashable as any normalised text query — never over a quantum state itself, which has no stable classical hash to compute before it is measured and destroyed. A Static Vault entry for a quantum computation is, precisely, a cached, pre-verified *answer distribution* for a specific, hashed circuit, not a cached quantum state.

26.7.4 A Hybrid Architecture: Classical, Quantum, and Deterministic

Every credible near-term QNN application is a **hybrid architecture**: classical layers (Chapter 8) handle pre-processing (encoding classical data into a quantum circuit's initial state) and post-processing (interpreting measurement statistics back into a classical answer), with a quantum core sitting between them, anchored by exactly Section 26.7.3's circuit-hash-based islands wherever a specific, safety-relevant quantum computation's answer distribution needs to be reproducible rather than re-run. Chapter 12's Arbiter extends to this setting with one addition: alongside deciding *which* deterministic mechanism should answer a query, it now also decides whether a query needs the quantum core invoked at all, or resolves entirely

within the classical ILM Chapter 22 already built.

26.7.5 Potential Applications

Quantum chemistry simulation (naturally suited to QNNs, since simulating quantum systems is exactly what a quantum computer is structurally closest to doing efficiently) could eventually inform Volume I and Volume II's own material-science and neutron-cross-section questions. **Optimisation of complex systems** extends Chapter 17, Section 17.2's MPC and Section 17.3's RL to problems whose search space grows exponentially classically. **Cryptography and secure communication** cuts both ways for this book's own architecture: Chapter 24, Section 24.2.1's SHA-256 collision-resistance argument rests on 2^{128} effective security against the birthday bound; a sufficiently large fault-tolerant quantum computer applying Grover's algorithm would reduce a 256-bit hash's brute-force *preimage* search from 2^{256} to an effective 2^{128} — still astronomically outside any plausible attack, but the honest, specific reason NIST finalised post-quantum cryptographic standards (ML-KEM, ML-DSA, SLH-DSA, FN-DSA) in August 2024, already deployed in production TLS by major browsers and cloud providers as of this writing.

26.7.6 The Road Ahead

Current limitations are exactly Section 26.7.2's NISQ figures: no deployed system yet runs fault-tolerant, deep quantum circuits at a scale making Section 26.7.5's applications practical today. **The path to fault-tolerant quantum computing** runs through scaling today's tens-to-hundreds of logical qubits toward the thousands and millions larger algorithms require, with most credible industry roadmaps, as of 2026, placing production-scale fault tolerance in the 2030-2035 window. **The long-term vision** this section has been building toward is direct: Artificial Certainty, Chapter 1 through Chapter 25's entire subject, does not stop being relevant the moment a classical Transformer is replaced or augmented by a quantum circuit. The same problem — stochastic capability, deterministic guarantee where guarantees are required — simply recurs at the quantum layer, and Section 26.7.3 already showed the same category of solution recurs with it.

Honest Boundary

QNNs are in their infancy; deterministic islands at the quantum level are even more premature — but they are a natural extension of this book's philosophy. Every technique in Section 26.7.3 is a reasonable design sketch, not a deployed, evaluated, or certified mechanism; no chapter in this book has given quantum deterministic islands anything like Chapter 12's worked code or Chapter 24's mathematical proof, because the underlying quantum hardware and software ecosystem has not yet matured to the point where those would be more than speculation. This section is included because the *pattern* — bound the stochastic core, guarantee the narrow subset that must be guaranteed — is, this book argues, general enough to outlive the specific classical architecture Chapters 10 through 24 built it for, not because quantum Deterministic Islands are ready for a reader's next project.

26.8 The Long View: From Artificial Intelligence to Artificial Certainty

Chapter 1 opened with Shannon's entropy and a perceptron that could learn a straight line. Twenty-five chapters later, this book has built a system that learns far more than a straight line, and constrains what it learns, where constraint matters, with the same mathematical rigour Chapter 1 applied to a coding theorem. The distance between those two endpoints is this book's entire content; the connective thread running through all of it, restated one final time, is that **artificial intelligence and artificial certainty are not competing goals**. A system need not choose between being capable and being trustworthy. It must simply be built, as every chapter of this book has tried to show concretely rather than merely argue for, so that its capability and its guarantees live in different, clearly bounded parts of the architecture — each one doing the job only it can do, neither one asked to do the other's.

Honest Boundary

The future is uncertain; the physics is not. Every technological trajectory Section 26.6 and Section 26.7 described — grid modernisation, quantum error correction timelines, the pace of regulatory maturity Chapter 19 and Chapter 25 tracked even within the writing of this book — is a forecast, and forecasts of technology adoption have a long, humbling history of being wrong in both directions. What is not uncertain, and does not depend on any of these forecasts resolving one way or another, is the mathematics: Chapter 24's birthday bound, Chapter 7's backpropagation, Chapter 21's Born rule, and Chapter 2's Bayes' theorem will be exactly as true a decade from now as they are today, regardless of which specific products, regulations, or hardware platforms exist by then. This book has tried, throughout, to build its arguments on that second, steadier foundation wherever possible, and to be explicit, every time it could not, about which claims were resting on the first, less certain one instead.

From Certainty to Responsibility

This book began with Claude Shannon asking how much information a noisy channel could carry, and ended twenty-six chapters later with a quantum circuit's hash sitting in the same kind of lookup table as a pump's rated pressure. Between these two points lies the entire arc this volume set out to walk: entropy became the loss function a perceptron descends; the perceptron became the Transformer; gradient descent became backpropagation, and backpropagation became the engine training every model this book has built. None of that arc was optional scaffolding on the way to the real subject. It *was* the real subject, because the central claim this book has made — that a stochastic model can sit beside an absolute guarantee, in one architecture, without either compromising the other — only means anything to a reader who has seen, concretely, what the stochastic side is actually capable of and where, precisely, its guarantees run out. Chapter 12's Deterministic Island is not a hedge against everything that came before it in this book. It is the specific, engineered answer to the specific question those earlier chapters raised and left open.

That answer required more than one discipline to give. Chapter 1's information theory belongs to physics and mathematics. Chapters 3 through 9 belong to statistics and computer science. Chapters 10 through 19 belong to systems engineering, safety engineering, and regulatory law. Chapters 20 through 24 belong to software architecture. Chapter 14 and Chapter 25 belong, unambiguously, to ethics — not as a closing formality appended to a technical book, but as a discipline this book has treated as no less rigorous, and no less load-bearing, than any of the others. **No single one of these disciplines, alone, could have built Chapter 12's Deterministic Island.** A physicist without software architecture writes a beautiful proof no production system ever runs. A software engineer without domain physics builds a system that is well-structured and quietly wrong about the world it models. An engineer without an ethicist's discipline builds a system that works exactly as specified and still, in the end, cannot say who is accountable when it doesn't. This book's insistence on moving between these disciplines, chapter by chapter, rather than treating any one of them as sufficient, was not a stylistic choice. It was the only way to reach the destination honestly.

This is Volume III of the NEXUS-1 series, a planned twenty-volume project bridging foundational science with the practical engineering and computation that turns science into working, trustworthy systems. *From Grid to Core* and *From Core to Quantum* came before this volume; further volumes, including *From Flood to Cause* and *From Trial to Policy*, continue the same project into new domains, each one asking the same underlying question this volume asked of artificial intelligence: not merely "does it work," but "what, precisely, can we guarantee, and what have we honestly left as a probability instead."

None of this, finally, is only a technical undertaking. A Deterministic Island is a

piece of code, and it is also a promise: that a specific fact, asked about by a specific operator, on a specific shift, will be answered the same way every time, because someone, by name, verified it and stood behind that verification. Chapter 14's accountability argument and this epilogue's closing line are the same claim, stated twice. Building AI for critical infrastructure is a human responsibility before it is a technical achievement, and every mechanism in this book — every hash, every invariant, every audit event — exists only to make that responsibility possible to keep.

Mathematical Review

A compact reference for the mathematics this book assumed throughout, gathered in one place rather than re-derived at each point of use.

Vectors and Operations

A vector $\mathbf{v} \in \mathbb{R}^n$ is an ordered tuple of numbers. Addition and scalar multiplication are componentwise; the **dot product** $\mathbf{u} \cdot \mathbf{v} = \sum_i u_i v_i$ measures alignment and underlies Chapter 4's neuron; the **norm** $\|\mathbf{v}\| = \sqrt{\mathbf{v} \cdot \mathbf{v}}$ measures length and underlies every loss function in this book from Chapter 3 onward.

Matrices and Operations

A matrix $A \in \mathbb{R}^{m \times n}$ represents a linear map. Matrix **multiplication** $(AB)_{ij} = \sum_k A_{ik} B_{kj}$ composes two such maps (Chapter 6's layer composition). A square matrix's **inverse** A^{-1} , when it exists, satisfies $AA^{-1} = I$. An **eigenvalue** λ and **eigenvector** $\mathbf{v} \neq 0$ satisfy $A\mathbf{v} = \lambda\mathbf{v}$ — the directions a linear map only stretches, never rotates, and the basis Chapter 2, Section 2.5's multivariate Gaussian covariance ellipse is built from.

Complex Numbers

A complex number $z = a + bi$ ($i^2 = -1$) has **conjugate** $\bar{z} = a - bi$ and **modulus** $|z| = \sqrt{z\bar{z}} = \sqrt{a^2 + b^2}$. Chapter 21's quantum amplitudes are complex; the Born rule's probability $|z|^2 = z\bar{z}$ is exactly the modulus squared.

Differential Equations

An **ordinary differential equation** relates a function to its own derivatives; Chapter 2, Section 2.8's exponential decay, $dN/dt = -\lambda N$, solved by $N(t) = N_0 e^{-\lambda t}$, is the single differential equation this book leans on most heavily, underlying radioactive decay, RC circuit discharge, and any first-order decay process alike.

Probability

The three Kolmogorov axioms (Chapter 2, Section 2.2): probabilities are non-negative, the sample space has probability 1, and probabilities of disjoint events add. A **distribution** assigns probability (discrete) or density (continuous) across outcomes; the Gaussian, exponential, and Poisson distributions recur throughout Chapters 2 through 16.

Information Theory

Entropy $H(X) = -\sum_x P(x) \log_2 P(x)$ (Chapter 1) measures average uncertainty in bits. **Mutual information** $I(X; Y) = H(X) - H(X | Y)$ measures how much observing Y reduces uncertainty about X . **KL divergence** $D_{\text{KL}}(P \| Q) = \sum_x P(x) \log_2 \frac{P(x)}{Q(x)}$ (Chapter 2, Section 2.9) measures how one distribution diverges from a reference, and reappears as the regularising term in Chapter 8's VAE loss.

Table of Physical Constants

Constant	Symbol	Value
Speed of light (exact, SI)	c	299,792,458 m/s
Planck's constant (exact, SI)	h	$6.62607015 \times 10^{-34}$ J·s
Reduced Planck's constant	$\hbar = h/2\pi$	$1.054571817 \times 10^{-34}$ J·s
Elementary charge (exact, SI)	e	$1.602176634 \times 10^{-19}$ C
Electron mass	m_e	$9.1093837015 \times 10^{-31}$ kg
Proton mass	m_p	$1.67262192369 \times 10^{-27}$ kg
Neutron mass	m_n	$1.67492749804 \times 10^{-27}$ kg
Gravitational constant	G	6.67430×10^{-11} m ³ kg ⁻¹ s ⁻²
Boltzmann constant (exact, SI)	k_B	1.380649×10^{-23} J/K
Avogadro number (exact, SI)	N_A	$6.02214076 \times 10^{23}$ mol ⁻¹

Values marked "exact, SI" are exact by definition under the 2019 redefinition of the SI base units; all others are the current CODATA-recommended values, known to a precision far beyond what any worked example in this book required.

Keepin Six-Group Delayed Neutron Data

Reproduced from Volume I, *From Grid to Core*, for convenience, and referenced by Chapter 23’s Nuclear Context example (IDelayedNeutronDataRepository).

Group	$T_{1/2}$ (s)	$\lambda = \ln(2)/T_{1/2}$ (s ⁻¹)	Rel. Abundance (%)	Typical Precursors
1	55.72	0.01244	0.021	Br-87, Br-88, Kr-88
2	22.72	0.03051	0.142	I-137, I-138, Kr-89
3	6.22	0.11143	0.128	Br-89, Kr-90, Xe-139
4	2.30	0.30137	0.278	Xe-140, Kr-91
5	0.61	1.13617	0.080	Rb-93, Rb-94
6	0.23	3.01363	0.002	I-140, Cs-143

Architecture of a Deterministic Island

The query-resolution pipeline underlying Chapter 12's Static Vault, Chapter 22's Arbiter, and every worked example built from them, in one figure.




Figure 26.1: The complete query-resolution pipeline: every query is hashed, looked up against the Static Vault, and either resolved immediately and absolutely (a match) or handed to the stochastic generator under Arbiter, Causality Lock, and Human-in-the-Loop governance (no match) — the entire architecture of Chapters 12 and 22, in one diagram.

Comparison Table of AI Paradigms in Critical Systems

Paradigm	Strength	Weakness	Best Use Case
Rule-Based Systems	Deterministic, auditable	Brittle, cannot handle novelty	Safety interlocks
Classical ML	Interpretable, fast	Limited to structured data	Anomaly detection
Deep Learning	Powerful, flexible	Black box, data-hungry	Perception, prediction
LLMs / ILMs	General, language-native	Hallucinations, non-deterministic	Decision support, documentation
Deterministic Islands	Reproducible, certifiable	Limited scope, maintenance cost	Safety-critical queries
Quantum NN (with Islands)	Exponential capacity	Noise, immaturity	Future: chemistry, optimization

This table is a map, not a ranking: Part IV through Part VI argued at length that a real ILM combines several of these rows deliberately (Chapter 9, Section 9.7’s own version of this table; Chapter 26’s NeuralConstitution), rather than treating any single row as sufficient on its own.

Installation and Execution Guide for the Code (C#)

System Requirements

.NET 8 SDK or later. Any of Visual Studio 2022, JetBrains Rider, or Visual Studio Code with the C# extension. NuGet packages: MathNet.Numerics, xUnit, Microsoft.Extensions.Logging, Microsoft.ML (ML.NET, Chapter 23, Section 23.7), and System.Numerics (Chapter 21's Complex type).

Solution Structure

Every class this book built lives in exactly the Bounded Context (Chapter 20, Section 20.4.2) it belongs to — the solution layout itself is the DDD map:

AINexus.sln

```

1 /AINexus.sln
2   /src/
3     /QuantumContext/           (Chapter 21)
4       StateVector.cs Operator.cs Qubit.cs Measurement.cs
5     /NuclearContext/           (Volume I; Chapter 23, Section 23.1)
6       Isotope.cs Decay.cs Fission.cs
7     /Services/
8       DecayCalculationService.cs FissionFragmentService.cs
9     /ReactorContext/           (Volume I; Chapter 23, Section 23.1)
10      DelayedNeutronGroup.cs NeutronPopulation.cs ReactorCore.
11      cs
12    /Services/
13      FissionSimulationService.cs DelayedNeutronService.cs
14    /AIContext/                 (Chapter 22)
15      Tensor.cs Layer.cs Model.cs Prompt.cs Island.cs ILM.cs
16    /Services/
17      TrainingService.cs InferenceService.cs IslandService.
18      cs
19    /Infrastructure/            (Chapter 23, Section 23.3.2)
20      /Repositories/
21        IsotopeRepository.cs KeepinRepository.cs
22        ModelRepository.cs   IslandRepository.cs
23    /Application/              (Chapter 23, Sections 23.3.1, 23.5)
24      /Cli/      Program.cs
25      /Api/      Startup.cs
26    /tests/                (Chapter 23, Section 23.6)
27      QuantumContext.Tests/ NuclearContext.Tests/
28      ReactorContext.Tests/ AIContext.Tests/

```

Running the Examples

Shell

```
1 dotnet run --project src/Application/Cli -- infer \  
2   --prompt "What is the maximum allowed pressure?"  
3  
4 dotnet test
```

The first command exercises exactly Chapter 22, Section 22.8's end-to-end worked example, against the identical query, hash, and expected deterministic answer walked through there by hand.

GLOSSARY

Appendix A.3: Glossary of Terms

Each entry below states the term’s core meaning in one sentence, then adds the context, contrast, or consequence that makes it more than a dictionary definition — and points to the chapter where it was built from first principles.

Physics

- Quantum state** The complete description of a quantum system, represented as a vector in a complex Hilbert space (Chapter 21). Unlike a classical state (a definite position and velocity), a quantum state can be a superposition of several classical-looking outcomes at once, and only reveals one of them, probabilistically, upon measurement.
- Wavefunction** A quantum state’s representation in position (or another continuous) basis — the same physical state Chapter 21’s finite-dimensional `StateVector` represents for qubits, generalised to the continuous, infinite-dimensional case Chapter 21’s closing *Honest Boundary* explicitly excluded from this book’s scope.
- Operator** A linear map acting on a quantum state, representing a physical action performed on the system: a measurement, or a gate (Chapter 21). Chapter 21’s `Operator` class implements exactly this concept as a matrix multiplying a `StateVector`.
- Collapse** The irreversible update of a quantum state to a single, definite outcome upon measurement (Chapter 21, Section 21.4). Before collapse, a state may assign non-zero probability to several outcomes at once; after collapse, exactly one of them is real, and the information about the others is gone, not merely hidden.
- Entanglement** A joint quantum state of two or more systems that cannot be written as a product of each system’s own, independent state (Chapter 21). Chapter 21’s Bell-state worked example demonstrated this concretely: measuring one of two entangled qubits instantly determines what measuring the other will show, with no classical mechanism carrying that correlation between them.
- Decoherence** The loss of a quantum system’s coherent superposition through uncontrolled interaction with its environment (Chapter 26, Section 26.7.2). Decoherence is the single largest practical obstacle to building useful quantum computers, and the reason Chapter 26 classified current hardware as still being in the “NISQ” (Noisy Intermediate-Scale Quantum) era.
- Nucleus, isotope** The nucleus is an atom’s dense, central mass, built from protons

and neutrons; an isotope is one of several forms of a given chemical element, sharing the same proton count (and therefore the same chemistry) but differing in neutron count (and therefore in mass and nuclear stability). Chapter 23's `IsotopeRepository` treats isotopes as the basic data unit of the Nuclear Context.

Beta decay A radioactive decay process converting a neutron into a proton (or a proton into a neutron), with the emission of an electron or positron and a neutrino. Beta decay is the mechanism behind many of Appendix A.4's delayed-neutron precursor isotopes.

Fission, fusion Fission is the splitting of a heavy atomic nucleus into lighter fragments, releasing energy and additional neutrons; fusion is the joining of light nuclei into a heavier one, also releasing energy. Nuclear reactors, the physical system this entire series is ultimately written for, are built around a controlled, sustained fission chain reaction.

Delayed neutrons

Neutrons emitted not directly by fission itself, but by the subsequent radioactive decay of certain fission products, seconds to tens of seconds later (Appendix A.4). Delayed neutrons are what make a nuclear reactor's power level controllable on a human timescale at all; without them, reactor control would require responding on the sub-millisecond timescale of prompt neutrons alone.

AI/ML

Neuron A single computational unit computing a weighted sum of its inputs followed by a non-linear activation function (Chapter 4). Every architecture in this book, from Chapter 5's perceptron to Chapter 8's Transformer, is built by composing neurons according to different wiring patterns, not by replacing the neuron with something more exotic.

Perceptron The earliest trainable neuron, using a step-function activation and a simple mistake-driven update rule proven, by Novikoff's theorem, to converge in finitely many steps whenever the data is linearly separable (Chapter 5). Its inability to solve the XOR problem (Chapter 5, Section 5.6) is the historical reason research shifted toward the multi-layer networks Chapter 6 built.

Backpropagation

The algorithm that computes the gradient of a loss function with respect to every weight in a network, by applying the chain rule systematically backward through the network's computation graph (Chapter 7). Backpropagation is not a separate mathematical idea from ordinary calculus; it is the chain rule, made computationally efficient enough to scale to billions of parameters.

Gradient descent

The iterative optimisation rule that updates each parameter in

the direction that most reduces a loss function, scaled by a learning rate (Chapter 4). Every training process in this book, from Chapter 5's perceptron rule to Chapter 11's fine-tuning of the ILM, is a variation on this one idea.

- Overfitting** A model fitting the noise specific to its training data rather than the underlying signal that generalises to new data (Chapter 3). Chapter 3's bias-variance decomposition gives overfitting its precise mathematical name: excess variance, traded against insufficient bias.
- Regularization** Any technique that constrains a model during training specifically to reduce overfitting, such as an L1 or L2 penalty on weight magnitude (Chapter 3). Regularization is the general answer to the bias-variance tradeoff: deliberately accepting a small increase in bias in exchange for a larger reduction in variance.
- Transformer** The attention-based neural network architecture (Chapter 8, Section 8.5) underlying essentially every modern large language model, including this book's own ILM. Its defining advantage over the recurrent architectures that preceded it is that self-attention processes an entire sequence in parallel, rather than one timestep at a time.
- Attention** A mechanism that lets a model weight different parts of its input by learned relevance to the task at hand, computed freshly for every input rather than fixed in advance (Chapter 8, Section 8.4). Chapter 8's fully worked two-token example computed an actual attention weight matrix by hand, from query, key, and value projections.
- LLM, ILM** LLM: Large Language Model, a Transformer-based model trained on very large text corpora for general-purpose language understanding and generation. ILM: Industry Language Model, this book's own term (Chapter 10) for an LLM specifically adapted to an industrial domain and augmented with the Deterministic Islands (Chapter 12) that make it safe to consult for safety-relevant facts.
- Deterministic Island** A bounded mechanism guaranteeing identical, reproducible output for a defined, enumerable class of query, implemented without relying on the ILM's ordinary stochastic generation at all (Chapter 12). This book's central technical contribution, formalised mathematically in Chapter 24 and expressed in full DDD structure in Chapter 22.
- Hallucination** Fluent, confident, but factually ungrounded output from a language model — not a lie, since lying requires knowing the truth, but the statistically plausible continuation of a prompt with no internal distinction between "recalled fact" and "shape a plausible answer usually takes" (Chapter 10, Section 10.2.1).
- RAG** Retrieval-Augmented Generation: a technique that retrieves relevant documents from an external corpus at query time and inserts them into a language model's context before generation,

grounding the model's answer in specific, checkable source material rather than only its trained-in parameters (Chapter 11, Section 11.6).

Software Engineering

- OOP** Object-Oriented Programming: a programming paradigm built around classes (blueprints), objects (instances), encapsulation, inheritance, and polymorphism (Chapter 20, Section 20.3). This book's entire C# codebase, from Chapter 4's earliest neuron examples to Chapter 22's ILM aggregate, is written in this paradigm.
- DDD** Domain-Driven Design: a software design discipline for keeping code and domain expertise in permanent, checkable correspondence, built around a shared Ubiquitous Language and an explicit vocabulary of Entities, Value Objects, Aggregates, Repositories, Domain Services, and Domain Events (Chapter 20). Chapters 21 through 23 apply this vocabulary, in full, to three concrete Bounded Contexts.
- Entity** An object whose identity persists through change, as opposed to being defined by its current attribute values (Chapter 20, Section 20.4.3). Chapter 22's `Model` and `TrainingSession` classes are Entities: the same trained model remains "the same model" even after its weights are updated by further training.
- Value object** An object with no identity at all, defined entirely by its attributes, such that two value objects with identical attributes simply *are* the same value (Chapter 20, Section 20.4.3). Chapter 20's `Pressure` struct and Chapter 22's `Tensor`, `Prompt`, and `Island` types are all Value Objects.
- Aggregate, aggregate root** An aggregate is a cluster of related Entities and Value Objects treated as a single consistency boundary; the aggregate root is the one Entity through which every change to that cluster must pass, so that the aggregate's invariants can never be violated piecemeal (Chapter 20, Section 20.4.4). Chapter 22's ILM class is this book's central Aggregate Root.
- Repository** An abstraction that hides how an Aggregate is actually stored (in memory, in a file, in a database) behind a simple retrieve-and-save interface (Chapter 20, Section 20.4.5). Chapter 12's `IVectorStore` was a Repository three chapters before this book gave the pattern its name.
- Domain service** Logic that coordinates across several Entities or Value Objects and therefore does not belong naturally to any one of them (Chapter 20, Section 20.4.6). Chapter 12's `Arbiter` and `VectorResolver` are both Domain Services: deciding which mechanism should answer a query is a coordination rule, not a property of any single `ILMVector`.
- Domain event** A record that something significant has happened in the domain, published for other parts of the system to react to without the

event's source needing to know who, if anyone, is listening (Chapter 20, Section 20.4.7). Chapter 22's `IslandTriggered` event is this book's most consequential Domain Event, underlying Chapter 19's entire auditability argument.

Bounded context

The boundary within which a Ubiquitous Language, and the model built from it, applies consistently, with no obligation to share vocabulary across the boundary (Chapter 20, Section 20.4.2). This book builds three: the Quantum Context (Chapter 21), the Nuclear and Reactor Control Contexts (Volume I), and the AI Context (Chapter 22).

Ubiquitous language

A vocabulary shared, without translation, between domain experts and the developers writing code for them, so that a class named `Trip` means exactly what an operator means by "trip" (Chapter 20, Section 20.4.1).

Quantum Computing**Qubit**

A two-level quantum system, the basic unit of quantum information, capable of existing in a superposition of its two basis states rather than being confined to one or the other as a classical bit is (Chapter 21, Section 21.3). Chapter 21's `Qubit` class gives it identity within a `QuantumSystem`, separate from the `StateVector` value it contributes to.

Quantum gate

A unitary operator applied to one or more qubits, the quantum-computing analogue of a classical logic gate (Chapter 21). Chapter 21's worked example applied exactly two gates, Hadamard and CNOT, to build a complete entangled Bell state by hand.

Superposition

A quantum state expressed as a linear combination of two or more basis states at once, each with its own complex amplitude, rather than being in exactly one basis state (Chapter 21). Superposition is the source of Chapter 26, Section 26.7.1's exponential state-space capacity, and the reason a QNN's advantage is real but far more subtle than "trying every answer simultaneously."

Entanglement

See **Entanglement** under Physics, above; the term names an identical phenomenon in both contexts, since quantum computing is, at bottom, applied quantum mechanics.

Quantum algorithm

A computation expressed as a specific sequence of quantum gates, engineered so that constructive and destructive interference between amplitudes concentrates measurement probability on the desired answer (Chapter 26, Section 26.7.1). A quantum algorithm's difficulty lies almost entirely in designing this interference pattern, not in the superposition itself.

Quantum error correction

The technique of encoding one logical, robust qubit across many

physical, noisy qubits, so that errors on individual physical qubits can be detected and corrected before they corrupt the logical computation (Chapter 26, Section 26.7.2). Google's Willow processor's 2024 demonstration that logical error rates fall as more physical qubits are added is, as of this writing, the field's most significant milestone toward making this practical at scale.

INDEX

- Adam optimizer, [57](#)
- Aggregate (DDD), [135](#)
- Aggregate Root (DDD), [135](#)
- Alarm fatigue, [117](#)
- Arbiter, [92](#)
- Attention (mechanism), [62](#)
- AUC-ROC, [96](#)
- AUPRC, [116](#)

- Backpropagation, [50](#)
- Batch normalization, [57](#)
- Bayes' theorem, [12](#)
- Bell state, [141](#)
- Beta decay, [178](#)
- Birthday bound, [158](#)
- Born rule, [138](#)
- Bounded Context (DDD), [134](#)
- BPFO (Ball Pass Frequency, Outer race), [110](#)

- Causality Lock, [92](#)
- CNN, [59](#)
- CoolantLoop (class), [136](#)
- Crest factor, [109](#)
- Cross-validation, [21](#)

- DDD, [180](#)
- DDPG, [120](#)
- Decay heat, [150](#)
- Decoherence, [168](#)
- Delayed neutrons, [178](#)
- DeterminismMetadata, [79](#)
- Deterministic Island, [3](#)
- Digital twin, [122](#)
- Dissolved gas analysis, [111](#)
- Domain Event (DDD), [136](#)
- Domain Service (DDD), [136](#)
- Dropout, [57](#)

- EASA, [127](#)
- Entanglement, [141](#)
- Entity (DDD), [134](#)
- Entropy, [3](#)
- Equipment (class), [133](#)
- EU AI Act, [104](#)

- F1 score, [96](#)
- FDA, [128](#)
- Fine-tuning, [67](#)
- Fission, [150](#)
- Fusion, [178](#)

- GAN, [65](#)
- GMLP, [128](#)
- Gradient descent, [29](#)
- GRU, [62](#)

- Hallucination, [75](#)
- Human-in-the-Loop, [92](#)
- Human-on-the-loop, [105](#)

- IAEA, [127](#)
- IEC 61508, [100](#)
- ILM, [179](#)
- ILM (aggregate root class), [145](#)
- ILMVector, [79](#)
- InferenceSession (class), [144](#)
- Island (value object), [143](#)
- ISO 26262, [129](#)
- Isolation Forest, [115](#)
- Isotope, [150](#), [177](#)
- IVectorStore, [136](#)

- K-means clustering, [68](#)
- KL divergence, [17](#)
- Kurtosis, [109](#)

- Layer (value object), [143](#)
- LIME, [99](#)
- LLM, [179](#)
- Local Outlier Factor (LOF), [114](#)
- LSTM, [61](#)

- MAP estimation, [16](#)
- Measurement collapse, [140](#)
- MLE, [16](#)
- Model (class), [144](#)
- Model Predictive Control (MPC), [119](#)
- Mutual information, [4](#)

- Neural Constitution, [167](#)
- Neuron, [26](#)

NISQ, [169](#)
 NIST AI RMF, [128](#)
 No Free Lunch theorem, [24](#)
 No-cloning theorem, [169](#)
 Noisy Channel Coding Theorem, [4](#)
 Novikoff's theorem, [37](#)
 Nucleus, [177](#)

 One-Class SVM, [115](#)
 OOP, [180](#)
 Operator, [177](#)
 Overfitting, [20](#)

 PCA, [7](#)
 PCCP, [128](#)
 Perceptron, [32](#)
 PID control, [67](#)
 Policy gradient, [70](#)
 PPO, [120](#)
 Precision (metric), [116](#)
 Pressure (value object), [137](#)
 Prompt (value object), [143](#)
 Pump (class), [133](#)

 Q-learning, [70](#)
 Quantum algorithm, [181](#)
 Quantum error correction, [181](#)
 Quantum gate, [181](#)
 QuantumSystem (class), [139](#)
 Qubit, [139](#)

 RAG (Retrieval-Augmented Generation),
 [84](#)
 Recall (metric), [96](#)
 Regularization, [16](#)
 REINFORCE, [71](#)
 Repository (DDD), [136](#)
 RMS, [109](#)
 RNN, [61](#)
 RUL (Remaining Useful Life), [110](#)

 Safety case, [128](#)
 SCRAM setpoint, [161](#)
 SGD (Stochastic Gradient Descent), [30](#)
 Shannon, Claude, [2](#)
 SHAP, [99](#)
 SOTIF, [129](#)
 StateVector (class), [138](#)
 Static Vault, [89](#)
 Superposition, [138](#)

 Tensor (value object), [143](#)
 TPLC, [128](#)
 TrainingSession (class), [144](#)

 Transfer learning, [72](#)
 Transformer, [64](#)

 Ubiquitous Language (DDD), [134](#)
 Universal Approximation Theorem, [5](#)

 VAE, [65](#)
 Value Object (DDD), [134](#)
 VectorResolver, [100](#)

 Wavefunction, [177](#)

 XOR problem, [31](#)

SOURCES AND REFERENCE NOTES

Bibliography

Foundational Texts on Information and Probability

- Shannon, C. E. (1948). *A Mathematical Theory of Communication*. Bell System Technical Journal.
- Cover, T. M., & Thomas, J. A. (2006). *Elements of Information Theory*. Wiley-Interscience.
- MacKay, D. J. C. (2003). *Information Theory, Inference, and Learning Algorithms*. Cambridge University Press.

Machine Learning and Deep Learning

- Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- Murphy, K. P. (2022). *Probabilistic Machine Learning: An Introduction*. MIT Press.
- Vaswani, A., et al. (2017). *Attention Is All You Need*. NeurIPS.

Reinforcement Learning and Control

- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction*. MIT Press.
- Bertsekas, D. P. (2019). *Reinforcement Learning and Optimal Control*. Athena Scientific.

AI Safety, Ethics, and Governance

- Russell, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking.
- European Union (2024). *Regulation (EU) 2024/1689 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union.
- NIST (2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. National Institute of Standards and Technology.

Software Engineering and Domain-Driven Design

- Evans, E. (2003). *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Addison-Wesley.
- Vernon, V. (2013). *Implementing Domain-Driven Design*. Addison-Wesley.
- Fowler, M. (2002). *Patterns of Enterprise Application Architecture*. Addison-Wesley.
- Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1994). *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley.

Nuclear and Quantum Physics

- Krane, K. S. (1988). *Introductory Nuclear Physics*. Wiley.
- Griffiths, D. J. (2005). *Introduction to Quantum Mechanics*. Pearson.
- Keepin, G. R. (1965). *Physics of Nuclear Kinetics*. Addison-Wesley.
- Nielsen, M. A., & Chuang, I. L. (2010). *Quantum Computation and Quantum Information*. Cambridge University Press.

Companion Volumes in the NEXUS-1 Series

- Agathangelidis, G. K. *From Grid to Core* (Volume I).
- Agathangelidis, G. K. *From Core to Quantum* (Volume II).
- Agathangelidis, G. K. *The NEXUS-1 Domain-Driven Design Volumes*.